



**Python for Data
Analysis, 3rd Edition**

[Wes McKinney](#)

Published by [O'Reilly
Media, Inc.](#)

<https://github.com/wesm/pydata-book>

Chapter 8. Data Wrangling: Join, Combine, and Reshape

In many applications, data may be spread across a number of files or databases, or be arranged in a form that is not convenient to analyze. This chapter focuses on tools to help combine, join, and rearrange data.

First, I introduce the concept of *hierarchical indexing* in pandas, which is used extensively in some of these operations. I then dig into the particular data manipulations. You can see various applied usages of these tools in [Chapter 13](#).

8.1 Hierarchical Indexing

Hierarchical indexing is an important feature of pandas that enables you to have multiple (two or more) index *levels* on an axis. Another way of thinking about it is that it provides a way for you to work with higher dimensional data in a lower dimensional form. Let's start with a simple example: create a Series with a list of lists (or arrays) as the index:

In [11]: `data = pd.Series(np.random.uniform(size=9),`

```
....:     index=[["a", "a", "a", "b", "b", "c", "c", "d", "d"],  
....:             [1, 2, 3, 1, 3, 1, 2, 2, 3]])
```

In [12]: `data`

Out[12]:

a 1 0.929616

2 0.316376

3 0.183919

b 1 0.204560

3 0.567725

c 1 0.595545

2 0.964515

```
d 2 0.653177  
3 0.748907  
dtype: float64
```

What you’re seeing is a prettified view of a Series with a MultiIndex as its index. The “gaps” in the index display mean “use the label directly above”:

In [13]: `data.index`

Out[13]:

```
MultiIndex([( 'a', 1),
```

```
          ('a', 2),
```

```
          ('a', 3),
```

```
          ('b', 1),
```

```
          ('b', 3),
```

```
          ('c', 1),
```

```
          ('c', 2),
```

```
          ('d', 2),
```

```
          ('d', 3)],
```

```
)
```

With a hierarchically indexed object, so-called *partial* indexing is possible, enabling you to concisely select subsets of the data:

In [14]: `data["b"]`

Out[14]:

```
1 0.204560
```

```
3 0.567725
```

dtype: float64

In [15]: `data["b":"c"]`

```
Out[15]:
```

```
b 1  0.204560  
    3  0.567725  
c 1  0.595545  
    2  0.964515  
dtype: float64
```

```
In [16]: data.loc[["b", "d"]]
```

```
Out[16]:
```

```
b 1  0.204560  
    3  0.567725  
d 2  0.653177  
    3  0.748907  
dtype: float64
```

Selection is even possible from an “inner” level. Here I select all of the values having the value 2 from the second index level:

```
In [17]: data.loc[:, 2]
```

```
Out[17]:
```

```
a  0.316376  
c  0.964515  
d  0.653177  
dtype: float64
```

Hierarchical indexing plays an important role in reshaping data and in group-based operations like forming a pivot table. For example, you can rearrange this data into a DataFrame using its unstack method:

```
In [18]: data.unstack()
```

Out[18]:

```
1    2    3  
a  0.929616 0.316376 0.183919  
b  0.204560  NaN  0.567725  
c  0.595545 0.964515  NaN  
d  NaN  0.653177 0.748907
```

The inverse operation of unstack is stack:

In [19]: data.unstack().stack()

Out[19]:

```
a 1  0.929616  
   2  0.316376  
   3  0.183919  
  
b 1  0.204560  
   3  0.567725  
  
c 1  0.595545  
   2  0.964515  
  
d 2  0.653177  
   3  0.748907  
  
dtype: float64
```

stack and unstack will be explored in more detail later in [Section 8.3, “Reshaping and Pivoting.”](#)

With a DataFrame, either axis can have a hierarchical index:

```
In [20]: frame = pd.DataFrame(np.arange(12).reshape((4, 3)),  
.....           index=[["a", "a", "b", "b"], [1, 2, 1, 2]],  
.....           columns=[["Ohio", "Ohio", "Colorado"]],
```

```
....: ["Green", "Red", "Green"]])
```

In [21]: frame

Out[21]:

```
Ohio  Colorado
Green Red  Green
a 1  0  1  2
      2  3  4  5
b 1  6  7  8
      2  9 10  11
```

The hierarchical levels can have names (as strings or any Python objects). If so, these will show up in the console output:

In [22]: frame.index.names = ["key1", "key2"]

In [23]: frame.columns.names = ["state", "color"]

In [24]: frame

Out[24]:

```
state  Ohio  Colorado
color  Green Red  Green
key1 key2
a  1  0  1  2
      2  3  4  5
b  1  6  7  8
      2  9 10  11
```

These names supersede the name attribute, which is used only with single-level indexes.

Caution

Be careful to note that the index names "state" and "color" are not part of the row labels (the frame.index values).

You can see how many levels an index has by accessing its nlevels attribute:

In [25]: frame.index.nlevels

Out[25]: 2

With partial column indexing you can similarly select groups of columns:

In [26]: frame["Ohio"]

Out[26]:

color Green Red

key1 key2

a 1 0 1

2 3 4

b 1 6 7

2 9 10

A MultiIndex can be created by itself and then reused; the columns in the preceding DataFrame with level names could also be created like this:

```
pd.MultiIndex.from_arrays([["Ohio", "Ohio", "Colorado"],
```

```
    ["Green", "Red", "Green"]],
```

```
    names=["state", "color"])
```

Reordering and Sorting Levels

At times you may need to rearrange the order of the levels on an axis or sort the data by the values in one specific level. The swaplevel method takes two level numbers or names and returns a new object with the levels interchanged (but the data is otherwise unaltered):

In [27]: frame.swaplevel("key1", "key2")

Out[27]:

```
state  Ohio  Colorado
color  Green Red  Green
key2 key1
1  a    0 1    2
2  a    3 4    5
1  b    6 7    8
2  b    9 10   11
```

sort_index by default sorts the data lexicographically using all the index levels, but you can choose to use only a single level or a subset of levels to sort by passing the level argument. For example:

In [28]: frame.sort_index(level=1)

Out[28]:

```
state  Ohio  Colorado
color  Green Red  Green
key1 key2
a  1    0 1    2
b  1    6 7    8
a  2    3 4    5
b  2    9 10   11
```

In [29]: frame.swaplevel(0, 1).sort_index(level=0)

Out[29]:

```
state  Ohio  Colorado
color  Green Red  Green
```

```
key2 key1  
1  a    0  1    2  
    b    6  7    8  
2  a    3  4    5  
    b    9 10   11
```

Note

Data selection performance is much better on hierarchically indexed objects if the index is lexicographically sorted starting with the outermost level—that is, the result of calling `sort_index(level=0)` or `sort_index()`.

Summary Statistics by Level

Many descriptive and summary statistics on DataFrame and Series have a `level` option in which you can specify the level you want to aggregate by on a particular axis. Consider the above DataFrame; we can aggregate by level on either the rows or columns, like so:

In [30]: `frame.groupby(level="key2").sum()`

Out[30]:

```
state Ohio  Colorado
```

```
color Green Red  Green
```

```
key2
```

```
1    6  8   10
```

```
2   12 14   16
```

In [31]: `frame.groupby(level="color", axis="columns").sum()`

Out[31]:

```
color  Green Red
```

```
key1 key2
```

```
a  1    2  1
```

```
2    8  4
```

```
b  1    14  7
```

```
2   20 10
```

We will discuss groupby in much more detail later in [Chapter 10](#).

Indexing with a DataFrame's columns

It's not unusual to want to use one or more columns from a DataFrame as the row index; alternatively, you may wish to move the row index into the DataFrame's columns. Here's an example DataFrame:

```
In [32]: frame = pd.DataFrame({"a": range(7), "b": range(7, 0, -1),
....:                 "c": ["one", "one", "one", "two", "two",
....:                     "two", "two"],
....:                 "d": [0, 1, 2, 0, 1, 2, 3]})
```

```
In [33]: frame
```

```
Out[33]:
```

```
   a  b  c  d
0  0  0  one  0
1  1  1  one  1
2  2  2  one  2
3  3  3  two  0
4  4  4  two  1
5  5  5  two  2
6  6  6  two  3
```

DataFrame's `set_index` function will create a new DataFrame using one or more of its columns as the index:

```
In [34]: frame2 = frame.set_index(["c", "d"])
```

```
In [35]: frame2
```

```
Out[35]:
```

```
a b  
c d  
one 0 0 7  
    1 1 6  
    2 2 5  
two 0 3 4  
    1 4 3  
    2 5 2  
    3 6 1
```

By default, the columns are removed from the DataFrame, though you can leave them in by passing drop=False to set_index:

```
In [36]: frame.set_index(["c", "d"], drop=False)
```

```
Out[36]:
```

```
a b c d  
c d  
one 0 0 7 one 0  
    1 1 6 one 1  
    2 2 5 one 2  
two 0 3 4 two 0  
    1 4 3 two 1  
    2 5 2 two 2  
    3 6 1 two 3
```

reset_index, on the other hand, does the opposite of set_index; the hierarchical index levels are moved into the columns:

```
In [37]: frame2.reset_index()
```

```
Out[37]:
```

```
c d a b  
0 one 0 0 7  
1 one 1 1 6  
2 one 2 2 5  
3 two 0 3 4  
4 two 1 4 3  
5 two 2 5 2  
6 two 3 6 1
```

8.2 Combining and Merging Datasets

Data contained in pandas objects can be combined in a number of ways:

pandas.merge

Connect rows in DataFrames based on one or more keys. This will be familiar to users of SQL or other relational databases, as it implements database *join* operations.

pandas.concat

Concatenate or “stack” objects together along an axis.

combine_first

Splice together overlapping data to fill in missing values in one object with values from another.

I will address each of these and give a number of examples. They’ll be utilized in examples throughout the rest of the book.

Database-Style DataFrame Joins

Merge or *join* operations combine datasets by linking rows using one or more *keys*. These operations are particularly important in relational databases (e.g., SQL-based).

The *pandas.merge* function in pandas is the main entry point for using these algorithms on your data.

Let’s start with a simple example:

```
In [38]: df1 = pd.DataFrame({"key": ["b", "b", "a", "c", "a", "a", "b"],  
....: "data1": pd.Series(range(7), dtype="Int64")})
```

```
In [39]: df2 = pd.DataFrame({"key": ["a", "b", "d"],
```

```
....:         "data2": pd.Series(range(3), dtype="Int64")})
```

In [40]: df1

Out[40]:

```
key  data1  
0   b   0  
1   b   1  
2   a   2  
3   c   3  
4   a   4  
5   a   5  
6   b   6
```

In [41]: df2

Out[41]:

```
key  data2  
0   a   0  
1   b   1  
2   d   2
```

Here I am using pandas's Int64 extension type for nullable integers, discussed in [Section 7.3, “Extension Data Types.”](#)

This is an example of a *many-to-one* join; the data in df1 has multiple rows labeled a and b, whereas df2 has only one row for each value in the key column. Calling `pandas.merge` with these objects, we obtain:

In [42]: `pd.merge(df1, df2)`

Out[42]:

```
key  data1  data2  
0   b     0     1  
1   b     1     1
```

```
2 b 6 1
3 a 2 0
4 a 4 0
5 a 5 0
```

Note that I didn't specify which column to join on. If that information is not specified, pandas.merge uses the overlapping column names as the keys. It's a good practice to specify explicitly, though:

```
In [43]: pd.merge(df1, df2, on="key")
```

```
Out[43]:
```

```
key data1 data2
0 b 0 1
1 b 1 1
2 b 6 1
3 a 2 0
4 a 4 0
5 a 5 0
```

In general, the order of column output in pandas.merge operations is unspecified.

If the column names are different in each object, you can specify them separately:

```
In [44]: df3 = pd.DataFrame({"lkey": ["b", "b", "a", "c", "a", "a", "b"],
....: "data1": pd.Series(range(7), dtype="Int64")})
```

```
In [45]: df4 = pd.DataFrame({"rkey": ["a", "b", "d"],
....: "data2": pd.Series(range(3), dtype="Int64")})
```

```
In [46]: pd.merge(df3, df4, left_on="lkey", right_on="rkey")
```

```
Out[46]:
```

```
lkey data1 rkey data2  
0 b 0 b 1  
1 b 1 b 1  
2 b 6 b 1  
3 a 2 a 0  
4 a 4 a 0  
5 a 5 a 0
```

You may notice that the "c" and "d" values and associated data are missing from the result. By default, pandas.merge does an "inner" join; the keys in the result are the intersection, or the common set found in both tables. Other possible options are "left", "right", and "outer". The outer join takes the union of the keys, combining the effect of applying both left and right joins:

In [47]: pd.merge(df1, df2, how="outer")

Out[47]:

```
key data1 data2  
0 b 0 1  
1 b 1 1  
2 b 6 1  
3 a 2 0  
4 a 4 0  
5 a 5 0  
6 c 3 <NA>  
7 d <NA> 2
```

In [48]: pd.merge(df3, df4, left_on="lkey", right_on="rkey", how="outer")

Out[48]:

```
lkey data1 rkey data2  
0 b 0 b 1  
1 b 1 b 1
```

```
2  b   6  b   1
3  a   2  a   0
4  a   4  a   0
5  a   5  a   0
6  c   3  NaN <NA>
7  NaN <NA>  d   2
```

In an outer join, rows from the left or right DataFrame objects that do not match on keys in the other DataFrame will appear with NA values in the other DataFrame's columns for the nonmatching rows.

See [Table 8-1](#) for a summary of the options for how.

Option	Behavior
how="inner"	Use only the key combinations observed in both tables
how="left"	Use all key combinations found in the left table
how="right"	Use all key combinations found in the right table
how="outer"	Use all key combinations observed in both tables together

Table 8-1. Different join types with the how argument

Many-to-many merges form the Cartesian product of the matching keys. Here's an example:

```
In [49]: df1 = pd.DataFrame({"key": ["b", "b", "a", "c", "a", "b"],
....:                 "data1": pd.Series(range(6), dtype="Int64")})
```

```
In [50]: df2 = pd.DataFrame({"key": ["a", "b", "a", "b", "d"],
....:                 "data2": pd.Series(range(5), dtype="Int64")})
```

```
In [51]: df1
```

```
Out[51]:
```

```
key data1  
0 b 0  
1 b 1  
2 a 2  
3 c 3  
4 a 4  
5 b 5
```

```
In [52]: df2
```

```
Out[52]:
```

```
key data2  
0 a 0  
1 b 1  
2 a 2  
3 b 3  
4 d 4
```

```
In [53]: pd.merge(df1, df2, on="key", how="left")
```

```
Out[53]:
```

```
key data1 data2  
0 b 0 1  
1 b 0 3  
2 b 1 1  
3 b 1 3  
4 a 2 0  
5 a 2 2  
6 c 3 <NA>  
7 a 4 0  
8 a 4 2
```

```
9  b  5  1  
10 b  5  3
```

Since there were three "b" rows in the left DataFrame and two in the right one, there are six "b" rows in the result. The join method passed to the how keyword argument affects only the distinct key values appearing in the result:

```
In [54]: pd.merge(df1, df2, how="inner")
```

```
Out[54]:
```

	key	data1	data2
0	b	0	1
1	b	0	3
2	b	1	1
3	b	1	3
4	b	5	1
5	b	5	3
6	a	2	0
7	a	2	2
8	a	4	0
9	a	4	2

To merge with multiple keys, pass a list of column names:

```
In [55]: left = pd.DataFrame({"key1": ["foo", "foo", "bar"],  
....:                 "key2": ["one", "two", "one"],  
....:                 "lval": pd.Series([1, 2, 3], dtype='Int64')})
```

```
In [56]: right = pd.DataFrame({"key1": ["foo", "foo", "bar", "bar"],  
....:                 "key2": ["one", "one", "one", "two"],  
....:                 "rval": pd.Series([4, 5, 6, 7], dtype='Int64')})
```

```
In [57]: pd.merge(left, right, on=["key1", "key2"], how="outer")
```

```
Out[57]:
```

	key1	key2	lval	rval
0	foo	one	1	4
1	foo	one	1	5
2	foo	two	2	<NA>
3	bar	one	3	6
4	bar	two	<NA>	7

To determine which key combinations will appear in the result depending on the choice of merge method, think of the multiple keys as forming an array of tuples to be used as a single join key.

Caution

When you're joining columns on columns, the indexes on the passed DataFrame objects are discarded. If you need to preserve the index values, you can use `reset_index` to append the index to the columns.

A last issue to consider in merge operations is the treatment of overlapping column names. For example:

```
In [58]: pd.merge(left, right, on="key1")
```

```
Out[58]:
```

	key1	key2_x	lval	key2_y	rval
0	foo	one	1	one	4
1	foo	one	1	one	5
2	foo	two	2	one	4
3	foo	two	2	one	5
4	bar	one	3	one	6
5	bar	one	3	two	7

While you can address the overlap manually (see the section “[Renaming Axis Indexes](#)” for renaming axis labels), pandas.merge has a suffixes option for specifying strings to append to overlapping names in the left and right DataFrame objects:

```
In [59]: pd.merge(left, right, on="key1", suffixes=("_left", "_right"))
```

Out[59]:

```
key1 key2_left lval key2_right rval
0 foo    one   1    one   4
1 foo    one   1    one   5
2 foo    two   2    one   4
3 foo    two   2    one   5
4 bar    one   3    one   6
5 bar    one   3    two   7
```

See [Table 8-2](#) for an argument reference on pandas.merge. The next section covers joining using the DataFrame’s row index.

Argument	Description
left	DataFrame to be merged on the left side.
right	DataFrame to be merged on the right side.
how	Type of join to apply: one of "inner", "outer", "left", or "right"; defaults to "inner".
on	Column names to join on. Must be found in both DataFrame objects. If not specified and no other join keys given, will use the intersection of the column names in left and right as the join keys.
left_on	Columns in left DataFrame to use as join keys. Can be a single column name or a list of column names.

Argument	Description
right_on	Analogous to left_on for right DataFrame.
left_index	Use row index in left as its join key (or keys, if a MultiIndex).
right_index	Analogous to left_index.
sort	Sort merged data lexicographically by join keys; False by default.
suffixes	Tuple of string values to append to column names in case of overlap; defaults to ("_x", "_y") (e.g., if "data" in both DataFrame objects, would appear as "data_x" and "data_y" in result).
copy	If False, avoid copying data into resulting data structure in some exceptional cases; by default always copies.
validate	Verifies if the merge is of the specified type, whether one-to-one, one-to-many, or many-to-many. See the docstring for full details on the options.
indicator	Adds a special column _merge that indicates the source of each row; values will be "left_only", "right_only", or "both" based on the origin of the joined data in each row.

Table 8-2. pandas.merge function arguments

Merging on Index

In some cases, the merge key(s) in a DataFrame will be found in its index (row labels). In this case, you can pass left_index=True or right_index=True (or both) to indicate that the index should be used as the merge key:

```
In [60]: left1 = pd.DataFrame({"key": ["a", "b", "a", "a", "b", "c"],
....:                 "value": pd.Series(range(6), dtype="Int64")})
```

```
In [61]: right1 = pd.DataFrame({"group_val": [3.5, 7]}, index=["a", "b"])
```

```
In [62]: left1
```

```
Out[62]:
```

key	value
0	a 0
1	b 1
2	a 2
3	a 3
4	b 4
5	c 5

0	a 0
1	b 1
2	a 2
3	a 3
4	b 4
5	c 5

1	b 1
2	a 2
3	a 3
4	b 4
5	c 5

2	a 2
3	a 3
4	b 4
5	c 5

3	a 3
4	b 4
5	c 5

4	b 4
5	c 5

5	c 5
---	-----

```
In [63]: right1
```

```
Out[63]:
```

group_val
a 3.5
b 7.0

a 3.5
b 7.0

b 7.0

```
In [64]: pd.merge(left1, right1, left_on="key", right_index=True)
```

```
Out[64]:
```

key	value	group_val
-----	-------	-----------

0	a 0	3.5
---	-----	-----

2	a 2	3.5
---	-----	-----

3	a 3	3.5
---	-----	-----

1	b 1	7.0
---	-----	-----

4	b 4	7.0
---	-----	-----

Note

If you look carefully here, you will see that the index values for left1 have been preserved, whereas in other examples above, the indexes of the input DataFrame objects are dropped. Because the index of right1 is unique, this “many-to-one” merge (with the

default how="inner" method) can preserve the index values from left1 that correspond to rows in the output.

Since the default merge method is to intersect the join keys, you can instead form the union of them with an outer join:

```
In [65]: pd.merge(left1, right1, left_on="key", right_index=True, how="outer")
```

```
Out[65]:
```

	key	value	group_val
0	a	0	3.5
2	a	2	3.5
3	a	3	3.5
1	b	1	7.0
4	b	4	7.0
5	c	5	NaN

With hierarchically indexed data, things are more complicated, as joining on index is equivalent to a multiple-key merge:

```
In [66]: lefth = pd.DataFrame({"key1": ["Ohio", "Ohio", "Ohio",
.....:                 "Nevada", "Nevada"],
.....:                 "key2": [2000, 2001, 2002, 2001, 2002],
.....:                 "data": pd.Series(range(5), dtype="Int64")})
```

```
In [67]: righth_index = pd.MultiIndex.from_arrays(
```

```
.....: [
.....:     ["Nevada", "Nevada", "Ohio", "Ohio", "Ohio", "Ohio"],
.....:     [2001, 2000, 2000, 2000, 2001, 2002]
.....: ]
.....: )
```

```
In [68]: righth = pd.DataFrame({"event1": pd.Series([0, 2, 4, 6, 8, 10], dtype="I
nt64",
```

```
....:                         index=righth_index),  
....:             "event2": pd.Series([1, 3, 5, 7, 9, 11], dtype="I  
nt64",  
....:                         index=righth_index)})
```

In [69]: lefth

Out[69]:

```
key1 key2 data  
0 Ohio 2000 0  
1 Ohio 2001 1  
2 Ohio 2002 2  
3 Nevada 2001 3  
4 Nevada 2002 4
```

In [70]: righth

Out[70]:

```
event1 event2  
Nevada 2001 0 1  
      2000 2 3  
Ohio 2000 4 5  
      2000 6 7  
      2001 8 9  
      2002 10 11
```

In this case, you have to indicate multiple columns to merge on as a list (note the handling of duplicate index values with how="outer"):

In [71]: pd.merge(lefth, righth, left_on=["key1", "key2"], right_index=True)

Out[71]:

```
key1 key2 data event1 event2
```

```
0 Ohio 2000 0 4 5
0 Ohio 2000 0 6 7
1 Ohio 2001 1 8 9
2 Ohio 2002 2 10 11
3 Nevada 2001 3 0 1
```

In [72]: pd.merge(left, right, left_on=["key1", "key2"],

....: right_index=True, how="outer")

Out[72]:

```
key1 key2 data event1 event2
0 Ohio 2000 0 4 5
0 Ohio 2000 0 6 7
1 Ohio 2001 1 8 9
2 Ohio 2002 2 10 11
3 Nevada 2001 3 0 1
4 Nevada 2002 4 <NA> <NA>
4 Nevada 2000 <NA> 2 3
```

Using the indexes of both sides of the merge is also possible:

In [73]: left2 = pd.DataFrame([[1., 2.], [3., 4.], [5., 6.]],

....: index=["a", "c", "e"],

....: columns=["Ohio", "Nevada"]).astype("Int64")

In [74]: right2 = pd.DataFrame([[7., 8.], [9., 10.], [11., 12.], [13, 14]],

....: index=["b", "c", "d", "e"],

....: columns=["Missouri", "Alabama"]).astype("Int64")

In [75]: left2

Out[75]:

```
Ohio Nevada
```

```
a 1 2  
c 3 4  
e 5 6
```

```
In [76]: right2
```

```
Out[76]:
```

```
Missouri Alabama
```

```
b 7 8  
c 9 10  
d 11 12  
e 13 14
```

```
In [77]: pd.merge(left2, right2, how="outer", left_index=True, right_index=True)
```

```
Out[77]:
```

```
Ohio Nevada Missouri Alabama
```

```
a 1 2 <NA> <NA>  
b <NA> <NA> 7 8  
c 3 4 9 10  
d <NA> <NA> 11 12  
e 5 6 13 14
```

DataFrame has a join instance method to simplify merging by index. It can also be used to combine many DataFrame objects having the same or similar indexes but nonoverlapping columns. In the prior example, we could have written:

```
In [78]: left2.join(right2, how="outer")
```

```
Out[78]:
```

```
Ohio Nevada Missouri Alabama
```

```
a 1 2 <NA> <NA>  
b <NA> <NA> 7 8
```

```
c  3   4   9   10  
d <NA> <NA>  11   12  
e  5   6   13   14
```

Compared with pandas.merge, DataFrame's join method performs a left join on the join keys by default. It also supports joining the index of the passed DataFrame on one of the columns of the calling DataFrame:

```
In [79]: left1.join(right1, on="key")
```

```
Out[79]:
```

	key	value	group_val
0	a	0	3.5
1	b	1	7.0
2	a	2	3.5
3	a	3	3.5
4	b	4	7.0
5	c	5	NaN

You can think of this method as joining data “into” the object whose join method was called.

Lastly, for simple index-on-index merges, you can pass a list of DataFrames to join as an alternative to using the more general pandas.concat function described in the next section:

```
In [80]: another = pd.DataFrame([[7., 8.], [9., 10.], [11., 12.], [16., 17.]],  
....:           index=["a", "c", "e", "f"],  
....:           columns=["New York", "Oregon"])
```

```
In [81]: another
```

```
Out[81]:
```

```
New York Oregon
```

```
a    7.0  8.0  
c    9.0  10.0  
e   11.0 12.0  
f   16.0 17.0
```

In [82]: `left2.join([right2, another])`

Out[82]:

```
Ohio Nevada Missouri Alabama New York Oregon  
a  1    2    <NA>  <NA>    7.0  8.0  
c  3    4    9    10    9.0  10.0  
e  5    6    13   14   11.0 12.0
```

In [83]: `left2.join([right2, another], how="outer")`

Out[83]:

```
Ohio Nevada Missouri Alabama New York Oregon  
a  1    2    <NA>  <NA>    7.0  8.0  
c  3    4    9    10    9.0  10.0  
e  5    6    13   14   11.0 12.0  
b <NA>  <NA>    7    8    NaN  NaN  
d <NA>  <NA>    11   12    NaN  NaN  
f <NA>  <NA>  <NA>  <NA>  16.0 17.0
```

Concatenating Along an Axis

Another kind of data combination operation is referred to interchangeably as *concatenation* or *stacking*. NumPy's `concatenate` function can do this with NumPy arrays:

In [84]: `arr = np.arange(12).reshape((3, 4))`

In [85]: `arr`

Out[85]:

```
array([[ 0,  1,  2,  3],  
       [ 4,  5,  6,  7],  
       [ 8,  9, 10, 11]])
```

In [86]: np.concatenate([arr, arr], axis=1)

Out[86]:

```
array([[ 0,  1,  2,  3,  0,  1,  2,  3],  
       [ 4,  5,  6,  7,  4,  5,  6,  7],  
       [ 8,  9, 10, 11,  8,  9, 10, 11]])
```

In the context of pandas objects such as Series and DataFrame, having labeled axes enable you to further generalize array concatenation. In particular, you have a number of additional concerns:

- If the objects are indexed differently on the other axes, should we combine the distinct elements in these axes or use only the values in common?
- Do the concatenated chunks of data need to be identifiable as such in the resulting object?
- Does the “concatenation axis” contain data that needs to be preserved? In many cases, the default integer labels in a DataFrame are best discarded during concatenation.

The concat function in pandas provides a consistent way to address each of these questions. I’ll give a number of examples to illustrate how it works. Suppose we have three Series with no index overlap:

In [87]: s1 = pd.Series([0, 1], index=["a", "b"], dtype="Int64")

In [88]: s2 = pd.Series([2, 3, 4], index=["c", "d", "e"], dtype="Int64")

In [89]: s3 = pd.Series([5, 6], index=["f", "g"], dtype="Int64")

Calling pandas.concat with these objects in a list glues together the values and indexes:

```
In [90]: s1
```

```
Out[90]:
```

```
a 0
```

```
b 1
```

```
dtype: Int64
```

```
In [91]: s2
```

```
Out[91]:
```

```
c 2
```

```
d 3
```

```
e 4
```

```
dtype: Int64
```

```
In [92]: s3
```

```
Out[92]:
```

```
f 5
```

```
g 6
```

```
dtype: Int64
```

```
In [93]: pd.concat([s1, s2, s3])
```

```
Out[93]:
```

```
a 0
```

```
b 1
```

```
c 2
```

```
d 3
```

```
e 4
```

```
f 5
```

```
g 6
```

```
dtype: Int64
```

By default, pandas.concat works along axis="index", producing another Series. If you pass axis="columns", the result will instead be a DataFrame:

In [94]: pd.concat([s1, s2, s3], axis="columns")

Out[94]:

```
0  1  2  
a  0 <NA> <NA>  
b  1 <NA> <NA>  
c <NA>  2 <NA>  
d <NA>  3 <NA>  
e <NA>  4 <NA>  
f <NA> <NA>  5  
g <NA> <NA>  6
```

In this case there is no overlap on the other axis, which as you can see is the union (the "outer" join) of the indexes. You can instead intersect them by passing join="inner":

In [95]: s4 = pd.concat([s1, s3])

In [96]: s4

Out[96]:

```
a  0  
b  1  
f  5  
g  6  
dtype: Int64
```

In [97]: pd.concat([s1, s4], axis="columns")

Out[97]:

```
0  1  
a  0 0
```

```
b 1 1  
f <NA> 5  
g <NA> 6
```

In [98]: `pd.concat([s1, s4], axis="columns", join="inner")`

Out[98]:

```
0 1  
a 0 0  
b 1 1
```

In this last example, the "f" and "g" labels disappeared because of the `join="inner"` option.

A potential issue is that the concatenated pieces are not identifiable in the result.

Suppose instead you wanted to create a hierarchical index on the concatenation axis. To do this, use the `keys` argument:

In [99]: `result = pd.concat([s1, s1, s3], keys=["one", "two", "three"])`

In [100]: `result`

Out[100]:

```
one  a  0  
     b  1  
two  a  0  
     b  1  
three f  5  
      g  6
```

`dtype: Int64`

In [101]: `result.unstack()`

Out[101]:

```
   a   b   f   g  
one  0   1 <NA> <NA>
```

```
two    0  1 <NA> <NA>
three <NA> <NA>  5  6
```

In the case of combining Series along axis="columns", the keys become the DataFrame column headers:

```
In [102]: pd.concat([s1, s2, s3], axis="columns", keys=["one", "two", "three"])
```

```
Out[102]:
```

```
   one  two  three
a  0  <NA> <NA>
b  1  <NA> <NA>
c <NA>  2  <NA>
d <NA>  3  <NA>
e <NA>  4  <NA>
f <NA> <NA>  5
g <NA> <NA>  6
```

The same logic extends to DataFrame objects:

```
In [103]: df1 = pd.DataFrame(np.arange(6).reshape(3, 2), index=["a", "b", "c"],
.....:           columns=["one", "two"])
```

```
In [104]: df2 = pd.DataFrame(5 + np.arange(4).reshape(2, 2), index=["a", "c"],
.....:           columns=["three", "four"])
```

```
In [105]: df1
```

```
Out[105]:
```

```
   one  two
a  0  1
```

```
b 2 3  
c 4 5
```

In [106]: df2

Out[106]:

```
three four  
a 5 6  
c 7 8
```

In [107]: pd.concat([df1, df2], axis="columns", keys=["level1", "level2"])

Out[107]:

```
level1  level2  
one two three four  
a 0 1 5.0 6.0  
b 2 3 NaN NaN  
c 4 5 7.0 8.0
```

Here the keys argument is used to create a hierarchical index where the first level can be used to identify each of the concatenated DataFrame objects.

If you pass a dictionary of objects instead of a list, the dictionary's keys will be used for the keys option:

In [108]: pd.concat({"level1": df1, "level2": df2}, axis="columns")

Out[108]:

```
level1  level2  
one two three four  
a 0 1 5.0 6.0  
b 2 3 NaN NaN  
c 4 5 7.0 8.0
```

There are additional arguments governing how the hierarchical index is created (see [Table 8-3](#)). For example, we can name the created axis levels with the names argument:

```
In [109]: pd.concat([df1, df2], axis="columns", keys=["level1", "level2"],
```

```
.....:      names=["upper", "lower"])
```

```
Out[109]:
```

```
upper level1  level2  
lower  one two three four  
a      0  1  5.0 6.0  
b      2  3  NaN NaN  
c      4  5  7.0 8.0
```

A last consideration concerns DataFrames in which the row index does not contain any relevant data:

```
In [110]: df1 = pd.DataFrame(np.random.standard_normal((3, 4)),
```

```
.....:      columns=["a", "b", "c", "d"])
```

```
In [111]: df2 = pd.DataFrame(np.random.standard_normal((2, 3)),
```

```
.....:      columns=["b", "d", "a"])
```

```
In [112]: df1
```

```
Out[112]:
```

```
      a      b      c      d  
0  1.248804  0.774191 -0.319657 -0.624964  
1  1.078814  0.544647  0.855588  1.343268  
2 -0.267175  1.793095 -0.652929 -1.886837
```

```
In [113]: df2
```

```
Out[113]:
```

```
b      d      a  
0 1.059626 0.644448 -0.007799  
1 -0.449204 2.448963 0.667226
```

In this case, you can pass `ignore_index=True`, which discards the indexes from each DataFrame and concatenates the data in the columns only, assigning a new default index:

In [114]: `pd.concat([df1, df2], ignore_index=True)`

Out[114]:

```
a      b      c      d  
0 1.248804 0.774191 -0.319657 -0.624964  
1 1.078814 0.544647 0.855588 1.343268  
2 -0.267175 1.793095 -0.652929 -1.886837  
3 -0.007799 1.059626    NaN 0.644448  
4 0.667226 -0.449204    NaN 2.448963
```

[Table 8-3](#) describes the `pandas.concat` function arguments.

Argument	Description
<code>objs</code>	List or dictionary of pandas objects to be concatenated; this is the only required argument
<code>axis</code>	Axis to concatenate along; defaults to concatenating along rows (<code>axis="index"</code>)
<code>join</code>	Either "inner" or "outer" ("outer" by default); whether to intersect (inner) or union (outer) indexes along the other axes

Argument	Description
keys	Values to associate with objects being concatenated, forming a hierarchical index along the concatenation axis; can be a list or array of arbitrary values, an array of tuples, or a list of arrays (if multiple-level arrays passed in levels)
levels	Specific indexes to use as hierarchical index level or levels if keys passed
names	Names for created hierarchical levels if keys and/or levels passed
verify_integrity	Check new axis in concatenated object for duplicates and raise an exception if so; by default (False) allows duplicates
ignore_index	Do not preserve indexes along concatenation axis, instead produce a new range(total_length) index

Table 8-3. pandas.concat function arguments

Combining Data with Overlap

There is another data combination situation that can't be expressed as either a merge or concatenation operation. You may have two datasets with indexes that overlap in full or in part. As a motivating example, consider NumPy's where function, which performs the array-oriented equivalent of an if-else expression:

```
In [115]: a = pd.Series([np.nan, 2.5, 0.0, 3.5, 4.5, np.nan],
```

```
.....:      index=["f", "e", "d", "c", "b", "a"])
```

```
In [116]: b = pd.Series([0., np.nan, 2., np.nan, np.nan, 5.],
```

```
.....:      index=["a", "b", "c", "d", "e", "f"])
```

```
In [117]: a
```

```
Out[117]:
```

```
f  NaN
```

```
e  2.5
```

```
d 0.0  
c 3.5  
b 4.5  
a NaN  
dtype: float64
```

In [118]: b

Out[118]:
a 0.0
b NaN
c 2.0
d NaN
e NaN
f 5.0
dtype: float64

In [119]: np.where(pd.isna(a), b, a)

Out[119]: array([0., 2.5, 0., 3.5, 4.5, 5.])

Here, whenever values in a are null, values from b are selected, otherwise the non-null values from a are selected. Using numpy.where does not check whether the index labels are aligned or not (and does not even require the objects to be the same length), so if you want to line up values by index, use the Series combine_first method:

In [120]: a.combine_first(b)

Out[120]:
a 0.0
b 4.5
c 3.5
d 0.0
e 2.5

```
f 5.0  
dtype: float64
```

With DataFrames, `combine_first` does the same thing column by column, so you can think of it as “patching” missing data in the calling object with data from the object you pass:

```
In [121]: df1 = pd.DataFrame({"a": [1., np.nan, 5., np.nan],  
.....: "b": [np.nan, 2., np.nan, 6.],  
.....: "c": range(2, 18, 4)})
```

```
In [122]: df2 = pd.DataFrame({"a": [5., 4., np.nan, 3., 7.],  
.....: "b": [np.nan, 3., 4., 6., 8.]})
```

```
In [123]: df1
```

```
Out[123]:
```

```
   a   b   c  
0  1.0  NaN  2  
1  NaN  2.0  6  
2  5.0  NaN  10  
3  NaN  6.0  14
```

```
In [124]: df2
```

```
Out[124]:
```

```
   a   b  
0  5.0  NaN  
1  4.0  3.0  
2  NaN  4.0  
3  3.0  6.0  
4  7.0  8.0
```

```
In [125]: df1.combine_first(df2)
```

```
Out[125]:
```

	a	b	c
0	1.0	NaN	2.0
1	4.0	2.0	6.0
2	5.0	4.0	10.0
3	3.0	6.0	14.0
4	7.0	8.0	NaN

The output of `combine_first` with DataFrame objects will have the union of all the column names.

8.3 Reshaping and Pivoting

There are a number of basic operations for rearranging tabular data. These are referred to as *reshape* or *pivot* operations.

Reshaping with Hierarchical Indexing

Hierarchical indexing provides a consistent way to rearrange data in a DataFrame. There are two primary actions:

stack

This “rotates” or pivots from the columns in the data to the rows.

unstack

This pivots from the rows into the columns.

I'll illustrate these operations through a series of examples. Consider a small DataFrame with string arrays as row and column indexes:

```
In [126]: data = pd.DataFrame(np.arange(6).reshape((2, 3)),
```

```
.....:           index=pd.Index(["Ohio", "Colorado"], name="state"),
.....:           columns=pd.Index(["one", "two", "three"],
.....:                         name="number"))
```

```
In [127]: data
```

```
Out[127]:
```

```
number  one two three
state
Ohio    0  1  2
Colorado 3  4  5
```

Using the `stack` method on this data pivots the columns into the rows, producing a Series:

```
In [128]: result = data.stack()
```

```
In [129]: result
```

```
Out[129]:
```

```
state  number
```

```
Ohio   one   0
```

```
      two   1
```

```
      three  2
```

```
Colorado one   3
```

```
      two   4
```

```
      three  5
```

```
dtype: int64
```

From a hierarchically indexed Series, you can rearrange the data back into a DataFrame with `unstack`:

```
In [130]: result.unstack()
```

```
Out[130]:
```

```
number  one two three
```

```
state
```

```
Ohio    0  1  2
```

```
Colorado 3  4  5
```

By default, the innermost level is unstacked (same with stack). You can unstack a different level by passing a level number or name:

```
In [131]: result.unstack(level=0)
```

```
Out[131]:
```

```
state  Ohio  Colorado
```

```
number
```

```
one    0    3
```

```
two    1    4
```

```
three   2    5
```

```
In [132]: result.unstack(level="state")
```

```
Out[132]:
```

```
state  Ohio  Colorado
```

```
number
```

```
one    0    3
```

```
two    1    4
```

```
three   2    5
```

Unstacking might introduce missing data if all of the values in the level aren't found in each subgroup:

```
In [133]: s1 = pd.Series([0, 1, 2, 3], index=["a", "b", "c", "d"], dtype="Int64")
```

```
In [134]: s2 = pd.Series([4, 5, 6], index=["c", "d", "e"], dtype="Int64")
```

```
In [135]: data2 = pd.concat([s1, s2], keys=["one", "two"])
```

```
In [136]: data2
```

```
Out[136]:
```

```
one a 0
```

```
    b 1
```

```
    c 2
```

```
    d 3
```

```
two c 4
```

```
    d 5
```

```
    e 6
```

```
dtype: Int64
```

Stacking filters out missing data by default, so the operation is more easily invertible:

```
In [137]: data2.unstack()
```

```
Out[137]:
```

```
    a b c d e
```

```
one 0 1 2 3 <NA>
```

```
two <NA> <NA> 4 5 6
```

```
In [138]: data2.unstack().stack()
```

```
Out[138]:
```

```
one a 0
```

```
    b 1
```

```
    c 2
```

```
    d 3
```

```
two c 4
```

```
    d 5
```

```
    e 6
```

```
dtype: Int64
```

```
In [139]: data2.unstack().stack(dropna=False)
```

```
Out[139]:
```

```
one a    0
```

```
    b    1
```

```
    c    2
```

```
    d    3
```

```
    e  <NA>
```

```
two a  <NA>
```

```
    b  <NA>
```

```
    c    4
```

```
    d    5
```

```
    e    6
```

```
dtype: Int64
```

When you unstack in a DataFrame, the level unstacked becomes the lowest level in the result:

```
In [140]: df = pd.DataFrame({"left": result, "right": result + 5},
```

```
.....:      columns=pd.Index(["left", "right"], name="side"))
```

```
In [141]: df
```

```
Out[141]:
```

```
side      left right
```

```
state  number
```

```
Ohio  one    0    5
```

```
    two    1    6
```

```
    three   2    7
```

```
Colorado one    3    8
```

```
    two    4    9
```

```
    three   5   10
```

```
In [142]: df.unstack(level="state")  
Out[142]:  
  
side  left      right  
state  Ohio Colorado  Ohio Colorado  
  
number  
  
one    0      3      5      8  
two    1      4      6      9  
three   2      5      7     10
```

As with unstack, when calling stack we can indicate the name of the axis to stack:

```
In [143]: df.unstack(level="state").stack(level="side")  
Out[143]:  
  
state      Colorado Ohio  
  
number side  
  
one  left      3      0  
      right     8      5  
two  left      4      1  
      right     9      6  
three left      5      2  
      right    10      7
```

Pivoting “Long” to “Wide” Format

A common way to store multiple time series in databases and CSV files is what is sometimes called *long* or *stacked* format. In this format, individual values are represented by a single row in a table rather than multiple values per row.

Let’s load some example data and do a small amount of time series wrangling and other data cleaning:

```
In [144]: data = pd.read_csv("examples/macrodata.csv")
```

```
In [145]: data = data.loc[:, ["year", "quarter", "realgdp", "infl", "unemp"]]
```

```
In [146]: data.head()
```

```
Out[146]:
```

	year	quarter	realgdp	infl	unemp
0	1959	1	2710.349	0.00	5.8
1	1959	2	2778.801	2.34	5.1
2	1959	3	2775.488	2.74	5.3
3	1959	4	2785.204	0.27	5.6
4	1960	1	2847.699	2.31	5.2

First, I use pandas.PeriodIndex (which represents time intervals rather than points in time), discussed in more detail in [Chapter 11](#), to combine the year and quarter columns to set the index to consist of datetime values at the end of each quarter:

```
In [147]: periods = pd.PeriodIndex(year=data.pop("year"),
.....:                 quarter=data.pop("quarter"),
.....:                 name="date")
```

```
In [148]: periods
```

```
Out[148]:
```

```
PeriodIndex(['1959Q1', '1959Q2', '1959Q3', '1959Q4', '1960Q1', '1960Q2',
 '1960Q3', '1960Q4', '1961Q1', '1961Q2',
 ...
 '2007Q2', '2007Q3', '2007Q4', '2008Q1', '2008Q2', '2008Q3',
 '2008Q4', '2009Q1', '2009Q2', '2009Q3'],
 dtype='period[Q-DEC]', name='date', length=203)
```

```
In [149]: data.index = periods.to_timestamp("D")
```

```
In [150]: data.head()
```

```
Out[150]:
```

```
realgdp infl unemp  
date  
1959-01-01 2710.349 0.00 5.8  
1959-04-01 2778.801 2.34 5.1  
1959-07-01 2775.488 2.74 5.3  
1959-10-01 2785.204 0.27 5.6  
1960-01-01 2847.699 2.31 5.2
```

Here I used the pop method on the DataFrame, which returns a column while deleting it from the DataFrame at the same time.

Then, I select a subset of columns and give the columns index the name "item":

```
In [151]: data = data.reindex(columns=["realgdp", "infl", "unemp"])
```

```
In [152]: data.columns.name = "item"
```

```
In [153]: data.head()
```

```
Out[153]:
```

```
item    realgdp infl unemp  
date  
1959-01-01 2710.349 0.00 5.8  
1959-04-01 2778.801 2.34 5.1  
1959-07-01 2775.488 2.74 5.3  
1959-10-01 2785.204 0.27 5.6  
1960-01-01 2847.699 2.31 5.2
```

Lastly, I reshape with stack, turn the new index levels into columns with reset_index, and finally give the column containing the data values the name "value":

```
In [154]: long_data = (data.stack()  
.....:     .reset_index()  
.....:     .rename(columns={0: "value"}))
```

Now, ldata looks like:

```
In [155]: long_data[:10]
```

```
Out[155]:
```

	date	item	value
0	1959-01-01	realgdp	2710.349
1	1959-01-01	infl	0.000
2	1959-01-01	unemp	5.800
3	1959-04-01	realgdp	2778.801
4	1959-04-01	infl	2.340
5	1959-04-01	unemp	5.100
6	1959-07-01	realgdp	2775.488
7	1959-07-01	infl	2.740
8	1959-07-01	unemp	5.300
9	1959-10-01	realgdp	2785.204

In this so-called *long* format for multiple time series, each row in the table represents a single observation.

Data is frequently stored this way in relational SQL databases, as a fixed schema (column names and data types) allows the number of distinct values in the item column to change as data is added to the table. In the previous example, date and item would usually be the primary keys (in relational database parlance), offering both relational integrity and easier joins. In some cases, the data may be more difficult to work with in this format; you might prefer to have a DataFrame containing one column per distinct item value indexed by

timestamps in the date column. DataFrame's pivot method performs exactly this transformation:

```
In [156]: pivoted = long_data.pivot(index="date", columns="item",
.....:                 values="value")
```

```
In [157]: pivoted.head()
```

Out[157]:

```
item    infl  realgdp  unemp
date
1959-01-01  0.00  2710.349  5.8
1959-04-01  2.34  2778.801  5.1
1959-07-01  2.74  2775.488  5.3
1959-10-01  0.27  2785.204  5.6
1960-01-01  2.31  2847.699  5.2
```

The first two values passed are the columns to be used, respectively, as the row and column index, then finally an optional value column to fill the DataFrame. Suppose you had two value columns that you wanted to reshape simultaneously:

```
In [158]: long_data["value2"] = np.random.standard_normal(len(long_data))
```

```
In [159]: long_data[:10]
```

Out[159]:

```
date  item  value  value2
0 1959-01-01  realgdp  2710.349  0.802926
1 1959-01-01  infl    0.000  0.575721
2 1959-01-01  unemp   5.800  1.381918
3 1959-04-01  realgdp  2778.801  0.000992
4 1959-04-01  infl    2.340 -0.143492
5 1959-04-01  unemp   5.100 -0.206282
6 1959-07-01  realgdp  2775.488 -0.222392
```

```
7 1959-07-01 infl 2.740 -1.682403  
8 1959-07-01 unemp 5.300 1.811659  
9 1959-10-01 realgdp 2785.204 -0.351305
```

By omitting the last argument, you obtain a DataFrame with hierarchical columns:

```
In [160]: pivoted = long_data.pivot(index="date", columns="item")
```

```
In [161]: pivoted.head()
```

```
Out[161]:
```

```
      value          value2  
item    infl  realgdp  unemp    infl  realgdp  unemp  
date  
1959-01-01  0.00 2710.349  5.8  0.575721  0.802926 1.381918  
1959-04-01  2.34 2778.801  5.1 -0.143492  0.000992 -0.206282  
1959-07-01  2.74 2775.488  5.3 -1.682403 -0.222392  1.811659  
1959-10-01  0.27 2785.204  5.6  0.128317 -0.351305 -1.313554  
1960-01-01  2.31 2847.699  5.2 -0.615939  0.498327  0.174072
```

```
In [162]: pivoted["value"].head()
```

```
Out[162]:
```

```
      item    infl  realgdp  unemp  
date  
1959-01-01  0.00 2710.349  5.8  
1959-04-01  2.34 2778.801  5.1  
1959-07-01  2.74 2775.488  5.3  
1959-10-01  0.27 2785.204  5.6  
1960-01-01  2.31 2847.699  5.2
```

Note that pivot is equivalent to creating a hierarchical index using set_index followed by a call to unstack:

```
In [163]: unstacked = long_data.set_index(["date", "item"]).unstack(level="item")
```

```
In [164]: unstacked.head()
```

```
Out[164]:
```

	value	value2						
item	infl	realgdp	unemp	infl	realgdp	unemp		
date								
1959-01-01	0.00	2710.349	5.8	0.575721	0.802926	1.381918		
1959-04-01	2.34	2778.801	5.1	-0.143492	0.000992	-0.206282		
1959-07-01	2.74	2775.488	5.3	-1.682403	-0.222392	1.811659		
1959-10-01	0.27	2785.204	5.6	0.128317	-0.351305	-1.313554		
1960-01-01	2.31	2847.699	5.2	-0.615939	0.498327	0.174072		

Pivoting “Wide” to “Long” Format

An inverse operation to pivot for DataFrames is pandas.melt. Rather than transforming one column into many in a new DataFrame, it merges multiple columns into one, producing a DataFrame that is longer than the input. Let’s look at an example:

```
In [166]: df = pd.DataFrame({"key": ["foo", "bar", "baz"],
```

```
.....:     "A": [1, 2, 3],  
.....:     "B": [4, 5, 6],  
.....:     "C": [7, 8, 9]})
```

```
In [167]: df
```

```
Out[167]:
```

	key	A	B	C
0	foo	1	4	7
1	bar	2	5	8

```
2 baz 3 6 9
```

The "key" column may be a group indicator, and the other columns are data values. When using pandas.melt, we must indicate which columns (if any) are group indicators. Let's use "key" as the only group indicator here:

```
In [168]: melted = pd.melt(df, id_vars="key")
```

```
In [169]: melted
```

```
Out[169]:
```

	key	variable	value
0	foo	A	1
1	bar	A	2
2	baz	A	3
3	foo	B	4
4	bar	B	5
5	baz	B	6
6	foo	C	7
7	bar	C	8
8	baz	C	9

Using pivot, we can reshape back to the original layout:

```
In [170]: reshaped = melted.pivot(index="key", columns="variable",
.....:                 values="value")
```

```
In [171]: reshaped
```

```
Out[171]:
```

```
variable A B C
```

```
key
```

```
bar    2 5 8  
baz    3 6 9  
foo    1 4 7
```

Since the result of pivot creates an index from the column used as the row labels, we may want to use reset_index to move the data back into a column:

In [172]: reshaped.reset_index()

Out[172]:

```
variable key A B C  
0      bar 2 5 8  
1      baz 3 6 9  
2      foo 1 4 7
```

You can also specify a subset of columns to use as value columns:

In [173]: pd.melt(df, id_vars="key", value_vars=["A", "B"])

Out[173]:

```
key variable value  
0 foo      A   1  
1 bar      A   2  
2 baz      A   3  
3 foo      B   4  
4 bar      B   5  
5 baz      B   6
```

pandas.melt can be used without any group identifiers, too:

```
In [174]: pd.melt(df, value_vars=["A", "B", "C"])
```

```
Out[174]:
```

```
variable value
```

```
0    A    1
1    A    2
2    A    3
3    B    4
4    B    5
5    B    6
6    C    7
7    C    8
8    C    9
```

```
In [175]: pd.melt(df, value_vars=["key", "A", "B"])
```

```
Out[175]:
```

```
variable value
```

```
0  key  foo
1  key  bar
2  key  baz
3    A    1
4    A    2
5    A    3
6    B    4
7    B    5
8    B    6
```

8.4 Conclusion

Now that you have some pandas basics for data import, cleaning, and reorganization under your belt, we are ready to move on to data visualization with matplotlib. We will

return to explore other areas of pandas later in the book when we discuss more advanced analytics.