

Python for Data Analysis, 3rd Edition

[Wes McKinney](#)

Published by [O'Reilly Media, Inc.](#)

<https://github.com/wesm/pydata-book>

Chapter 9. Plotting and Visualization

Making informative visualizations (sometimes called *plots*) is one of the most important tasks in data analysis. It may be a part of the exploratory process—for example, to help identify outliers or needed data transformations, or as a way of generating ideas for models. For others, building an interactive visualization for the web may be the end goal. Python has many add-on libraries for making static or dynamic visualizations, but I'll be mainly focused on [matplotlib](#) and libraries that build on top of it.

matplotlib is a desktop plotting package designed for creating plots and figures suitable for publication. The project was started by John Hunter in 2002 to enable a MATLAB-like plotting interface in Python. The matplotlib and IPython communities have collaborated to simplify interactive plotting from the IPython shell (and now, Jupyter notebook). matplotlib supports various GUI backends on all operating systems and can export visualizations to all of the common vector and raster graphics formats (PDF, SVG, JPG, PNG, BMP, GIF, etc.). With the exception of a few diagrams, nearly all of the graphics in this book were produced using matplotlib.

Over time, matplotlib has spawned a number of add-on toolkits for data visualization that use matplotlib for their underlying plotting. One of these is [seaborn](#), which we explore later in this chapter.

The simplest way to follow the code examples in the chapter is to output plots in the Jupyter notebook. To set this up, execute the following statement in a Jupyter notebook:

%matplotlib inline

Explain this code

Translate into

Translate into

Note

Since this book's first edition in 2012, many new data visualization libraries have been created, some of which (like Bokeh and Altair) take advantage of modern web technology to create interactive visualizations that integrate well with the Jupyter notebook. Rather than use multiple visualization tools in this book, I decided to stick with matplotlib for teaching the fundamentals, in particular since pandas has good integration with matplotlib. You can adapt the principles from this chapter to learn how to use other visualization libraries as well.

9.1 A Brief matplotlib API Primer

With matplotlib, we use the following import convention:

In [13]: `import matplotlib.pyplot as plt`

Explain this code

Translate into

Translate into

After running %matplotlib notebook in Jupyter (or simply %matplotlib in IPython), we can try creating a simple plot. If everything is set up right, a line plot like [Figure 9-1](#) should appear:

In [14]: `data = np.arange(10)`

In [15]: `data`

Out[15]: `array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])`

In [16]: `plt.plot(data)`

Explain this code

Translate into

Translate into

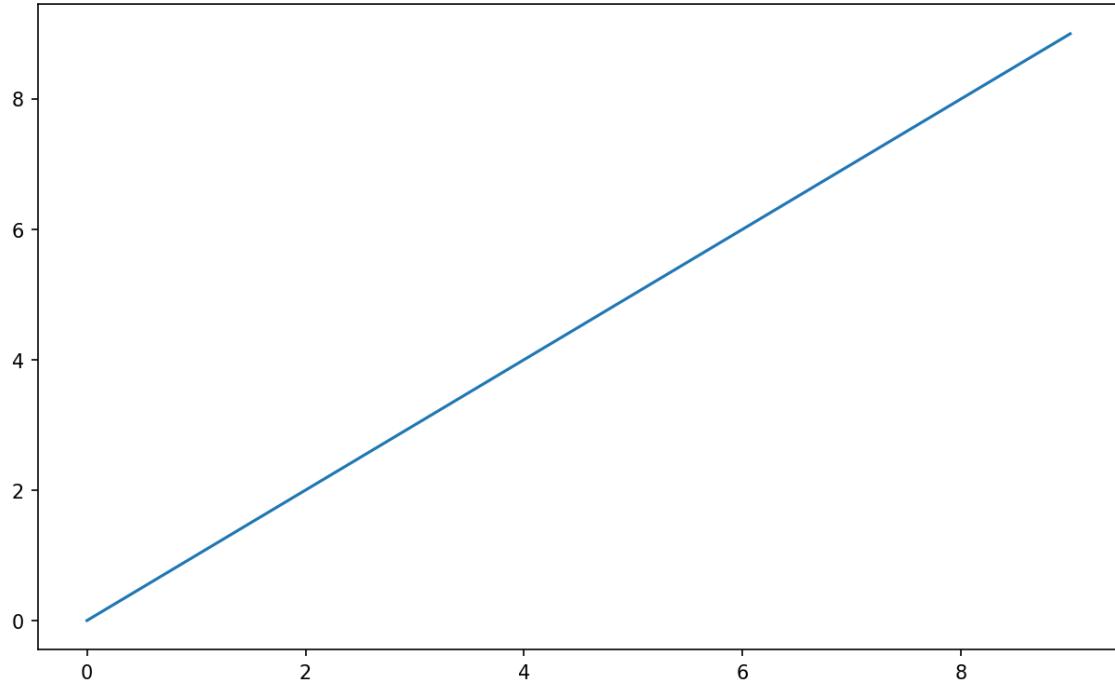


Figure 9-1. Simple line plot

While libraries like seaborn and pandas's built-in plotting functions will deal with many of the mundane details of making plots, should you wish to customize them beyond the function options provided, you will need to learn a bit about the matplotlib API.

Note

There is not enough room in the book to give comprehensive treatment of the breadth and depth of functionality in matplotlib. It should be enough to teach you the ropes to get up and running. The matplotlib gallery and documentation are the best resource for learning advanced features.

Figures and Subplots

Plots in matplotlib reside within a Figure object. You can create a new figure with plt.figure:

In [17]: fig = plt.figure()

Explain this code

Translate into

Translate into

In IPython, if you first run %matplotlib to set up the matplotlib integration, an empty plot window will appear, but in Jupyter nothing will be shown until we use a few more commands.

plt.figure has a number of options; notably, figsize will guarantee the figure has a certain size and aspect ratio if saved to disk.

You can't make a plot with a blank figure. You have to create one or more subplots using add_subplot:

In [18]: ax1 = fig.add_subplot(2, 2, 1)

Explain this code

Translate into

Translate into

This means that the figure should be 2×2 (so up to four plots in total), and we're selecting the first of four subplots (numbered from 1). If you create the next two subplots, you'll end up with a visualization that looks like [Figure 9-2](#):

In [19]: ax2 = fig.add_subplot(2, 2, 2)

In [20]: ax3 = fig.add_subplot(2, 2, 3)

Explain this code

Translate into

Translate into

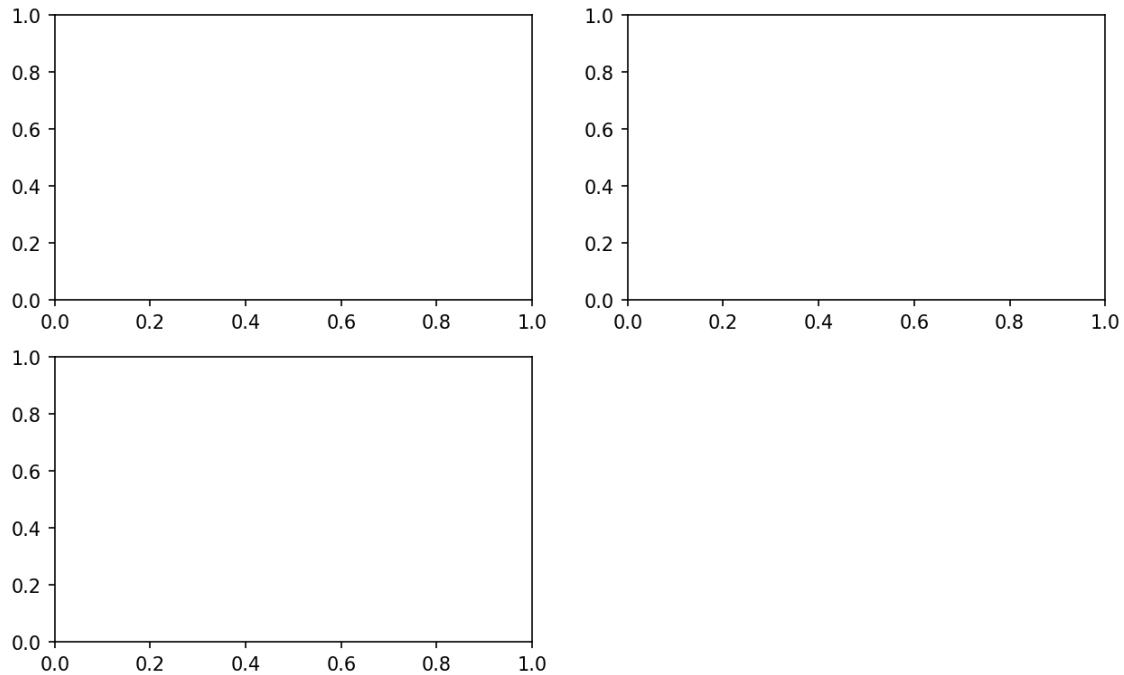


Figure 9-2. An empty matplotlib figure with three subplots

Tip

One nuance of using Jupyter notebooks is that plots are reset after each cell is evaluated, so you must put all of the plotting commands in a single notebook cell.

Here we run all of these commands in the same cell:

```
fig = plt.figure()
ax1 = fig.add_subplot(2, 2, 1)
ax2 = fig.add_subplot(2, 2, 2)
ax3 = fig.add_subplot(2, 2, 3)
```

Explain this code

Translate into

Translate into

These plot axis objects have various methods that create different types of plots, and it is preferred to use the axis methods over the top-level plotting functions like plt.plot. For example, we could make a line plot with the plot method (see [Figure 9-3](#)):

```
In [21]: ax3.plot(np.random.standard_normal(50).cumsum(), color="black",
....:      linestyle="dashed")
```

Explain this code

Translate into

Translate into

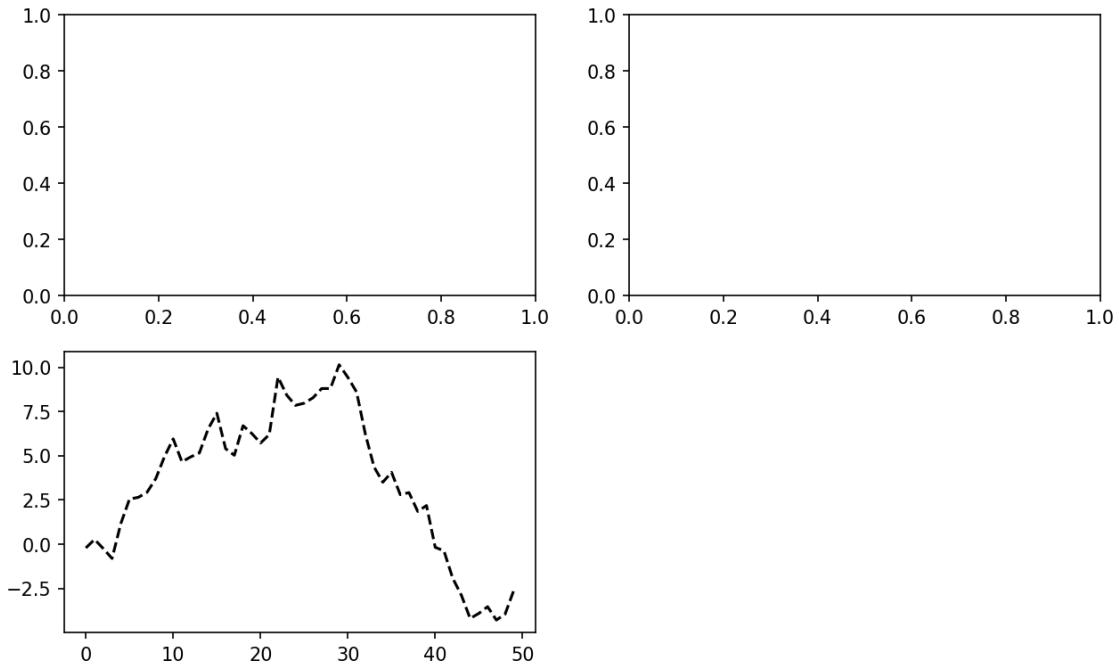


Figure 9-3. Data visualization after a single plot

You may notice output like <matplotlib.lines.Line2D at ...> when you run this. matplotlib returns objects that reference the plot subcomponent that was just added. A lot of the time you can safely ignore this output, or you can put a semicolon at the end of the line to suppress the output.

The additional options instruct matplotlib to plot a black dashed line. The objects returned by `fig.add_subplot` here are `AxesSubplot` objects, on which you can directly plot on the other empty subplots by calling each one's instance method (see [Figure 9-4](#)):

```
In [22]: ax1.hist(np.random.standard_normal(100), bins=20, color="black", alpha=0
.3);
```

```
In [23]: ax2.scatter(np.arange(30), np.arange(30) + 3 * np.random.standard_normal
(30));
```

Explain this code

Translate into

Translate into

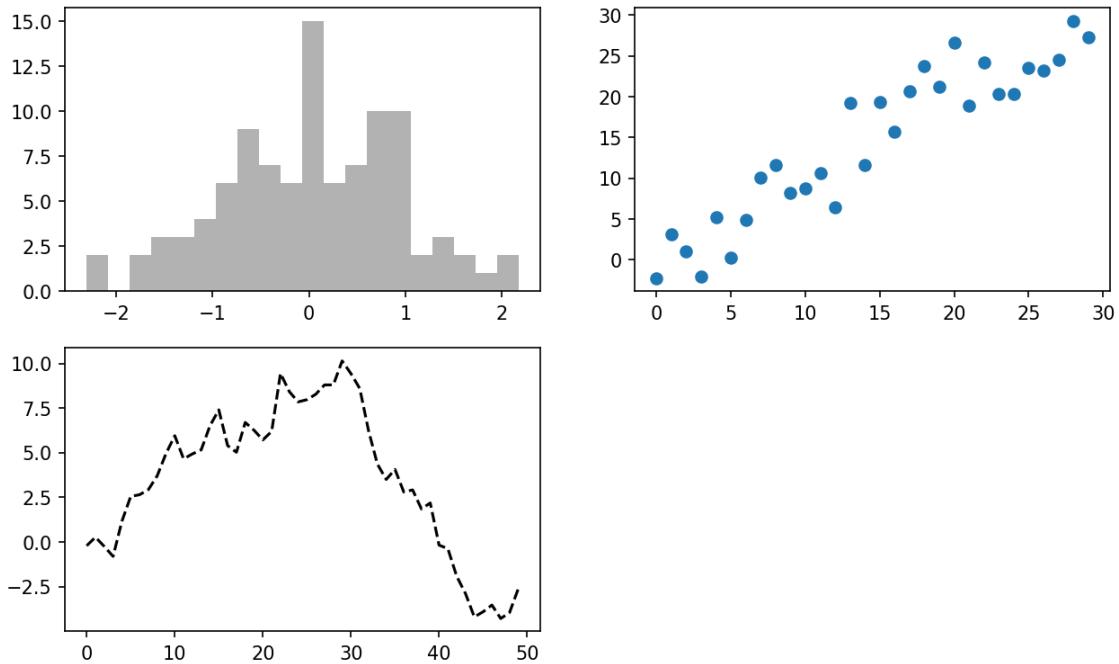


Figure 9-4. Data visualization after additional plots

The style option `alpha=0.3` sets the transparency of the overlaid plot.

You can find a comprehensive catalog of plot types in the [matplotlib documentation](#).

To make creating a grid of subplots more convenient, matplotlib includes a `plt.subplots` method that creates a new figure and returns a NumPy array containing the created subplot objects:

In [25]: `fig, axes = plt.subplots(2, 3)`

In [26]: `axes`

Out[26]:

```
array([[], [], []],
```

```
[<AxesSubplot>, <AxesSubplot>, <AxesSubplot>], dtype=object)
```

Explain this code

Translate into

Translate into

The `axes` array can then be indexed like a two-dimensional array; for example, `axes[0, 1]` refers to the subplot in the top row at the center. You can also indicate that subplots should have the same x- or y-axis using `sharex` and `sharey`, respectively. This can be useful when you're comparing data on the same scale; otherwise, matplotlib autoscales plot limits independently. See [Table 9-1](#) for more on this method.

Argument	Description
nrows	Number of rows of subplots
ncols	Number of columns of subplots
sharex	All subplots should use the same x-axis ticks (adjusting the xlim will affect all subplots)
sharey	All subplots should use the same y-axis ticks (adjusting the ylim will affect all subplots)
subplot_kw	Dictionary of keywords passed to add_subplot call used to create each subplot
**fig_kw	Additional keywords to subplots are used when creating the figure, such as plt.subplots(2, 2, figsize=(8, 6))

Table 9-1. matplotlib.pyplot.subplots options

Adjusting the spacing around subplots

By default, matplotlib leaves a certain amount of padding around the outside of the subplots and in spacing between subplots. This spacing is all specified relative to the height and width of the plot, so that if you resize the plot either programmatically or manually using the GUI window, the plot will dynamically adjust itself. You can change the spacing using the subplots_adjust method on Figure objects:

```
subplots_adjust(left=None, bottom=None, right=None, top=None,
                wspace=None, hspace=None)
```

Explain this code

Translate into

Translate into

wspace and hspace control the percent of the figure width and figure height, respectively, to use as spacing between subplots. Here is a small example you can execute in Jupyter where I shrink the spacing all the way to zero (see [Figure 9-5](#)):

```
fig, axes = plt.subplots(2, 2, sharex=True, sharey=True)
```

```

for i in range(2):
    for j in range(2):
        axes[i, j].hist(np.random.standard_normal(500), bins=50,
                        color="black", alpha=0.5)
fig.subplots_adjust(wspace=0, hspace=0)

```

Explain this code

Translate into

Translate into

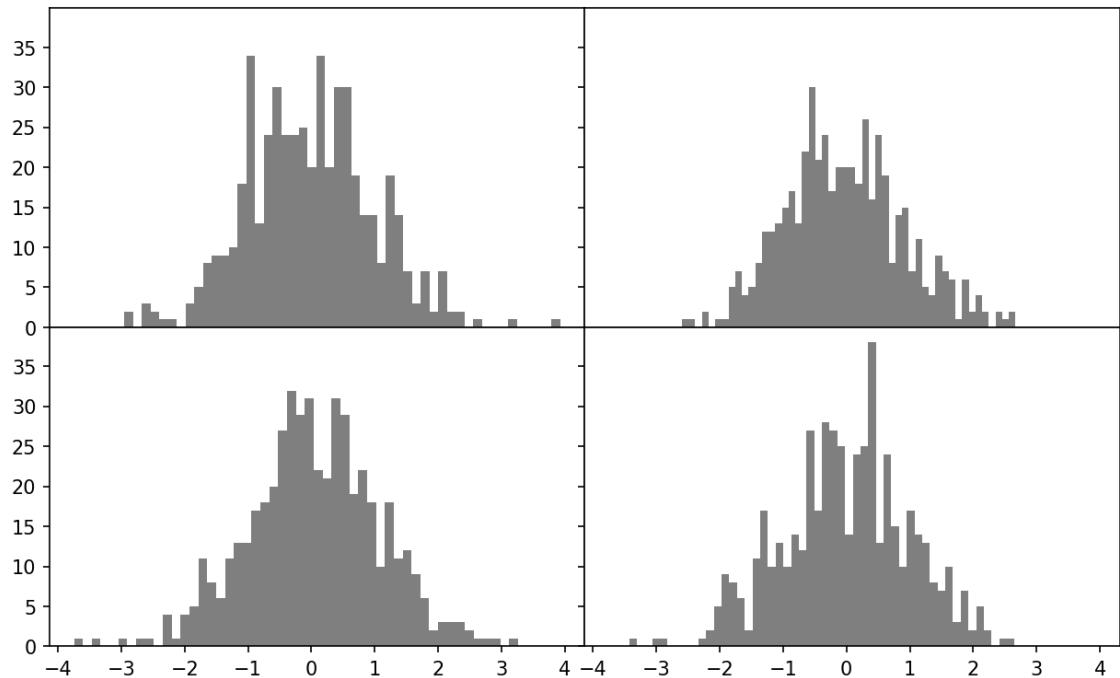


Figure 9-5. Data visualization with no inter-subplot spacing

You may notice that the axis labels overlap. matplotlib doesn't check whether the labels overlap, so in a case like this you would need to fix the labels yourself by specifying explicit tick locations and tick labels (we'll look at how to do this in the later section "[Ticks, Labels, and Legends](#)").

Colors, Markers, and Line Styles

matplotlib's line plot function accepts arrays of x and y coordinates and optional color styling options. For example, to plot x versus y with green dashes, you would execute:

```
ax.plot(x, y, linestyle="--", color="green")
```

Explain this code

Translate into

Translate into

A number of color names are provided for commonly used colors, but you can use any color on the spectrum by specifying its hex code (e.g., "#CECECE"). You can see some of the supported line styles by looking at the docstring for plt.plot (use plt.plot? in IPython or Jupyter). A more comprehensive reference is available in the online documentation.

Line plots can additionally have *markers* to highlight the actual data points. Since matplotlib's plot function creates a continuous line plot, interpolating between points, it can occasionally be unclear where the points lie. The marker can be supplied as an additional styling option (see [Figure 9-6](#)):

In [31]: `ax = fig.add_subplot()`

In [32]: `ax.plot(np.random.standard_normal(30).cumsum(), color="black",
....: linestyle="dashed", marker="o");`

Explain this code

Translate into

Translate into

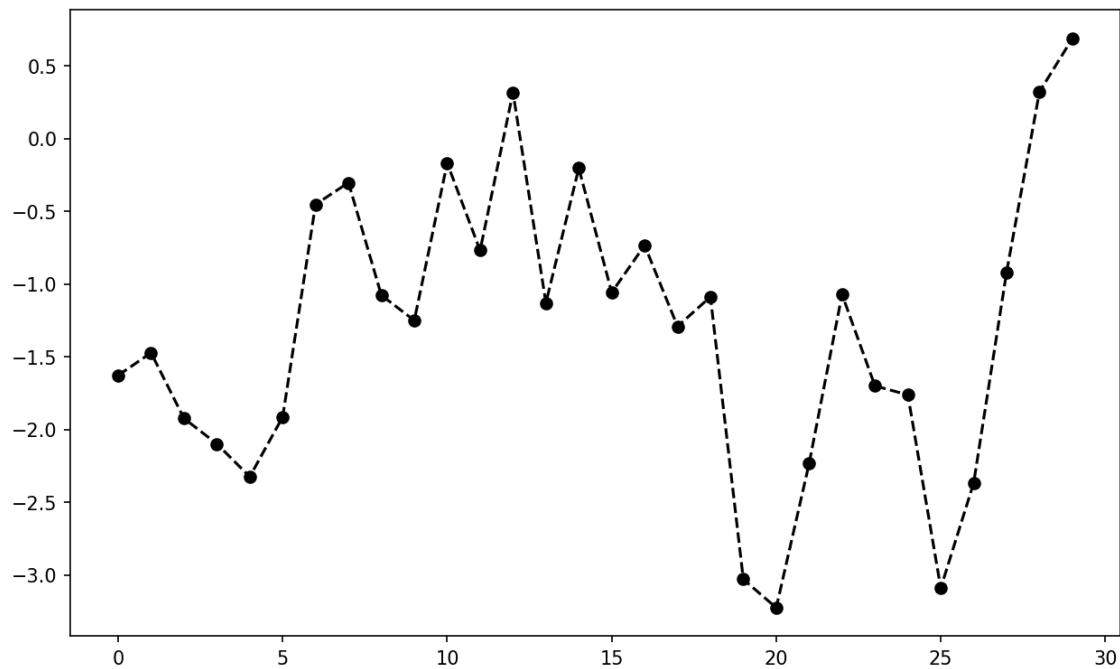


Figure 9-6. Line plot with markers

For line plots, you will notice that subsequent points are linearly interpolated by default. This can be altered with the drawstyle option (see [Figure 9-7](#)):

In [34]: `fig = plt.figure()`

In [35]: `ax = fig.add_subplot()`

```
In [36]: data = np.random.standard_normal(30).cumsum()
```

```
In [37]: ax.plot(data, color="black", linestyle="dashed", label="Default");
```

```
In [38]: ax.plot(data, color="black", linestyle="dashed",
```

```
....:     drawstyle="steps-post", label="steps-post");
```

```
In [39]: ax.legend()
```

Explain this code

Translate into

Translate into

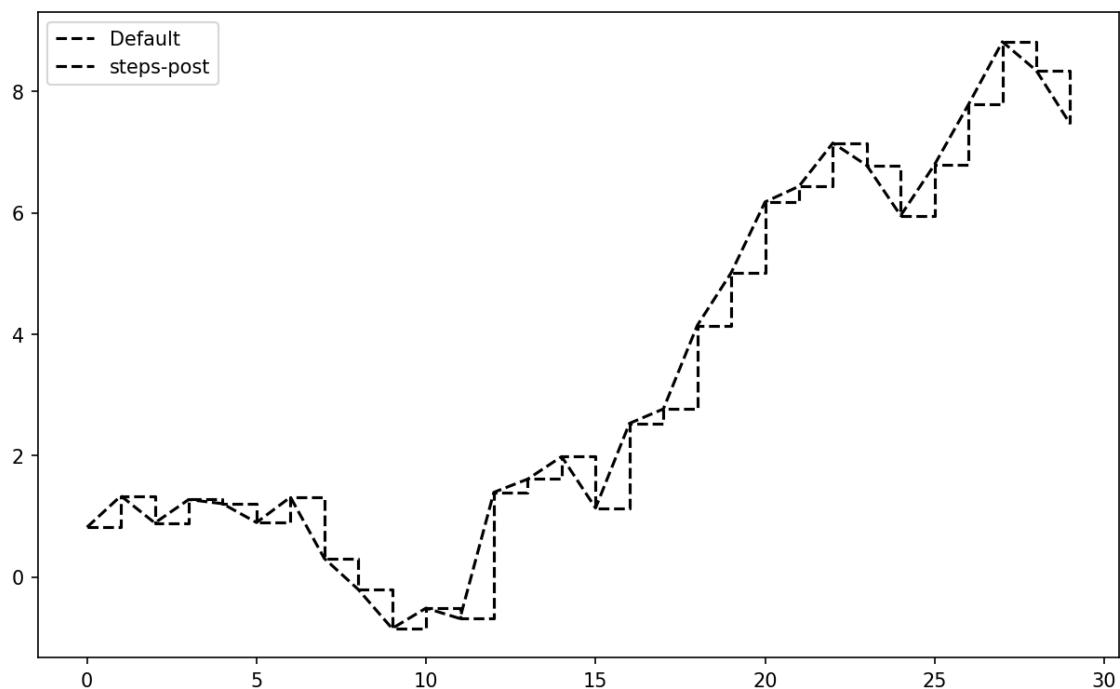


Figure 9-7. Line plot with different drawstyle options

Here, since we passed the label arguments to plot, we are able to create a plot legend to identify each line using ax.legend. I discuss legends more in [“Ticks, Labels, and Legends”](#).

Note

You must call ax.legend to create the legend, whether or not you passed the label options when plotting the data.

Ticks, Labels, and Legends

Most kinds of plot decorations can be accessed through methods on matplotlib axes objects. This includes methods like xlim, xticks, and xticklabels. These control the plot range, tick locations, and tick labels, respectively. They can be used in two ways:

- Called with no arguments returns the current parameter value (e.g., ax.xlim() returns the current x-axis plotting range)

- Called with parameters sets the parameter value (e.g., `ax.xlim([0, 10])`) sets the x-axis range to 0 to 10)

All such methods act on the active or most recently created `AxesSubplot`. Each corresponds to two methods on the subplot object itself; in the case of `xlim`, these are `ax.get_xlim` and `ax.set_xlim`.

Setting the title, axis labels, ticks, and tick labels

To illustrate customizing the axes, I'll create a simple figure and plot of a random walk (see [Figure 9-8](#)):

In [40]: `fig, ax = plt.subplots()`

In [41]: `ax.plot(np.random.standard_normal(1000).cumsum())`;

Explain this code

Translate into

Translate into

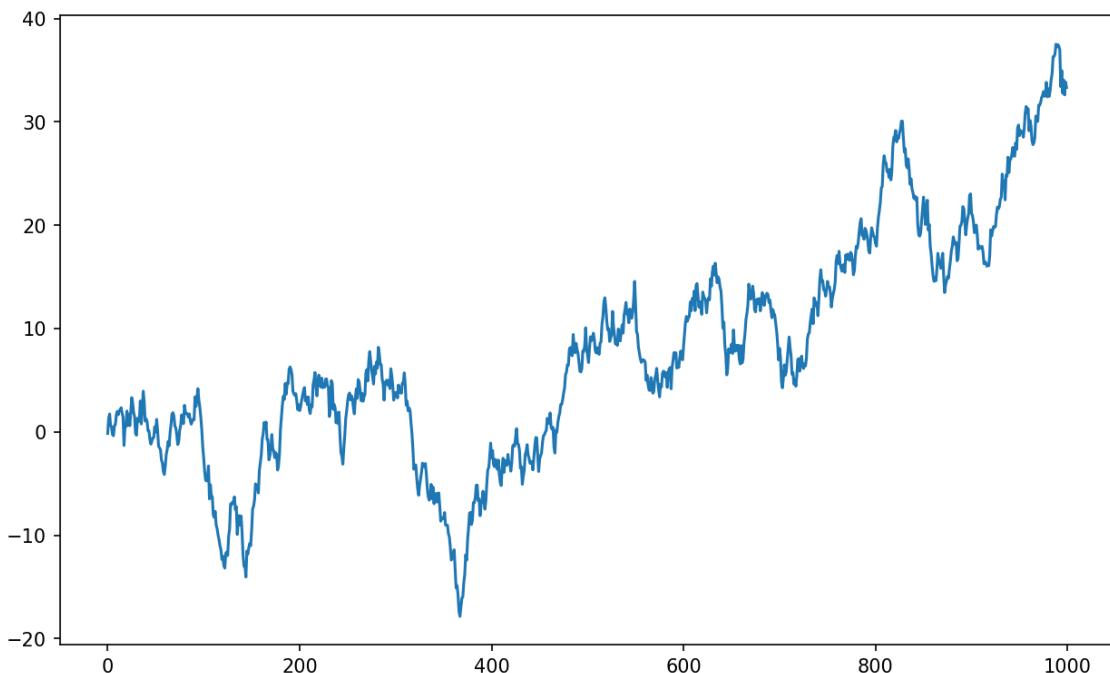


Figure 9-8. Simple plot for illustrating xticks (with default labels)

To change the x-axis ticks, it's easiest to use `set_xticks` and `set_xticklabels`. The former instructs matplotlib where to place the ticks along the data range; by default these locations will also be the labels. But we can set any other values as the labels using `set_xticklabels`:

In [42]: `ticks = ax.set_xticks([0, 250, 500, 750, 1000])`

In [43]: `labels = ax.set_xticklabels(["one", "two", "three", "four", "five"],`

```
....:         rotation=30, fontsize=8)
```

Explain this code

Translate into

Translate into

The rotation option sets the x tick labels at a 30-degree rotation. Lastly, set_xlabel gives a name to the x-axis, and set_title is the subplot title (see [Figure 9-9](#) for the resulting figure):

```
In [44]: ax.set_xlabel("Stages")
```

```
Out[44]: Text(0.5, 6.666666666666652, 'Stages')
```

```
In [45]: ax.set_title("My first matplotlib plot")
```

Explain this code

Translate into

Translate into

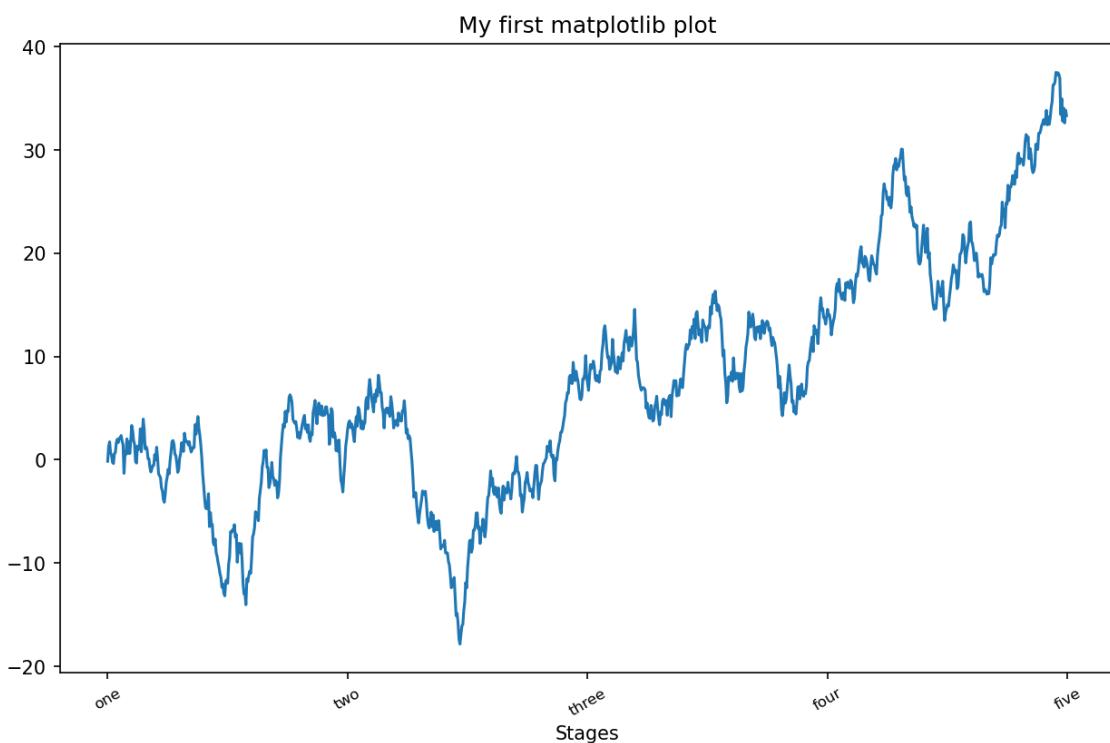


Figure 9-9. Simple plot for illustrating custom xticks

Modifying the y-axis consists of the same process, substituting y for x in this example. The axes class has a set method that allows batch setting of plot properties. From the prior example, we could also have written:

```
ax.set(title="My first matplotlib plot", xlabel="Stages")
```

Explain this code

Translate into

Translate into

Adding legends

Legends are another critical element for identifying plot elements. There are a couple of ways to add one. The easiest is to pass the `label` argument when adding each piece of the plot:

In [46]: `fig, ax = plt.subplots()`

In [47]: `ax.plot(np.random.randn(1000).cumsum(), color="black", label="one");`

In [48]: `ax.plot(np.random.randn(1000).cumsum(), color="black", linestyle="dashed`

`",`

`....: label="two");`

In [49]: `ax.plot(np.random.randn(1000).cumsum(), color="black", linestyle="dotted`

`",`

`....: label="three");`

Explain this code

Translate into

Translate into

Once you've done this, you can call `ax.legend()` to automatically create a legend. The resulting plot is in [Figure 9-10](#):

In [50]: `ax.legend()`

Explain this code

Translate into

Translate into

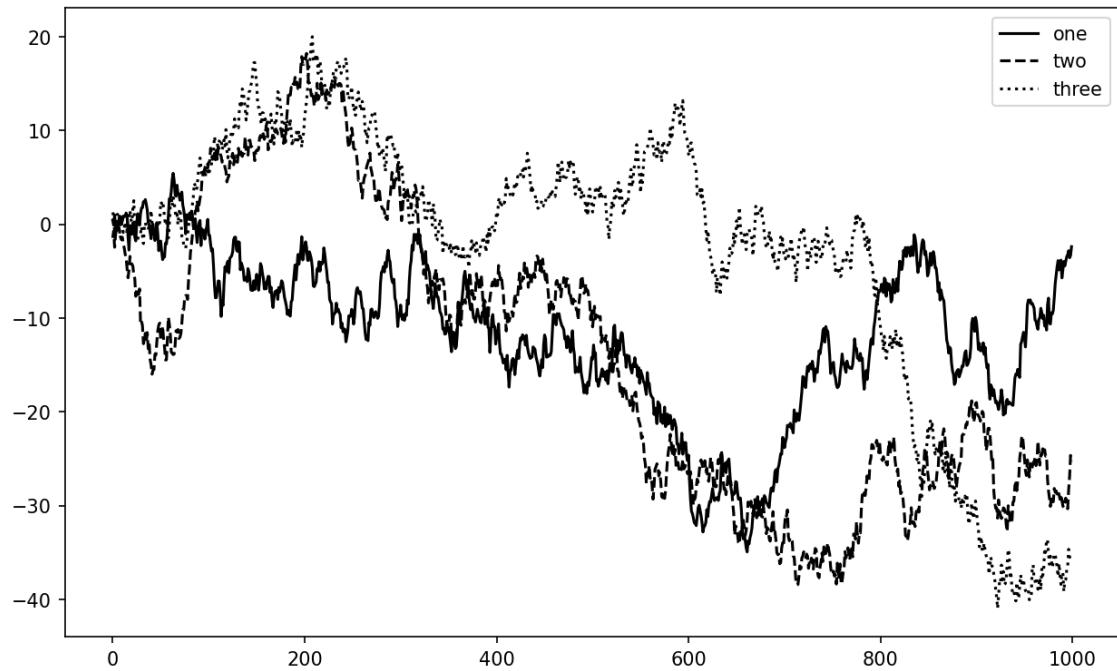


Figure 9-10. Simple plot with three lines and legend

The legend method has several other choices for the location loc argument. See the docstring (with `ax.legend?`) for more information.

The loc legend option tells matplotlib where to place the plot. The default is "best", which tries to choose a location that is most out of the way. To exclude one or more elements from the legend, pass no label or `label=_nolegend_`.

Annotations and Drawing on a Subplot

In addition to the standard plot types, you may wish to draw your own plot annotations, which could consist of text, arrows, or other shapes. You can add annotations and text using the `text`, `arrow`, and `annotate` functions. `text` draws text at given coordinates (`x`, `y`) on the plot with optional custom styling:

```
ax.text(x, y, "Hello world!",  
        family="monospace", fontsize=10)
```

Explain this code

Translate into

Translate into

Annotations can draw both text and arrows arranged appropriately. As an example, let's plot the closing S&P 500 index price since 2007 (obtained from Yahoo! Finance) and annotate it with some of the important dates from the 2008–2009 financial crisis. You can run this code example in a single cell in a Jupyter notebook. See [Figure 9-11](#) for the result:

```
from datetime import datetime
```

```
fig, ax = plt.subplots()

data = pd.read_csv("examples/spx.csv", index_col=0, parse_dates=True)
spx = data["SPX"]

spx.plot(ax=ax, color="black")

crisis_data = [
    (datetime(2007, 10, 11), "Peak of bull market"),
    (datetime(2008, 3, 12), "Bear Stearns Fails"),
    (datetime(2008, 9, 15), "Lehman Bankruptcy")
]

for date, label in crisis_data:
    ax.annotate(label, xy=(date, spx.asof(date) + 75),
               xytext=(date, spx.asof(date) + 225),
               arrowprops=dict(facecolor="black", headwidth=4, width=2,
                               headlength=4),
               horizontalalignment="left", verticalalignment="top")

# Zoom in on 2007-2010
ax.set_xlim(["1/1/2007", "1/1/2011"])
ax.set_ylim([600, 1800])
```

ax.set_title("Important dates in the 2008-2009 financial crisis")

Explain this code

Translate into

Translate into



Figure 9-11. Important dates in the 2008–2009 financial crisis

There are a couple of important points to highlight in this plot. The `ax.annotate` method can draw labels at the indicated x and y coordinates. We use the `set_xlim` and `set_ylim` methods to manually set the start and end boundaries for the plot rather than using matplotlib's default. Lastly, `ax.set_title` adds a main title to the plot.

See the online matplotlib gallery for many more annotation examples to learn from.

Drawing shapes requires some more care. matplotlib has objects that represent many common shapes, referred to as *patches*. Some of these, like `Rectangle` and `Circle`, are found in `matplotlib.pyplot`, but the full set is located in `matplotlib.patches`.

To add a shape to a plot, you create the `patch` object and add it to a subplot `ax` by passing the `patch` to `ax.add_patch` (see [Figure 9-12](#)):

```
fig, ax = plt.subplots()

rect = plt.Rectangle((0.2, 0.75), 0.4, 0.15, color="black", alpha=0.3)
circ = plt.Circle((0.7, 0.2), 0.15, color="blue", alpha=0.3)
pgon = plt.Polygon([[0.15, 0.15], [0.35, 0.4], [0.2, 0.6]],
                   color="green", alpha=0.5)

ax.add_patch(rect)
ax.add_patch(circ)
ax.add_patch(pgon)
```

Explain this code

Translate into

Translate into

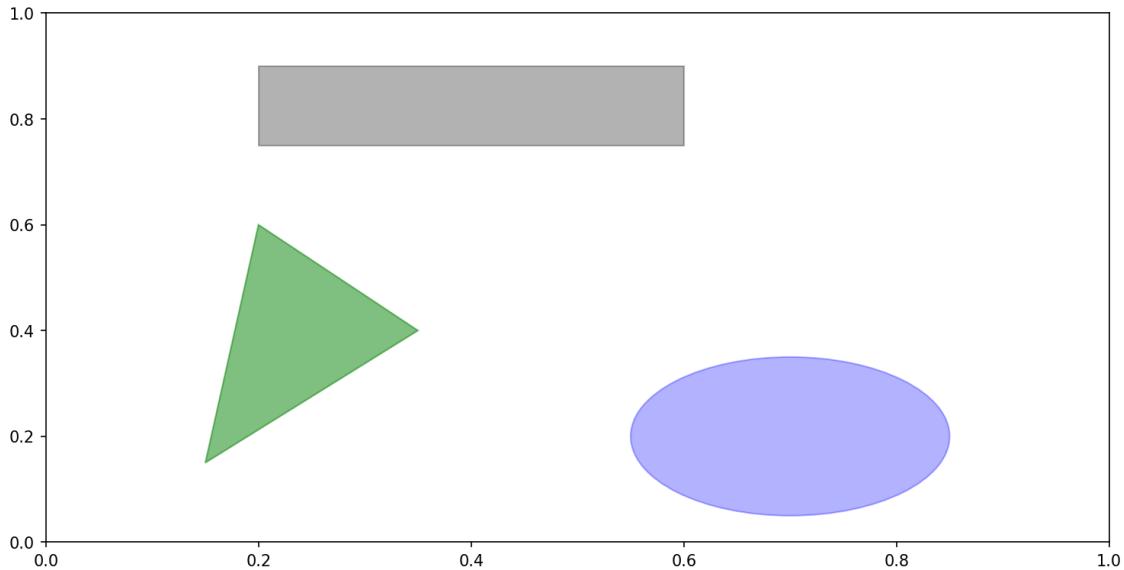


Figure 9-12. Data visualization composed from three different patches

If you look at the implementation of many familiar plot types, you will see that they are assembled from patches.

Saving Plots to File

You can save the active figure to file using the figure object's `savefig` instance method. For example, to save an SVG version of a figure, you need only type:

```
fig.savefig("figpath.svg")
```

Explain this code

Translate into

Translate into

The file type is inferred from the file extension. So if you used `.pdf` instead, you would get a PDF. One important option that I use frequently for publishing graphics is `dpi`, which controls the dots-per-inch resolution. To get the same plot as a PNG at 400 DPI, you would do:

```
fig.savefig("figpath.png", dpi=400)
```

Explain this code

Translate into

Translate into

See [Table 9-2](#) for a list of some other options for `savefig`. For a comprehensive listing, refer to the docstring in IPython or Jupyter.

Argument	Description
fname	String containing a filepath or a Python file-like object. The figure format is inferred from the file extension (e.g., .pdf for PDF or .png for PNG).
dpi	The figure resolution in dots per inch; defaults to 100 in IPython or 72 in Jupyter out of the box but can be configured.
facecolor, edgecolor	The color of the figure background outside of the subplots; "w" (white), by default.
format	The explicit file format to use ("png", "pdf", "svg", "ps", "eps", ...).

Table 9-2. Some fig.savefig options

matplotlib Configuration

matplotlib comes configured with color schemes and defaults that are geared primarily toward preparing figures for publication. Fortunately, nearly all of the default behavior can be customized via global parameters governing figure size, subplot spacing, colors, font sizes, grid styles, and so on. One way to modify the configuration programmatically from Python is to use the `rc` method; for example, to set the global default figure size to be 10×10 , you could enter:

```
plt.rc("figure", figsize=(10, 10))
```

Explain this code

Translate into

Translate into

All of the current configuration settings are found in the `plt.rcParams` dictionary, and they can be restored to their default values by calling the `plt.rcdefaults()` function.

The first argument to `rc` is the component you wish to customize, such as "figure", "axes", "xtick", "ytick", "grid", "legend", or many others. After that can follow a sequence of keyword arguments indicating the new parameters. A convenient way to write down the options in your program is as a dictionary:

```
plt.rc("font", family="monospace", weight="bold", size=8)
```

Explain this code

Translate into

Translate into

For more extensive customization and to see a list of all the options, matplotlib comes with a configuration file *matplotlibrc* in the *matplotlib/mpl-data* directory. If you customize this file and place it in your home directory titled *.matplotlibrc*, it will be loaded each time you use matplotlib.

As we'll see in the next section, the seaborn package has several built-in plot themes or *styles* that use matplotlib's configuration system internally.

9.2 Plotting with pandas and seaborn

matplotlib can be a fairly low-level tool. You assemble a plot from its base components: the data display (i.e., the type of plot: line, bar, box, scatter, contour, etc.), legend, title, tick labels, and other annotations.

In pandas, we may have multiple columns of data, along with row and column labels. pandas itself has built-in methods that simplify creating visualizations from DataFrame and Series objects. Another library is [seaborn](#), a high-level statistical graphics library built on matplotlib. seaborn simplifies creating many common visualization types.

Line Plots

Series and DataFrame have a `plot` attribute for making some basic plot types. By default, `plot()` makes line plots (see [Figure 9-13](#)):

```
In [61]: s = pd.Series(np.random.standard_normal(10).cumsum(), index=np.arange(0, 100, 10))
```

```
In [62]: s.plot()
```

Explain this code

Translate into

Translate into

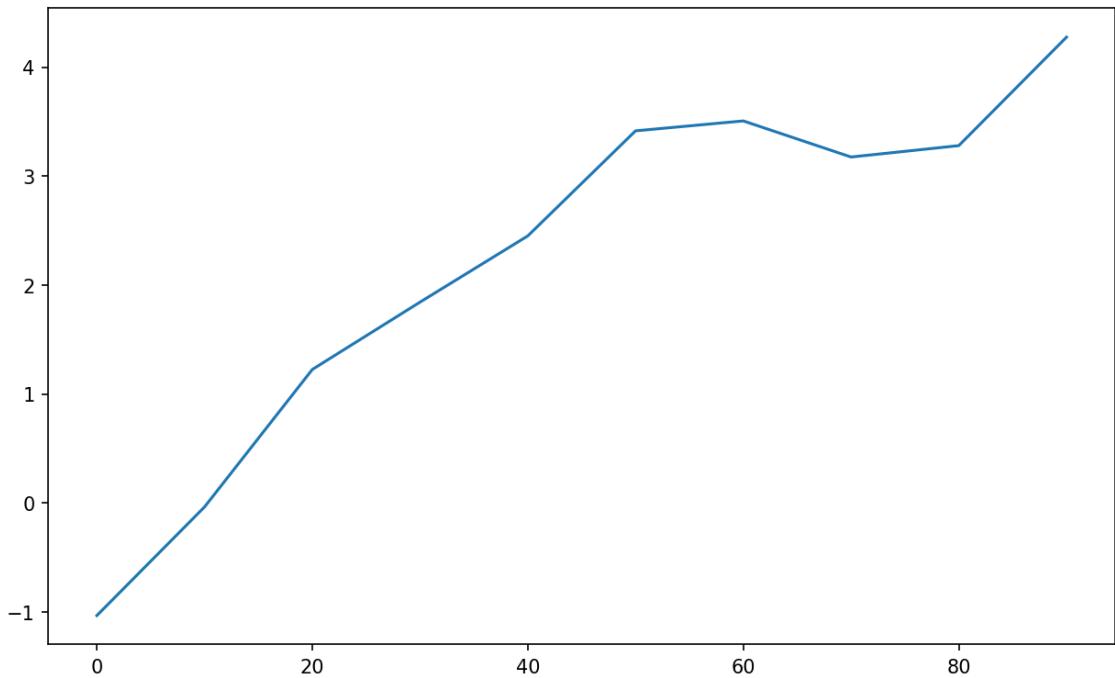


Figure 9-13. Simple Series plot

The Series object's index is passed to matplotlib for plotting on the x-axis, though you can disable this by passing `use_index=False`. The x-axis ticks and limits can be adjusted with the `xticks` and `xlim` options, and the y-axis respectively with `yticks` and `ylim`. See [Table 9-3](#) for a partial listing of plot options. I'll comment on a few more of them throughout this section and leave the rest for you to explore.

Argument	Description
label	Label for plot legend
ax	matplotlib subplot object to plot on; if nothing passed, uses active matplotlib subplot
style	Style string, like "ko--", to be passed to matplotlib
alpha	The plot fill opacity (from 0 to 1)
kind	Can be "area", "bar", "barh", "density", "hist", "kde", "line", or "pie"; defaults to "line"

Argument	Description
figsize	Size of the figure object to create
logx	Pass True for logarithmic scaling on the x axis; pass "sym" for symmetric logarithm that permits negative values
logy	Pass True for logarithmic scaling on the y axis; pass "sym" for symmetric logarithm that permits negative values
title	Title to use for the plot
use_index	Use the object index for tick labels
rot	Rotation of tick labels (0 through 360)
xticks	Values to use for x-axis ticks
yticks	Values to use for y-axis ticks
xlim	x-axis limits (e.g., [0, 10])
ylim	y-axis limits
grid	Display axis grid (off by default)

Table 9-3. Series.plot method arguments

Most of pandas's plotting methods accept an optional `ax` parameter, which can be a matplotlib subplot object. This gives you more flexible placement of subplots in a grid layout.

`DataFrame`'s `plot` method plots each of its columns as a different line on the same subplot, creating a legend automatically (see [Figure 9-14](#)):

```
In [63]: df = pd.DataFrame(np.random.standard_normal((10, 4)).cumsum(0),
```

```
....:     columns=["A", "B", "C", "D"],  
....:     index=np.arange(0, 100, 10))
```

In [64]: plt.style.use('grayscale')

In [65]: df.plot()

Explain this code

Translate into

Translate into

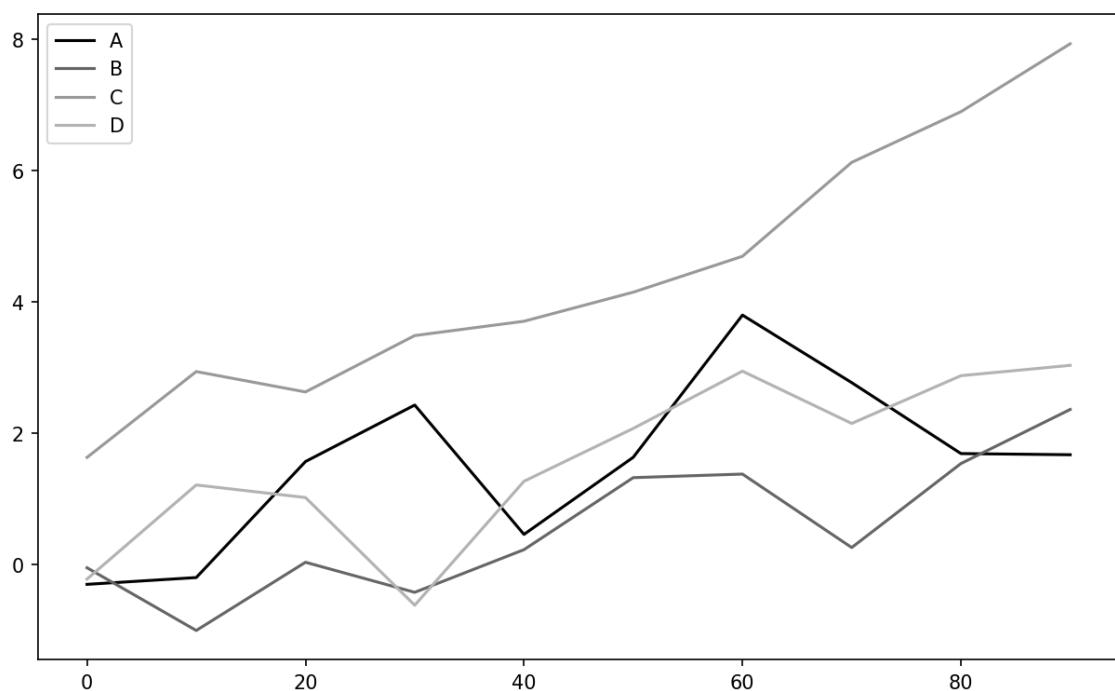


Figure 9-14. Simple DataFrame plot

Note

Here I used plt.style.use('grayscale') to switch to a color scheme more suitable for black and white publication, since some readers will not be able to see the full color plots.

The plot attribute contains a “family” of methods for different plot types. For example, df.plot() is equivalent to df.plot.line(). We’ll explore some of these methods next.

Note

Additional keyword arguments to plot are passed through to the respective matplotlib plotting function, so you can further customize these plots by learning more about the matplotlib API.

DataFrame has a number of options allowing some flexibility for how the columns are handled, for example, whether to plot them all on the same subplot or to create separate subplots. See [Table 9-4](#) for more on these.

Argument	Description
subplots	Plot each DataFrame column in a separate subplot
layouts	2-tuple (rows, columns) providing layout of subplots
sharex	If subplots=True, share the same x-axis, linking ticks and limits
sharey	If subplots=True, share the same y-axis
legend	Add a subplot legend (True by default)
sort_columns	Plot columns in alphabetical order; by default uses existing column order

Table 9-4. DataFrame-specific plot arguments

Note

For time series plotting, see [Chapter 11](#).

Bar Plots

The `plot.bar()` and `plot.banh()` make vertical and horizontal bar plots, respectively. In this case, the Series or DataFrame index will be used as the x (bar) or y (bahn) ticks (see [Figure 9-15](#)):

In [66]: `fig, axes = plt.subplots(2, 1)`

In [67]: `data = pd.Series(np.random.uniform(size=16), index=list("abcdefghijklmno
p"))`

In [68]: `data.plot.bar(ax=axes[0], color="black", alpha=0.7)`

Out[68]: <AxesSubplot:>

```
In [69]: data.plot.barh(ax=axes[1], color="black", alpha=0.7)
```

Explain this code

Translate into

Translate into

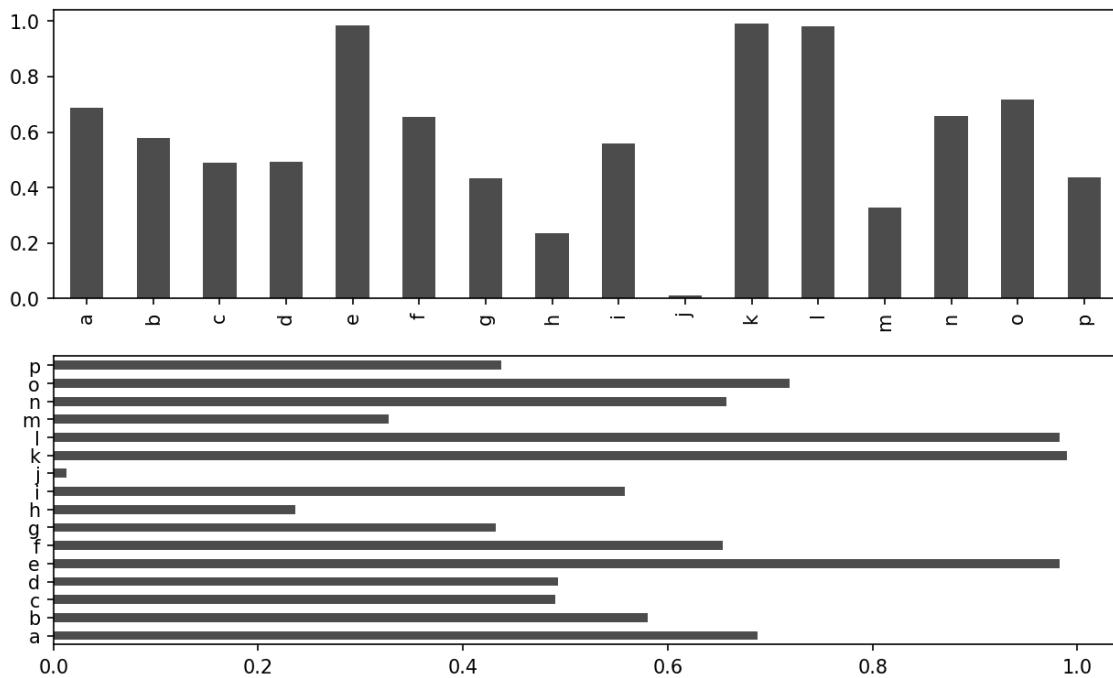


Figure 9-15. Horizontal and vertical bar plot

With a DataFrame, bar plots group the values in each row in bars, side by side, for each value. See [Figure 9-16](#):

```
In [71]: df = pd.DataFrame(np.random.uniform(size=(6, 4)),
```

```
....:     index=["one", "two", "three", "four", "five", "six"],  
....:     columns=pd.Index(["A", "B", "C", "D"], name="Genus"))
```

```
In [72]: df
```

```
Out[72]:
```

Genus	A	B	C	D
one	0.370670	0.602792	0.229159	0.486744
two	0.420082	0.571653	0.049024	0.880592
three	0.814568	0.277160	0.880316	0.431326
four	0.374020	0.899420	0.460304	0.100843

```
five 0.433270 0.125107 0.494675 0.961825  
six 0.601648 0.478576 0.205690 0.560547
```

In [73]: df.plot.bar()

Explain this code

Translate into

Translate into

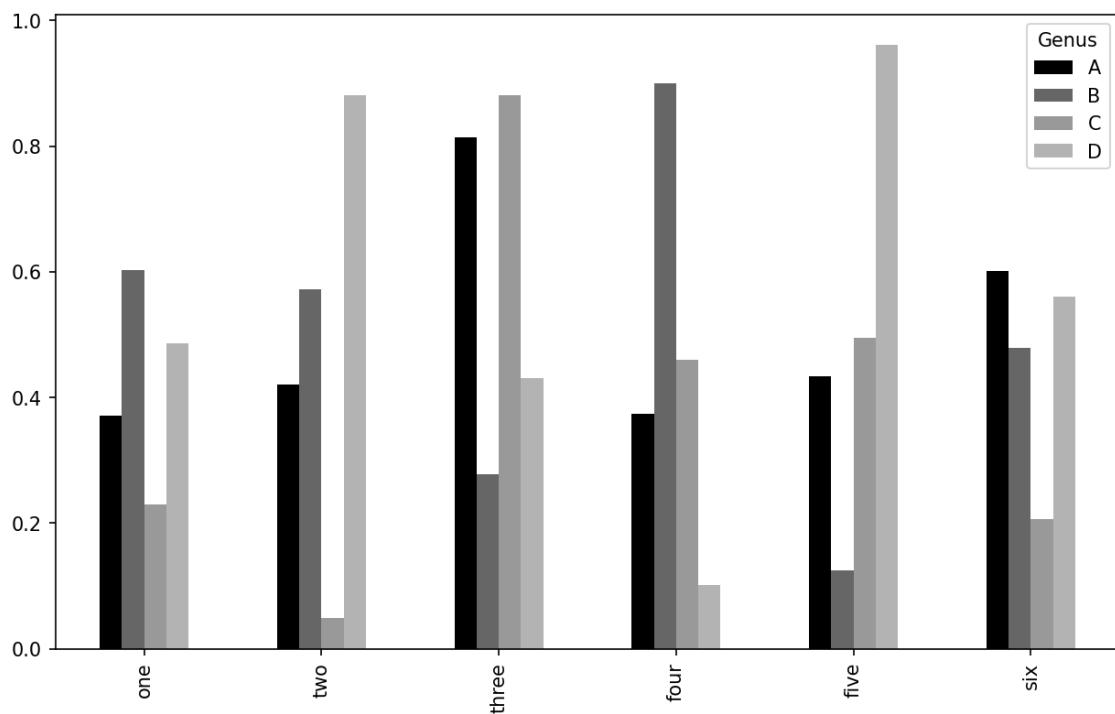


Figure 9-16. DataFrame bar plot

Note that the name “Genus” on the DataFrame’s columns is used to title the legend.

We create stacked bar plots from a DataFrame by passing stacked=True, resulting in the value in each row being stacked together horizontally (see [Figure 9-17](#)):

In [75]: df.plot.barh(stacked=True, alpha=0.5)

Explain this code

Translate into

Translate into

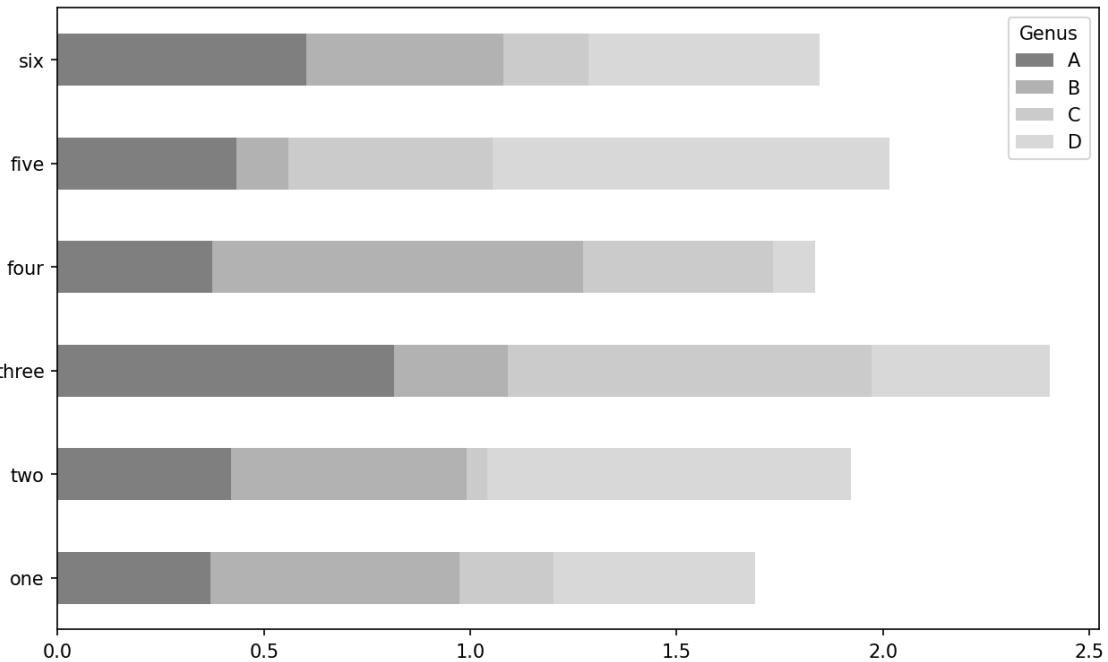


Figure 9-17. DataFrame stacked bar plot

Note

A useful recipe for bar plots is to visualize a Series's value frequency using `value_counts`: `s.value_counts().plot.bar()`.

Let's have a look at an example dataset about restaurant tipping. Suppose we wanted to make a stacked bar plot showing the percentage of data points for each party size for each day. I load the data using `read_csv` and make a cross-tabulation by day and party size. The `pandas.crosstab` function is a convenient way to compute a simple frequency table from two DataFrame columns:

```
In [77]: tips = pd.read_csv("examples/tips.csv")
```

```
In [78]: tips.head()
```

Out[78]:

	total_bill	tip	smoker	day	time	size
0	16.99	1.01	No	Sun	Dinner	2
1	10.34	1.66	No	Sun	Dinner	3
2	21.01	3.50	No	Sun	Dinner	3
3	23.68	3.31	No	Sun	Dinner	2
4	24.59	3.61	No	Sun	Dinner	4

```
In [79]: party_counts = pd.crosstab(tips["day"], tips["size"])
```

```
In [80]: party_counts = party_counts.reindex(index=["Thur", "Fri", "Sat", "Sun"])
```

In [81]: party_counts

Out[81]:

size 1 2 3 4 5 6

day

Thur 1 48 4 5 1 3

Fri 1 16 1 1 0 0

Sat 2 53 18 13 1 0

Sun 0 39 15 18 3 1

Explain this code

Translate into

Translate into

Since there are not many one- and six-person parties, I remove them here:

```
In [82]: party_counts = party_counts.loc[:, 2:5]
```

Explain this code

Translate into

Translate into

Then, normalize so that each row sums to 1, and make the plot (see [Figure 9-18](#)):

Normalize to sum to 1

```
In [83]: party_pcts = party_counts.div(party_counts.sum(axis="columns"),
.....:           axis="index")
```

```
In [84]: party_pcts
```

Out[84]:

size 2 3 4 5

day

Thur 0.827586 0.068966 0.086207 0.017241

Fri 0.888889 0.055556 0.055556 0.000000

Sat 0.623529 0.211765 0.152941 0.011765

```
Sun 0.520000 0.200000 0.240000 0.040000
```

```
In [85]: party_pcts.plot.bar(stacked=True)
```

Explain this code

Translate into

Translate into

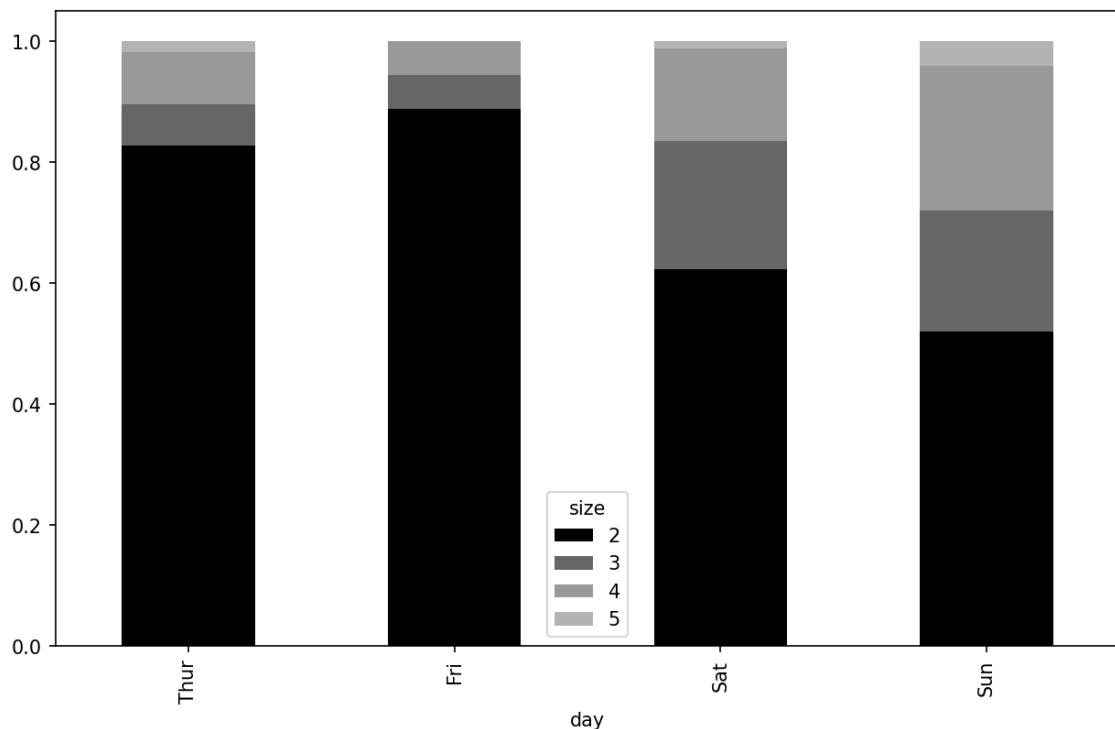


Figure 9-18. Fraction of parties by size within each day

So you can see that party sizes appear to increase on the weekend in this dataset.

With data that requires aggregation or summarization before making a plot, using the seaborn package can make things much simpler (install it with conda install seaborn). Let's look now at the tipping percentage by day with seaborn (see [Figure 9-19](#) for the resulting plot):

```
In [87]: import seaborn as sns
```

```
In [88]: tips["tip_pct"] = tips["tip"] / (tips["total_bill"] - tips["tip"])
```

```
In [89]: tips.head()
```

Out[89]:

```
total_bill tip smoker day time size tip_pct
```

```

0    16.99 1.01  No Sun Dinner  2  0.063204
1    10.34 1.66  No Sun Dinner  3  0.191244
2    21.01 3.50  No Sun Dinner  3  0.199886
3    23.68 3.31  No Sun Dinner  2  0.162494
4    24.59 3.61  No Sun Dinner  4  0.172069

```

In [90]: `sns.barplot(x="tip_pct", y="day", data=tips, orient="h")`

Explain this code

Translate into

Translate into

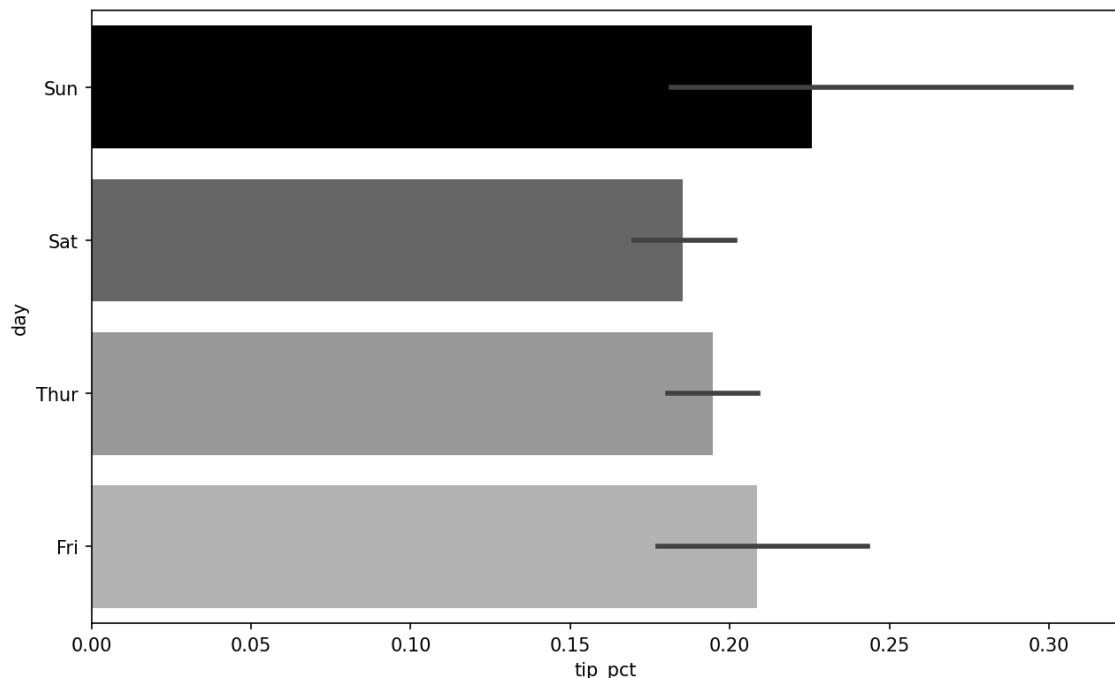


Figure 9-19. Tipping percentage by day with error bars

Plotting functions in seaborn take a `data` argument, which can be a pandas DataFrame. The other arguments refer to column names. Because there are multiple observations for each value in the day, the bars are the average value of `tip_pct`. The black lines drawn on the bars represent the 95% confidence interval (this can be configured through optional arguments).

`seaborn.barplot` has a `hue` option that enables us to split by an additional categorical value (see [Figure 9-20](#)):

In [92]: `sns.barplot(x="tip_pct", y="day", hue="time", data=tips, orient="h")`

Explain this code

Translate into

Translate into

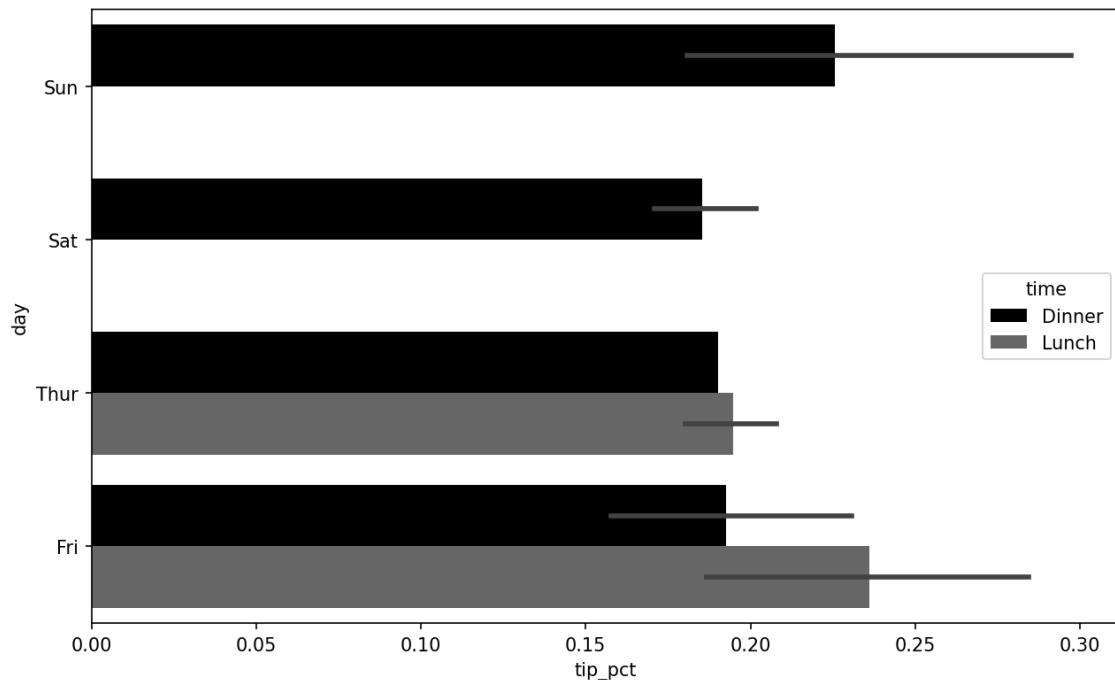


Figure 9-20. Tipping percentage by day and time

Notice that seaborn has automatically changed the aesthetics of plots: the default color palette, plot background, and grid line colors. You can switch between different plot appearances using `seaborn.set_style`:

In [94]: `sns.set_style("whitegrid")`

Explain this code

Translate into

Translate into

When producing plots for black-and-white print medium, you may find it useful to set a greyscale color palette, like so:

`sns.set_palette("Greys_r")`

Explain this code

Translate into

Translate into

Histograms and Density Plots

A *histogram* is a kind of bar plot that gives a discretized display of value frequency. The data points are split into discrete, evenly spaced bins, and the number of data points in each bin is plotted. Using the tipping data from before, we can make a histogram of tip percentages of the total bill using the `plot.hist` method on the Series (see [Figure 9-21](#)):

In [96]: `tips["tip_pct"].plot.hist(bins=50)`

Explain this code

Translate into

Translate into

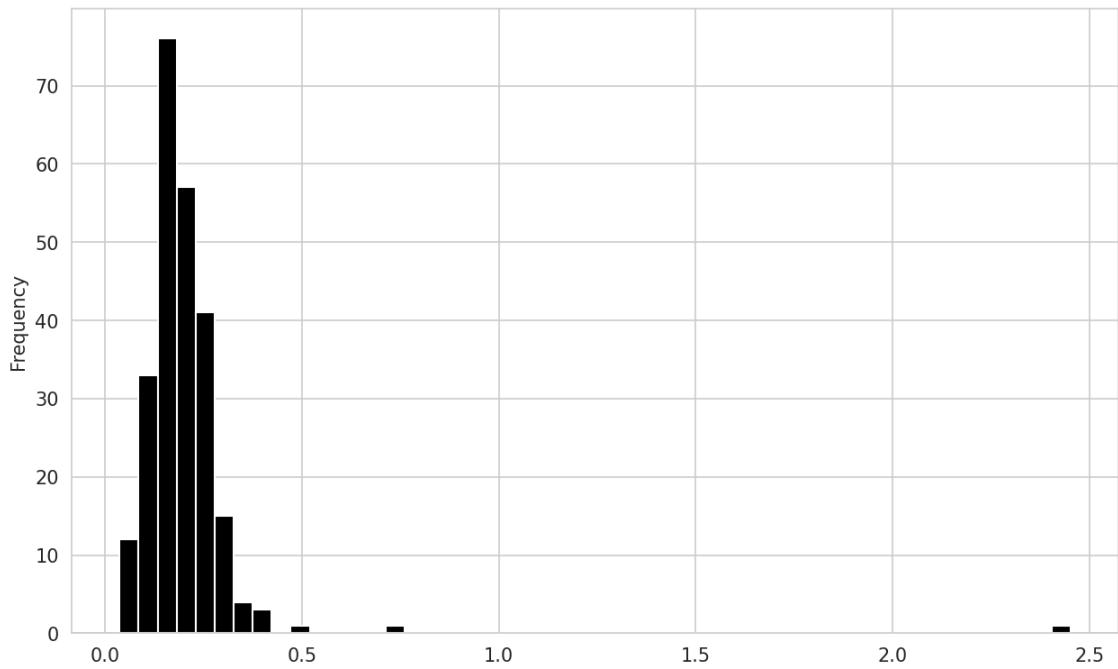


Figure 9-21. Histogram of tip percentages

A related plot type is a *density plot*, which is formed by computing an estimate of a continuous probability distribution that might have generated the observed data. The usual procedure is to approximate this distribution as a mixture of “kernels”—that is, simpler distributions like the normal distribution. Thus, density plots are also known as kernel density estimate (KDE) plots. Using `plot.density` makes a density plot using the conventional mixture-of-normals estimate (see [Figure 9-22](#)):

In [98]: `tips["tip_pct"].plot.density()`

Explain this code

Translate into

Translate into

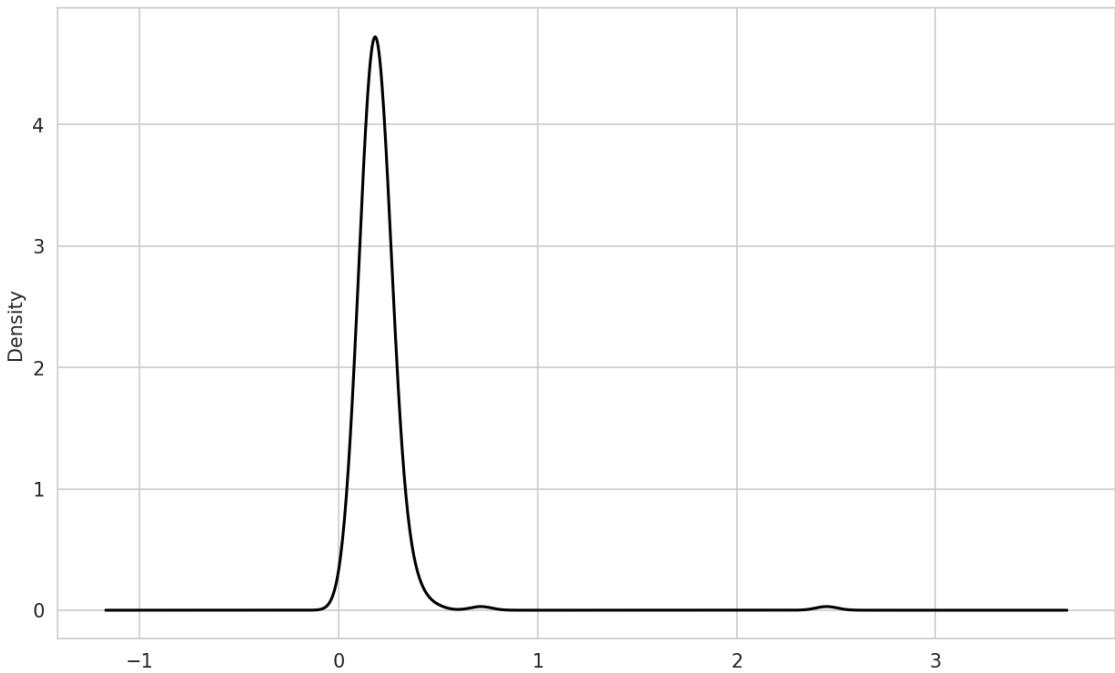


Figure 9-22. Density plot of tip percentages

This kind of plot requires SciPy, so if you do not have it installed already, you can pause and do that now:

conda install scipy

Explain this code

Translate into

Translate into

seaborn makes histograms and density plots even easier through its histplot method, which can plot both a histogram and a continuous density estimate simultaneously. As an example, consider a bimodal distribution consisting of draws from two different standard normal distributions (see [Figure 9-23](#)):

In [100]: `comp1 = np.random.standard_normal(200)`

In [101]: `comp2 = 10 + 2 * np.random.standard_normal(200)`

In [102]: `values = pd.Series(np.concatenate([comp1, comp2]))`

In [103]: `sns.histplot(values, bins=100, color="black")`

Explain this code

Translate into

Translate into

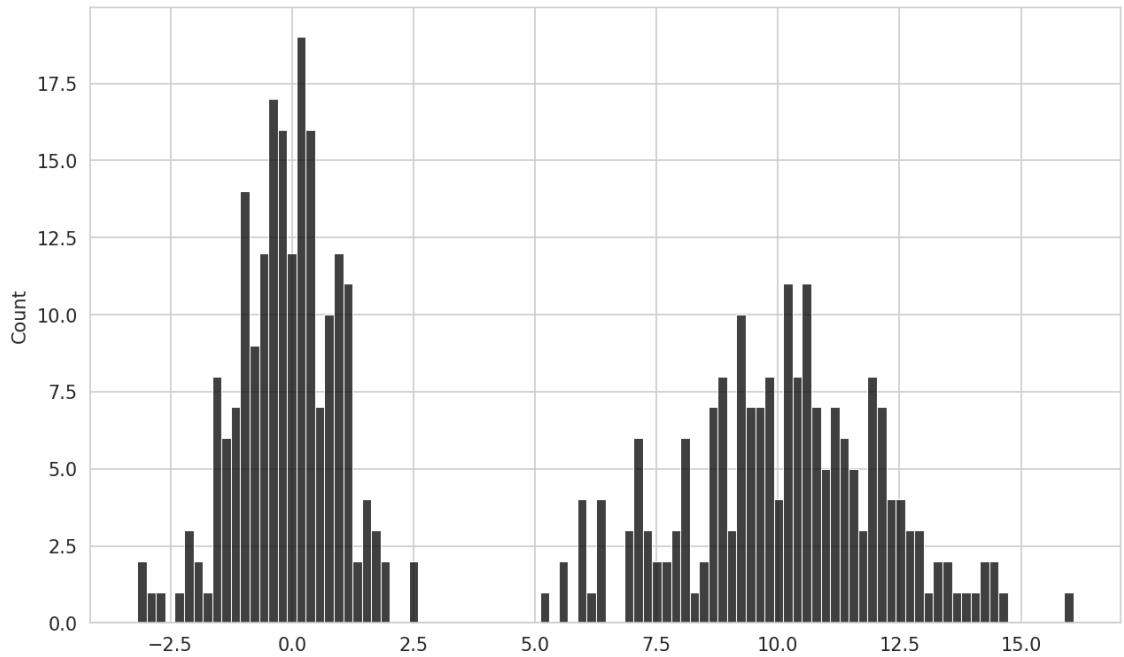


Figure 9-23. Normalized histogram of normal mixture

Scatter or Point Plots

Point plots or scatter plots can be a useful way of examining the relationship between two one-dimensional data series. For example, here we load the macrodata dataset from the statsmodels project, select a few variables, then compute log differences:

```
In [104]: macro = pd.read_csv("examples/macrodata.csv")
```

```
In [105]: data = macro[["cpi", "m1", "tbilrate", "unemp"]]
```

```
In [106]: trans_data = np.log(data).diff().dropna()
```

```
In [107]: trans_data.tail()
```

Out[107]:

	cpi	m1	tbilrate	unemp
198	-0.007904	0.045361	-0.396881	0.105361
199	-0.021979	0.066753	-2.277267	0.139762
200	0.002340	0.010286	0.606136	0.160343
201	0.008419	0.037461	-0.200671	0.127339
202	0.008894	0.012202	-0.405465	0.042560

Explain this code

Translate into

Translate into

We can then use seaborn's regplot method, which makes a scatter plot and fits a linear regression line (see [Figure 9-24](#)):

```
In [109]: ax = sns.regplot(x="m1", y="unemp", data=trans_data)
```

```
In [110]: ax.title("Changes in log(m1) versus log(unemp)")
```

Explain this code

Translate into

Translate into

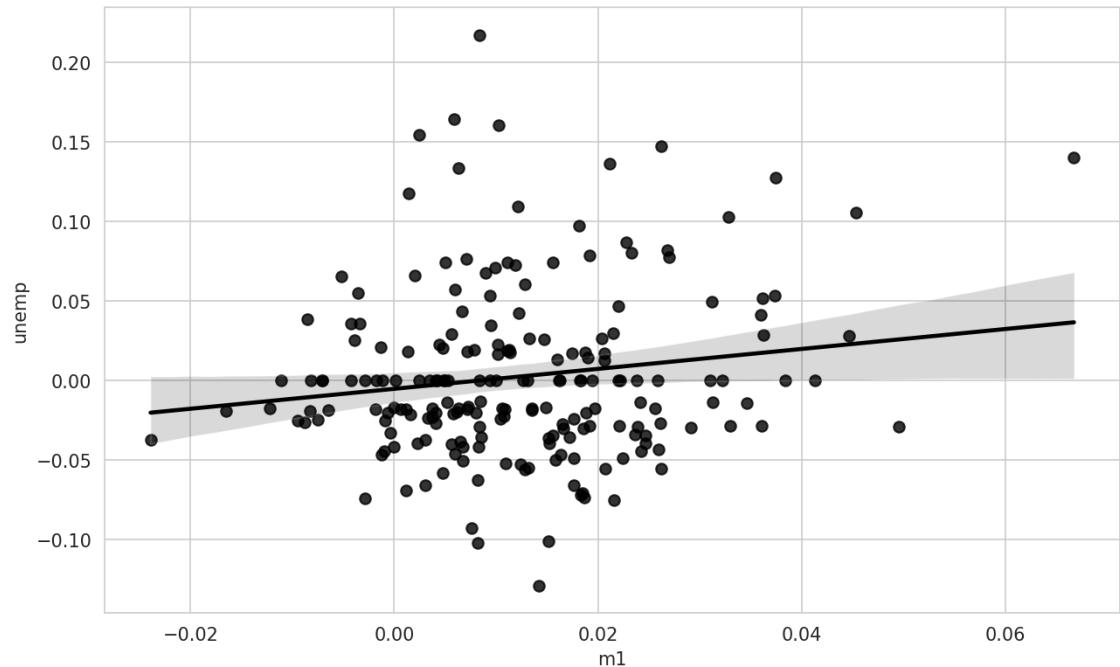


Figure 9-24. A seaborn regression/scatter plot

In exploratory data analysis, it's helpful to be able to look at all the scatter plots among a group of variables; this is known as a *pairs* plot or *scatter plot matrix*. Making such a plot from scratch is a bit of work, so seaborn has a convenient pairplot function that supports placing histograms or density estimates of each variable along the diagonal (see [Figure 9-25](#) for the resulting plot):

```
In [111]: sns.pairplot(trans_data, diag_kind="kde", plot_kws={"alpha": 0.2})
```

Explain this code

Translate into

Translate into

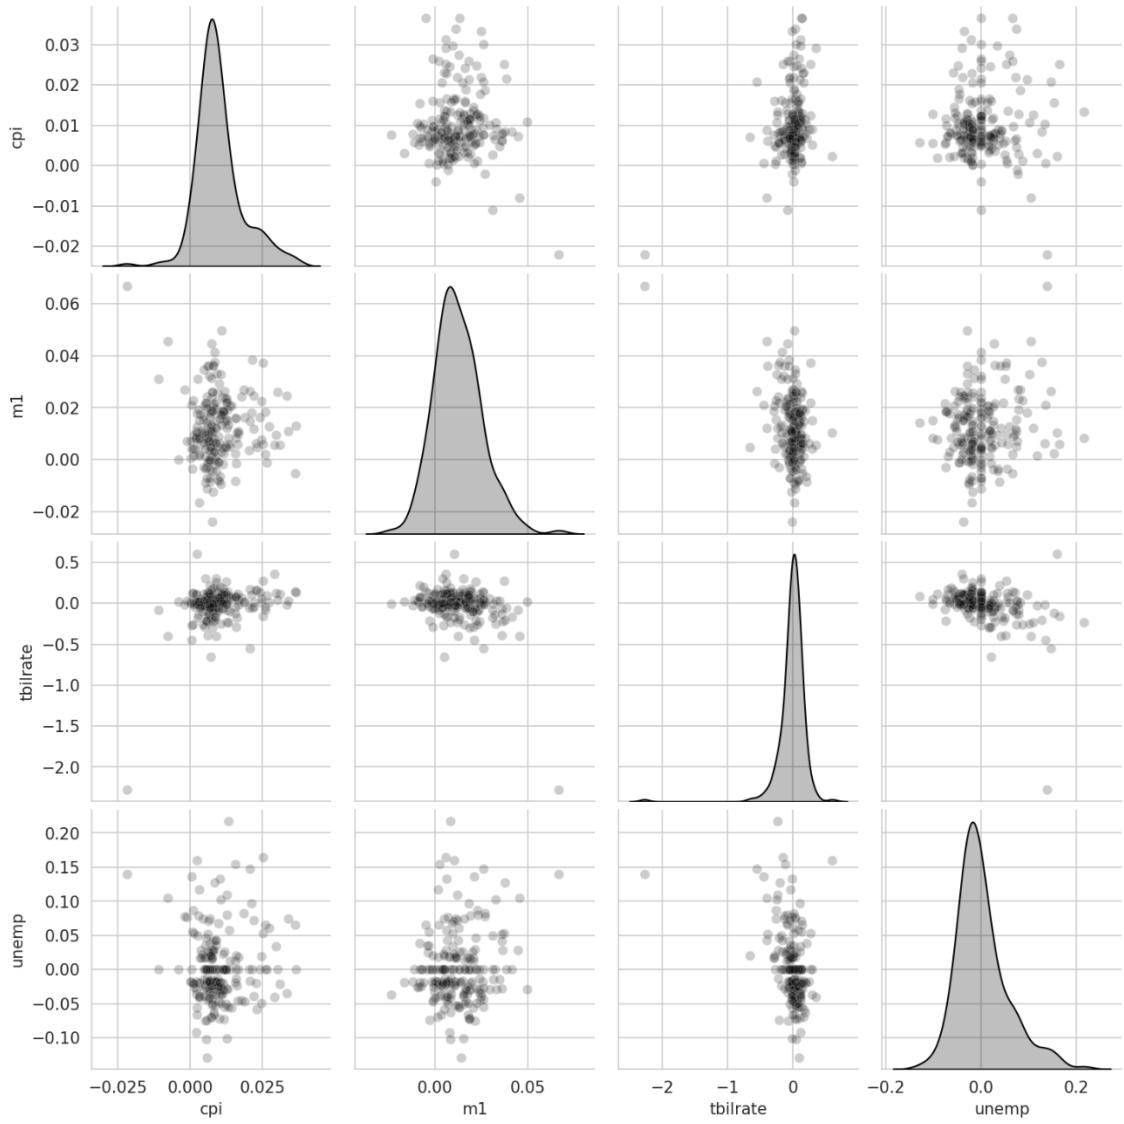


Figure 9-25. Pair plot matrix of statsmodels macro data

You may notice the `plot_kws` argument. This enables us to pass down configuration options to the individual plotting calls on the off-diagonal elements. Check out the `seaborn.pairplot` docstring for more granular configuration options.

Facet Grids and Categorical Data

What about datasets where we have additional grouping dimensions? One way to visualize data with many categorical variables is to use a *facet grid*, which is a two-dimensional layout of plots where the data is split across the plots on each axis based on the distinct values of a certain variable. `seaborn` has a useful built-in function `catplot` that simplifies making many kinds of faceted plots split by categorical variables (see [Figure 9-26](#) for the resulting plot):

```
In [112]: sns.catplot(x="day", y="tip_pct", hue="time", col="smoker",
.....:     kind="bar", data=tips[tips.tip_pct < 1])
```

Explain this code

Translate into

Translate into

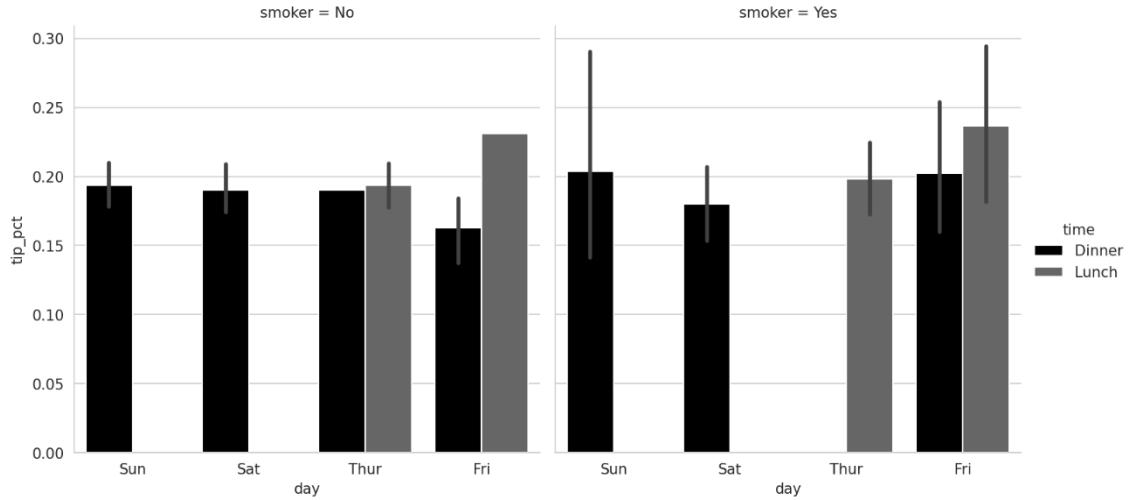


Figure 9-26. Tipping percentage by day/time/smoker

Instead of grouping by "time" by different bar colors within a facet, we can also expand the facet grid by adding one row per time value (see [Figure 9-27](#)):

```
In [113]: sns.catplot(x="day", y="tip_pct", row="time",  
.....:     col="smoker",  
.....:     kind="bar", data=tips[tips.tip_pct < 1])
```

Explain this code

Translate into

Translate into

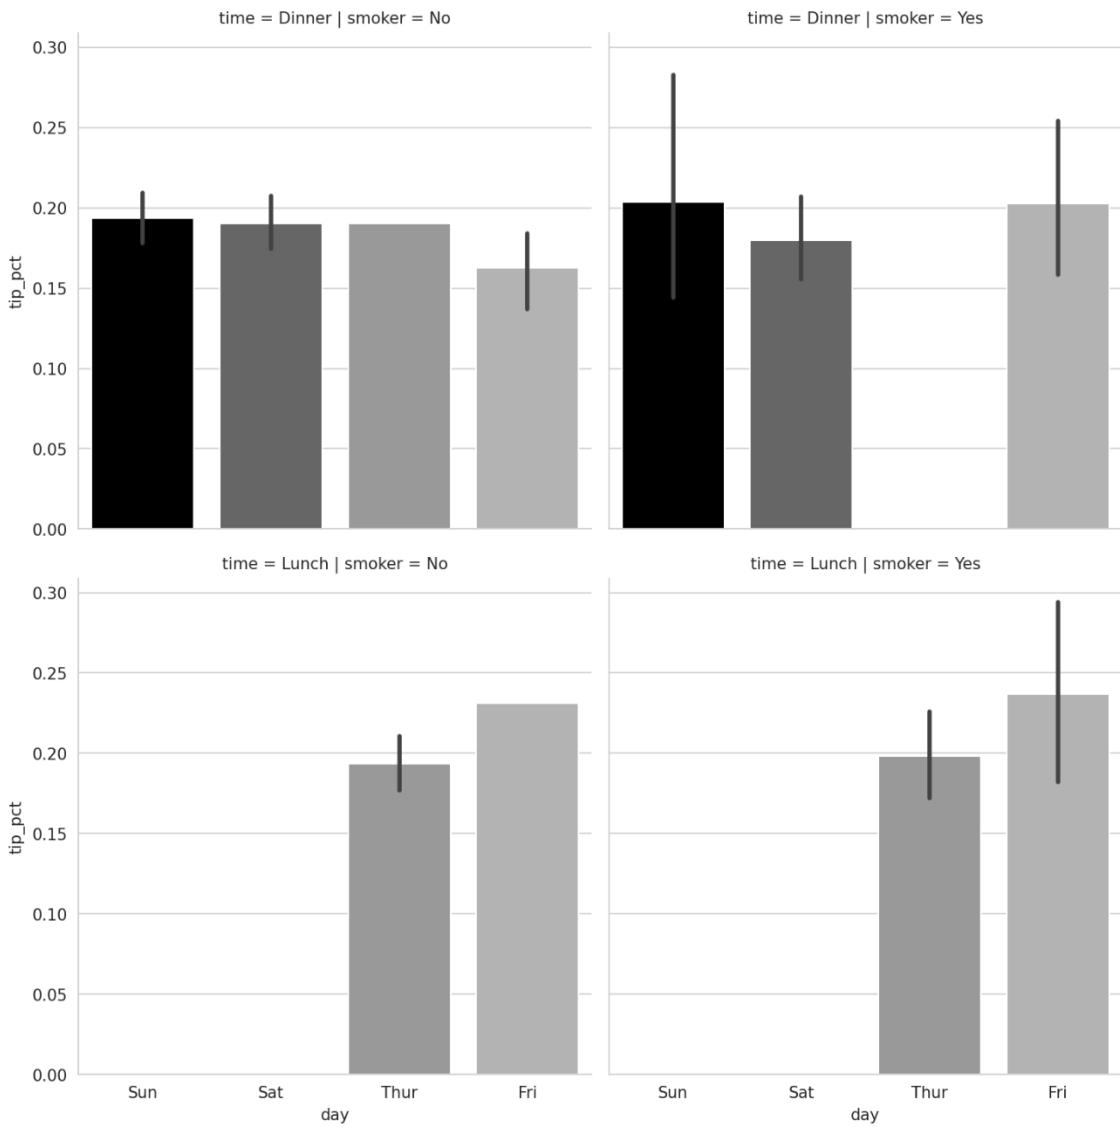


Figure 9-27. Tipping percentage by day split by time/smoker

catplot supports other plot types that may be useful depending on what you are trying to display. For example, *box plots* (which show the median, quartiles, and outliers) can be an effective visualization type (see [Figure 9-28](#)):

In [114]: `sns.catplot(x="tip_pct", y="day", kind="box",`

.....: `data=tips[tips.tip_pct < 0.5]`

Explain this code

Translate into

Translate into

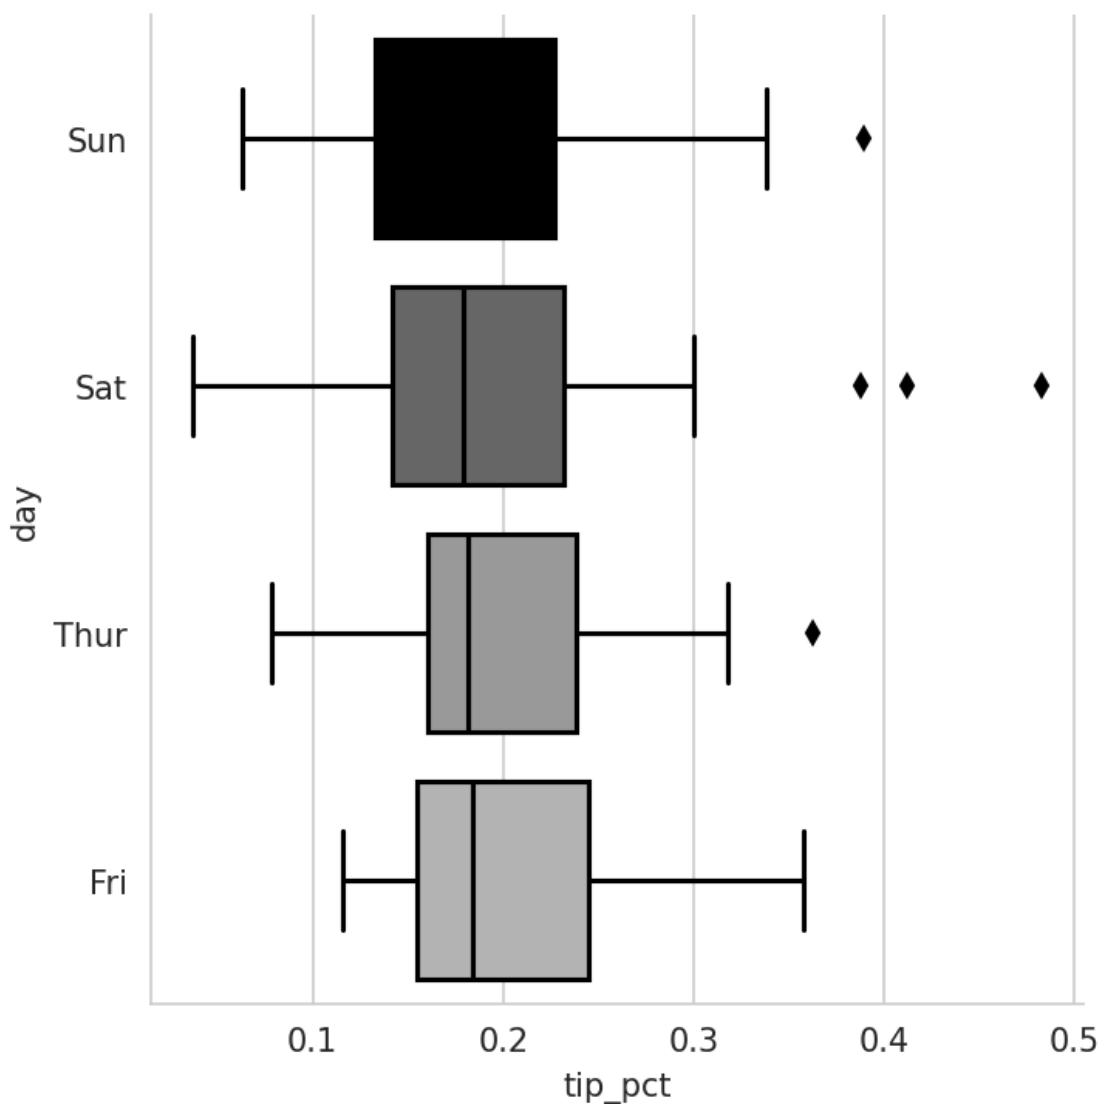


Figure 9-28. Box plot of tipping percentage by day

You can create your own facet grid plots using the more general `seaborn.FacetGrid` class. See the [seaborn documentation](#) for more.

9.3 Other Python Visualization Tools

As is common with open source, there many options for creating graphics in Python (too many to list). Since 2010, much development effort has been focused on creating interactive graphics for publication on the web. With tools like [Altair](#), [Bokeh](#), and [Plotly](#), it's now possible to specify dynamic, interactive graphics in Python that are intended for use with web browsers.

For creating static graphics for print or web, I recommend using `matplotlib` and libraries that build on `matplotlib`, like `pandas` and `seaborn`, for your needs. For other data visualization requirements, it may be useful to learn how to use one of the other available tools. I encourage you to explore the ecosystem as it continues to evolve and innovate into the future.

An excellent book on data visualization is *Fundamentals of Data Visualization* by Claus O. Wilke (O'Reilly), which is available in print or on Claus's website at <https://clauswilke.com/dataviz>.

9.4 Conclusion

The goal of this chapter was to get your feet wet with some basic data visualization using pandas, matplotlib, and seaborn. If visually communicating the results of data analysis is important in your work, I encourage you to seek out resources to learn more about effective data visualization. It is an active field of research, and you can practice with many excellent learning resources available online and in print.

In the next chapter, we turn our attention to data aggregation and group operations with pandas.