

**Python for Data
Analysis, 3rd Edition**

[Wes McKinney](#)

Published by O'Reilly
Media, Inc.

<https://github.com/wesm/pydata-book>

Chapter 10. Data Aggregation and Group Operations

Categorizing a dataset and applying a function to each group, whether an aggregation or transformation, can be a critical component of a data analysis workflow. After loading, merging, and preparing a dataset, you may need to compute group statistics or possibly *pivot tables* for reporting or visualization purposes. pandas provides a versatile groupby interface, enabling you to slice, dice, and summarize datasets in a natural way.

One reason for the popularity of relational databases and SQL (which stands for “structured query language”) is the ease with which data can be joined, filtered, transformed, and aggregated. However, query languages like SQL impose certain limitations on the kinds of group operations that can be performed. As you will see, with the expressiveness of Python and pandas, we can perform quite complex group operations by expressing them as custom Python functions that manipulate the data associated with each group. In this chapter, you will learn how to:

- Split a pandas object into pieces using one or more keys (in the form of functions, arrays, or DataFrame column names)
- Calculate group summary statistics, like count, mean, or standard deviation, or a user-defined function
- Apply within-group transformations or other manipulations, like normalization, linear regression, rank, or subset selection
- Compute pivot tables and cross-tabulations
- Perform quantile analysis and other statistical group analyses

Note

Time-based aggregation of time series data, a special use case of groupby, is referred to as *resampling* in this book and will receive separate treatment in [Chapter 11](#).

As with the rest of the chapters, we start by importing NumPy and pandas:

In [12]: `import numpy as np`

In [13]: `import pandas as pd`

10.1 How to Think About Group Operations

Hadley Wickham, an author of many popular packages for the R programming language, coined the term *split-apply-combine* for describing group operations. In the first stage of the process, data contained in a pandas object, whether a Series, DataFrame, or otherwise, is *split* into groups based on one or more *keys* that you provide. The splitting is performed on a particular axis of an object. For example, a DataFrame can be grouped on its rows (`axis="index"`) or its columns (`axis="columns"`). Once this is done, a function is *applied* to each group, producing a new value. Finally, the results of all those function applications are *combined* into a result object. The form of the resulting object will usually depend on what's being done to the data. See [Figure 10-1](#) for a mockup of a simple group aggregation.

Each grouping key can take many forms, and the keys do not have to be all of the same type:

- A list or array of values that is the same length as the axis being grouped
- A value indicating a column name in a DataFrame
- A dictionary or Series giving a correspondence between the values on the axis being grouped and the group names
- A function to be invoked on the axis index or the individual labels in the index

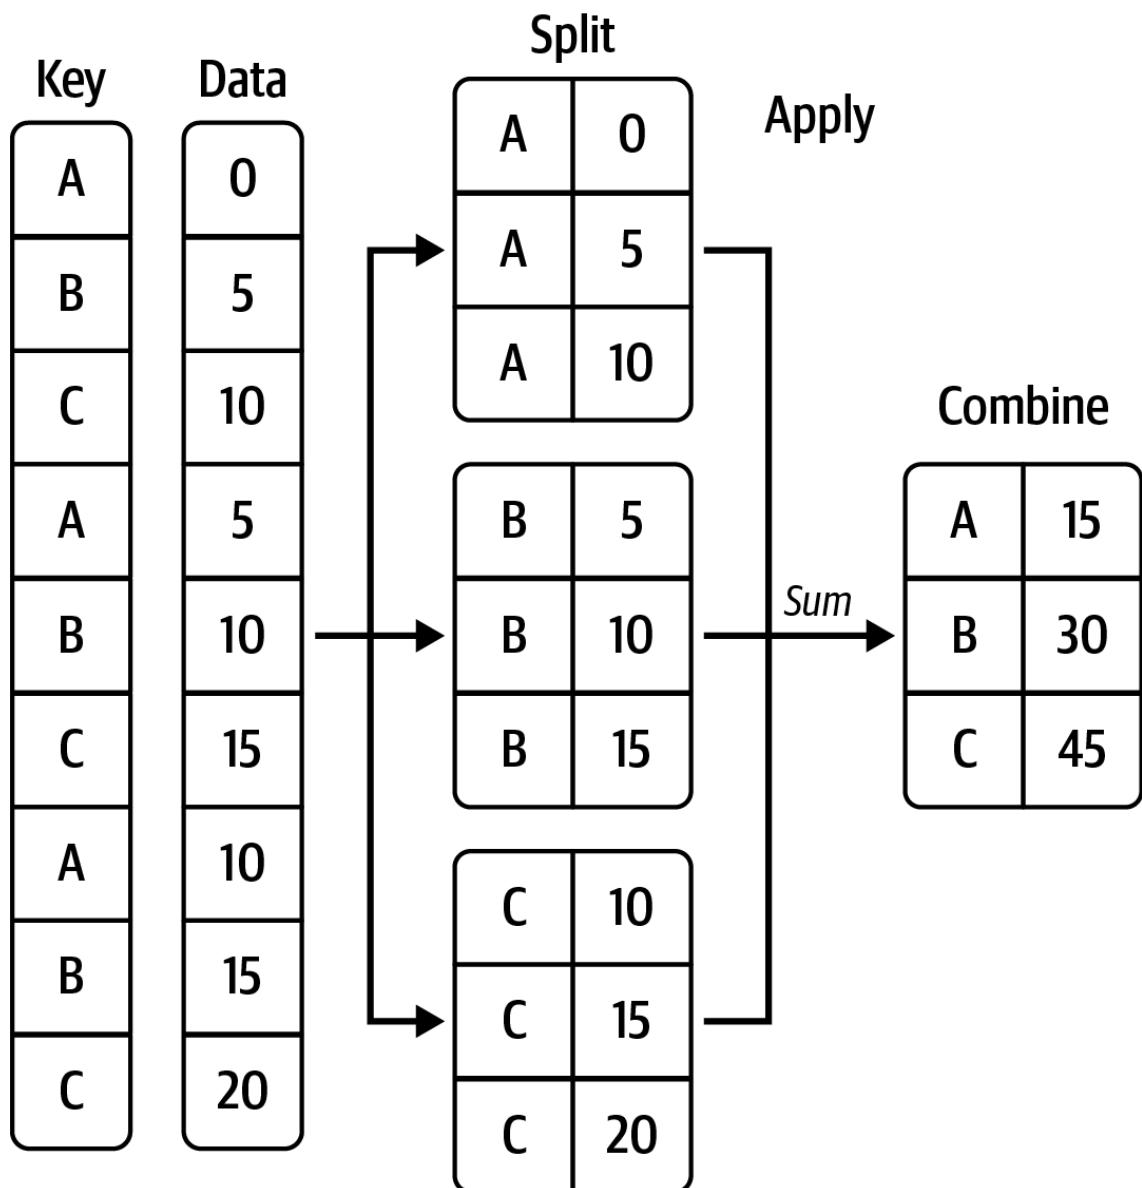


Figure 10-1. Illustration of a group aggregation

Note that the latter three methods are shortcuts for producing an array of values to be used to split up the object. Don't worry if this all seems abstract. Throughout this chapter, I will give many examples of all these methods. To get started, here is a small tabular dataset as a DataFrame:

```
In [14]: df = pd.DataFrame({"key1" : ["a", "a", None, "b", "b", "a", None],
....:                 "key2" : pd.Series([1, 2, 1, 2, 1, None, 1], dtype="I
nt64"),
....:                 "data1" : np.random.standard_normal(7),
....:                 "data2" : np.random.standard_normal(7)})
```

In [15]: df

Out[15]:

```
key1 key2  data1  data2
0  a  1 -0.204708  0.281746
1  a  2  0.478943  0.769023
2 None  1 -0.519439  1.246435
3  b  2 -0.555730  1.007189
4  b  1  1.965781 -1.296221
5  a <NA>  1.393406  0.274992
6 None  1  0.092908  0.228913
```

Suppose you wanted to compute the mean of the data1 column using the labels from key1. There are a number of ways to do this. One is to access data1 and call groupby with the column (a Series) at key1:

In [16]: grouped = df["data1"].groupby(df["key1"])

In [17]: grouped

Out[17]: <pandas.core.groupby.generic.SeriesGroupBy object at 0x7f4b76420a00>

This grouped variable is now a special “*GroupBy*” object. It has not actually computed anything yet except for some intermediate data about the group key df["key1"]. The idea is that this object has all of the information needed to then apply some operation to each of the groups. For example, to compute group means we can call the GroupBy's mean method:

In [18]: grouped.mean()

Out[18]:

key1

```
a  0.555881
b  0.705025
```

Name: data1, dtype: float64

Later in [Section 10.2, “Data Aggregation,”](#) I’ll explain more about what happens when you call `.mean()`. The important thing here is that the data (a Series) has been aggregated by splitting the data on the group key, producing a new Series that is now indexed by the unique values in the `key1` column. The result index has the name “`key1`” because the DataFrame column `df["key1"]` did.

If instead we had passed multiple arrays as a list, we’d get something different:

```
In [19]: means = df["data1"].groupby([df["key1"], df["key2"]]).mean()
```

```
In [20]: means
```

```
Out[20]:
```

```
key1 key2
```

```
a 1 -0.204708
```

```
2 0.478943
```

```
b 1 1.965781
```

```
2 -0.555730
```

```
Name: data1, dtype: float64
```

Here we grouped the data using two keys, and the resulting Series now has a hierarchical index consisting of the unique pairs of keys observed:

```
In [21]: means.unstack()
```

```
Out[21]:
```

```
key2    1    2
```

```
key1
```

```
a -0.204708 0.478943
```

```
b 1.965781 -0.555730
```

In this example, the group keys are all Series, though they could be any arrays of the right length:

```
In [22]: states = np.array(["OH", "CA", "CA", "OH", "OH", "CA", "OH"])
```

```
In [23]: years = [2005, 2005, 2006, 2005, 2006, 2005, 2006]
```

```
In [24]: df["data1"].groupby([states, years]).mean()
```

Out[24]:

```
CA 2005  0.936175
```

```
      2006 -0.519439
```

```
OH 2005 -0.380219
```

```
      2006  1.029344
```

```
Name: data1, dtype: float64
```

Frequently, the grouping information is found in the same DataFrame as the data you want to work on. In that case, you can pass column names (whether those are strings, numbers, or other Python objects) as the group keys:

```
In [25]: df.groupby("key1").mean()
```

Out[25]:

```
key2  data1  data2
```

```
key1
```

```
  a   1.5  0.555881  0.441920
```

```
  b   1.5  0.705025 -0.144516
```

```
In [26]: df.groupby("key2").mean()
```

Out[26]:

```
data1  data2
```

```
key2
```

```
  1   0.333636  0.115218
```

```
  2  -0.038393  0.888106
```

```
In [27]: df.groupby(["key1", "key2"]).mean()
```

```
Out[27]:
```

```
data1  data2  
key1 key2  
a 1 -0.204708 0.281746  
   2 0.478943 0.769023  
b 1 1.965781 -1.296221  
   2 -0.555730 1.007189
```

You may have noticed in the second case, `df.groupby("key2").mean()`, that there is no `key1` column in the result. Because `df["key1"]` is not numeric data, it is said to be a *nuisance column*, which is therefore automatically excluded from the result. By default, all of the numeric columns are aggregated, though it is possible to filter down to a subset, as you'll see soon.

Regardless of the objective in using `groupby`, a generally useful `GroupBy` method is `size`, which returns a `Series` containing group sizes:

```
In [28]: df.groupby(["key1", "key2"]).size()
```

```
Out[28]:
```

```
key1 key2  
a 1 1  
   2 1  
b 1 1  
   2 1  
dtype: int64
```

Note that any missing values in a group key are excluded from the result by default. This behavior can be disabled by passing `dropna=False` to `groupby`:

```
In [29]: df.groupby("key1", dropna=False).size()
```

```
Out[29]:
```

```
key1  
a    3  
b    2  
NaN  2  
dtype: int64
```

```
In [30]: df.groupby(["key1", "key2"], dropna=False).size()
```

```
Out[30]:
```

```
key1 key2  
a   1   1  
     2   1  
      <NA> 1  
b   1   1  
     2   1  
NaN  1   2  
dtype: int64
```

A group function similar in spirit to size is count, which computes the number of nonnull values in each group:

```
In [31]: df.groupby("key1").count()
```

```
Out[31]:
```

```
key2 data1 data2  
key1  
a   2   3   3  
b   2   2   2
```

Iterating over Groups

The object returned by groupby supports iteration, generating a sequence of 2-tuples containing the group name along with the chunk of data. Consider the following:

In [32]: for name, group in df.groupby("key1"):

```
....: print(name)
....: print(group)
....:

a
key1 key2  data1  data2
0  a  1 -0.204708  0.281746
1  a  2  0.478943  0.769023
5  a <NA>  1.393406  0.274992

b
key1 key2  data1  data2
3  b  2 -0.555730  1.007189
4  b  1  1.965781 -1.296221
```

In the case of multiple keys, the first element in the tuple will be a tuple of key values:

In [33]: for (k1, k2), group in df.groupby(["key1", "key2"]):

```
....: print((k1, k2))
....: print(group)
....:

('a', 1)
key1 key2  data1  data2
0  a  1 -0.204708  0.281746

('a', 2)
key1 key2  data1  data2
1  a  2  0.478943  0.769023

('b', 1)
```

```
key1 key2  data1  data2
4  b   1 1.965781 -1.296221
('b', 2)
key1 key2  data1  data2
3  b   2 -0.55573  1.007189
```

Of course, you can choose to do whatever you want with the pieces of data. A recipe you may find useful is computing a dictionary of the data pieces as a one-liner:

```
In [34]: pieces = {name: group for name, group in df.groupby("key1")}
```

```
In [35]: pieces["b"]
```

Out[35]:

```
key1 key2  data1  data2
3  b   2 -0.555730  1.007189
4  b   1 1.965781 -1.296221
```

By default groupby groups on axis="index", but you can group on any of the other axes. For example, we could group the columns of our example df here by whether they start with "key" or "data":

```
In [36]: grouped = df.groupby({"key1": "key", "key2": "key",
....:                      "data1": "data", "data2": "data"}, axis="columns")
```

We can print out the groups like so:

```
In [37]: for group_key, group_values in grouped:
```

```
....:     print(group_key)
....:     print(group_values)
```

```
....:  
data  
    data1  data2  
0 -0.204708 0.281746  
1 0.478943 0.769023  
2 -0.519439 1.246435  
3 -0.555730 1.007189  
4 1.965781 -1.296221  
5 1.393406 0.274992  
6 0.092908 0.228913  
  
key  
    key1 key2  
0  a  1  
1  a  2  
2 None 1  
3  b  2  
4  b  1  
5  a <NA>  
6 None 1
```

Selecting a Column or Subset of Columns

Indexing a GroupBy object created from a DataFrame with a column name or array of column names has the effect of column subsetting for aggregation. This means that:

```
df.groupby("key1")["data1"]  
df.groupby("key1")[[ "data2" ]]
```

are conveniences for:

```
df["data1"].groupby(df["key1"])
df[["data2"]].groupby(df["key1"])
```

Especially for large datasets, it may be desirable to aggregate only a few columns. For example, in the preceding dataset, to compute the means for just the data2 column and get the result as a DataFrame, we could write:

```
In [38]: df.groupby(["key1", "key2"])["data2"].mean()
```

```
Out[38]:
```

```
data2
key1 key2
a  1    0.281746
   2    0.769023
b  1    -1.296221
   2    1.007189
```

The object returned by this indexing operation is a grouped DataFrame if a list or array is passed, or a grouped Series if only a single column name is passed as a scalar:

```
In [39]: s_grouped = df.groupby(["key1", "key2"])["data2"]
```

```
In [40]: s_grouped
```

```
Out[40]: <pandas.core.groupby.generic.SeriesGroupBy object at 0x7f4b76423340>
```

```
In [41]: s_grouped.mean()
```

```
Out[41]:
```

```
key1 key2
a  1    0.281746
   2    0.769023
b  1    -1.296221
```

```
2    1.007189
```

```
Name: data2, dtype: float64
```

Grouping with Dictionaries and Series

Grouping information may exist in a form other than an array. Let's consider another example DataFrame:

```
In [42]: people = pd.DataFrame(np.random.standard_normal((5, 5)),  
....:             columns=["a", "b", "c", "d", "e"],  
....:             index=["Joe", "Steve", "Wanda", "Jill", "Trey"])
```

```
In [43]: people.iloc[2:3, [1, 2]] = np.nan # Add a few NA values
```

```
In [44]: people
```

```
Out[44]:
```

	a	b	c	d	e
Joe	1.352917	0.886429	-2.001637	-0.371843	1.669025
Steve	-0.438570	-0.539741	0.476985	3.248944	-1.021228
Wanda	-0.577087	NaN	NaN	0.523772	0.000940
Jill	1.343810	-0.713544	-0.831154	-2.370232	-1.860761
Trey	-0.860757	0.560145	-1.265934	0.119827	-1.063512

Now, suppose I have a group correspondence for the columns and want to sum the columns by group:

```
In [45]: mapping = {"a": "red", "b": "red", "c": "blue",  
....:             "d": "blue", "e": "red", "f": "orange"}
```

Now, you could construct an array from this dictionary to pass to groupby, but instead we can just pass the dictionary (I included the key "f" to highlight that unused grouping keys are OK):

```
In [46]: by_column = people.groupby(mapping, axis="columns")
```

```
In [47]: by_column.sum()
```

```
Out[47]:
```

```
blue    red  
Joe   -2.373480  3.908371  
Steve  3.725929 -1.999539  
Wanda  0.523772 -0.576147  
Jill   -3.201385 -1.230495  
Trey   -1.146107 -1.364125
```

The same functionality holds for Series, which can be viewed as a fixed-size mapping:

```
In [48]: map_series = pd.Series(mapping)
```

```
In [49]: map_series
```

```
Out[49]:
```

```
a    red  
b    red  
c    blue  
d    blue  
e    red  
f    orange  
dtype: object
```

```
In [50]: people.groupby(map_series, axis="columns").count()
```

```
Out[50]:
```

```
blue red
```

```
Joe    2  3
Steve   2  3
Wanda   1  2
Jill    2  3
Trey    2  3
```

Grouping with Functions

Using Python functions is a more generic way of defining a group mapping compared with a dictionary or Series. Any function passed as a group key will be called once per index value (or once per column value if using axis="columns"), with the return values being used as the group names. More concretely, consider the example DataFrame from the previous section, which has people's first names as index values. Suppose you wanted to group by name length. While you could compute an array of string lengths, it's simpler to just pass the len function:

```
In [51]: people.groupby(len).sum()
```

```
Out[51]:
```

```
      a      b      c      d      e
3  1.352917  0.886429 -2.001637 -0.371843  1.669025
4  0.483052 -0.153399 -2.097088 -2.250405 -2.924273
5 -1.015657 -0.539741  0.476985  3.772716 -1.020287
```

Mixing functions with arrays, dictionaries, or Series is not a problem, as everything gets converted to arrays internally:

```
In [52]: key_list = ["one", "one", "one", "two", "two"]
```

```
In [53]: people.groupby([len, key_list]).min()
```

```
Out[53]:
```

```
      a      b      c      d      e
3 one  1.352917  0.886429 -2.001637 -0.371843  1.669025
4 two -0.860757 -0.713544 -1.265934 -2.370232 -1.860761
```

```
5 one -0.577087 -0.539741 0.476985 0.523772 -1.021228
```

Grouping by Index Levels

A final convenience for hierarchically indexed datasets is the ability to aggregate using one of the levels of an axis index. Let's look at an example:

```
In [54]: columns = pd.MultiIndex.from_arrays([["US", "US", "US", "JP", "JP"],  
.....: [1, 3, 5, 1, 3]],  
.....: names=["cty", "tenor"])
```

```
In [55]: hier_df = pd.DataFrame(np.random.standard_normal((4, 5)), columns=columns)
```

```
In [56]: hier_df
```

```
Out[56]:
```

	cty	US		JP		
	tenor	1	3	5	1	3
0		0.332883	-2.359419	-0.199543	-1.541996	-0.970736
1		-1.307030	0.286350	0.377984	-0.753887	0.331286
2		1.349742	0.069877	0.246674	-0.011862	1.004812
3		1.327195	-0.919262	-1.549106	0.022185	0.758363

To group by level, pass the level number or name using the level keyword:

```
In [57]: hier_df.groupby(level="cty", axis="columns").count()
```

```
Out[57]:
```

	cty	JP	US
0	2	3	
1	2	3	

2 2 3

3 2 3

10.2 Data Aggregation

Aggregations refer to any data transformation that produces scalar values from arrays. The preceding examples have used several of them, including mean, count, min, and sum. You may wonder what is going on when you invoke mean() on a GroupBy object. Many common aggregations, such as those found in [Table 10-1](#), have optimized implementations. However, you are not limited to only this set of methods.

Function name	Description
any, all	Return True if any (one or more values) or all non-NA values are “truthy”
count	Number of non-NA values
cummin, cummax	Cumulative minimum and maximum of non-NA values
cumsum	Cumulative sum of non-NA values
cumprod	Cumulative product of non-NA values
first, last	First and last non-NA values
mean	Mean of non-NA values
median	Arithmetic median of non-NA values
min, max	Minimum and maximum of non-NA values

Function name	Description
nth	Retrieve value that would appear at position n with the data in sorted order
ohlc	Compute four “open-high-low-close” statistics for time series-like data
prod	Product of non-NA values
quantile	Compute sample quantile
rank	Ordinal ranks of non-NA values, like calling Series.rank
size	Compute group sizes, returning result as a Series
sum	Sum of non-NA values
std, var	Sample standard deviation and variance

Table 10-1. Optimized groupby methods

You can use aggregations of your own devising and additionally call any method that is also defined on the object being grouped. For example, the nsmallest Series method selects the smallest requested number of values from the data. While nsmallest is not explicitly implemented for GroupBy, we can still use it with a nonoptimized implementation. Internally, GroupBy slices up the Series, calls piece.nsmallest(n) for each piece, and then assembles those results into the result object:

In [58]: df

Out[58]:

```
key1 key2  data1  data2
0   a    1 -0.204708 0.281746
1   a    2  0.478943 0.769023
2 None  1 -0.519439 1.246435
```

```
3  b  2 -0.555730 1.007189
4  b  1 1.965781 -1.296221
5  a <NA> 1.393406 0.274992
6 None 1 0.092908 0.228913
```

In [59]: grouped = df.groupby("key1")

In [60]: grouped["data1"].nsmallest(2)

Out[60]:

```
key1
a    0   -0.204708
      1   0.478943
b    3   -0.555730
      4   1.965781
Name: data1, dtype: float64
```

To use your own aggregation functions, pass any function that aggregates an array to the aggregate method or its short alias agg:

In [61]: def peak_to_peak(arr):

```
....:     return arr.max() - arr.min()
```

In [62]: grouped.agg(peak_to_peak)

Out[62]:

```
key2  data1  data2
key1
a    1 1.598113 0.494031
b    1 2.521511 2.303410
```

You may notice that some methods, like `describe`, also work, even though they are not aggregations, strictly speaking:

In [63]: `grouped.describe()`

Out[63]:

```
key2          data1      ... \
count mean   std min 25% 50% 75% max count  mean ...
key1
...
a  2.0 1.5 0.707107 1.0 1.25 1.5 1.75 2.0 3.0 0.555881 ...
b  2.0 1.5 0.707107 1.0 1.25 1.5 1.75 2.0 2.0 0.705025 ...
data2          \
75%  max count  mean   std   min   25%
key1
a  0.936175 1.393406 3.0 0.441920 0.283299 0.274992 0.278369
b  1.335403 1.965781 2.0 -0.144516 1.628757 -1.296221 -0.720368

50% 75%  max
key1
a  0.281746 0.525384 0.769023
b  -0.144516 0.431337 1.007189
[2 rows x 24 columns]
```

I will explain in more detail what has happened here in [Section 10.3, “Apply: General split-apply-combine.”](#)

Note

Custom aggregation functions are generally much slower than the optimized functions found in [Table 10-1](#). This is because there is some extra overhead (function calls, data rearrangement) in constructing the intermediate group data chunks.

Column-Wise and Multiple Function Application

Let’s return to the tipping dataset used in the last chapter. After loading it with `pandas.read_csv`, we add a tipping percentage column:

```
In [64]: tips = pd.read_csv("examples/tips.csv")
```

```
In [65]: tips.head()
```

```
Out[65]:
```

```
total_bill tip smoker day time size  
0 16.99 1.01 No Sun Dinner 2  
1 10.34 1.66 No Sun Dinner 3  
2 21.01 3.50 No Sun Dinner 3  
3 23.68 3.31 No Sun Dinner 2  
4 24.59 3.61 No Sun Dinner 4
```

Now I will add a tip_pct column with the tip percentage of the total bill:

```
In [66]: tips["tip_pct"] = tips["tip"] / tips["total_bill"]
```

```
In [67]: tips.head()
```

```
Out[67]:
```

```
total_bill tip smoker day time size tip_pct  
0 16.99 1.01 No Sun Dinner 2 0.059447  
1 10.34 1.66 No Sun Dinner 3 0.160542  
2 21.01 3.50 No Sun Dinner 3 0.166587  
3 23.68 3.31 No Sun Dinner 2 0.139780  
4 24.59 3.61 No Sun Dinner 4 0.146808
```

As you've already seen, aggregating a Series or all of the columns of a DataFrame is a matter of using aggregate (or agg) with the desired function or calling a method like mean or std. However, you may want to aggregate using a different function, depending on the column, or multiple functions at once. Fortunately, this is possible to do, which I'll illustrate through a number of examples. First, I'll group the tips by day and smoker:

```
In [68]: grouped = tips.groupby(["day", "smoker"])
```

Note that for descriptive statistics like those in [Table 10-1](#), you can pass the name of the function as a string:

```
In [69]: grouped_pct = grouped["tip_pct"]
```

```
In [70]: grouped_pct.agg("mean")
```

```
Out[70]:
```

day	smoker	tip_pct
Fri	No	0.151650
	Yes	0.174783
Sat	No	0.158048
	Yes	0.147906
Sun	No	0.160113
	Yes	0.187250
Thur	No	0.160298
	Yes	0.163863

```
Name: tip_pct, dtype: float64
```

If you pass a list of functions or function names instead, you get back a DataFrame with column names taken from the functions:

```
In [71]: grouped_pct.agg(["mean", "std", "peak_to_peak"])
```

```
Out[71]:
```

		mean	std	peak_to_peak
	day	smoker		
Fri	No	0.151650	0.028123	0.067349
	Yes	0.174783	0.051293	0.159925
Sat	No	0.158048	0.039767	0.235193

```
Yes 0.147906 0.061375 0.290095
Sun No 0.160113 0.042347 0.193226
    Yes 0.187250 0.154134 0.644685
Thur No 0.160298 0.038774 0.193350
    Yes 0.163863 0.039389 0.151240
```

Here we passed a list of aggregation functions to `agg` to evaluate independently on the data groups.

You don't need to accept the names that `GroupBy` gives to the columns; notably, lambda functions have the name "`<lambda>`", which makes them hard to identify (you can see for yourself by looking at a function's `__name__` attribute). Thus, if you pass a list of (name, function) tuples, the first element of each tuple will be used as the `DataFrame` column names (you can think of a list of 2-tuples as an ordered mapping):

```
In [72]: grouped_pct.agg([("average", "mean"), ("stdev", np.std)])
```

Out[72]:

```
average stdev
day smoker
Fri No 0.151650 0.028123
    Yes 0.174783 0.051293
Sat No 0.158048 0.039767
    Yes 0.147906 0.061375
Sun No 0.160113 0.042347
    Yes 0.187250 0.154134
Thur No 0.160298 0.038774
    Yes 0.163863 0.039389
```

With a `DataFrame` you have more options, as you can specify a list of functions to apply to all of the columns or different functions per column. To start, suppose we wanted to compute the same three statistics for the `tip_pct` and `total_bill` columns:

```
In [73]: functions = ["count", "mean", "max"]
```

```
In [74]: result = grouped[["tip_pct", "total_bill"]].agg(functions)
```

```
In [75]: result
```

```
Out[75]:
```

		tip_pct			total_bill		
		count	mean	max	count	mean	max
	day	smoker					
Fri	No	4	0.151650	0.187735	4	18.420000	22.75
	Yes	15	0.174783	0.263480	15	16.813333	40.17
Sat	No	45	0.158048	0.291990	45	19.661778	48.33
	Yes	42	0.147906	0.325733	42	21.276667	50.81
Sun	No	57	0.160113	0.252672	57	20.506667	48.17
	Yes	19	0.187250	0.710345	19	24.120000	45.35
Thur	No	45	0.160298	0.266312	45	17.113111	41.19
	Yes	17	0.163863	0.241255	17	19.190588	43.11

As you can see, the resulting DataFrame has hierarchical columns, the same as you would get aggregating each column separately and using concat to glue the results together using the column names as the keys argument:

```
In [76]: result["tip_pct"]
```

```
Out[76]:
```

		count	mean	max
	day	smoker		
Fri	No	4	0.151650	0.187735
	Yes	15	0.174783	0.263480
Sat	No	45	0.158048	0.291990
	Yes	42	0.147906	0.325733
Sun	No	57	0.160113	0.252672
	Yes	19	0.187250	0.710345

```
Thur No    45 0.160298 0.266312
```

```
Yes     17 0.163863 0.241255
```

As before, a list of tuples with custom names can be passed:

```
In [77]: ftuples = [("Average", "mean"), ("Variance", np.var)]
```

```
In [78]: grouped[["tip_pct", "total_bill"]].agg(ftuples)
```

```
Out[78]:
```

	tip_pct	total_bill	Average	Variance	Average	Variance
day	smoker					
Fri	No	0.151650	0.000791	18.420000	25.596333	
Yes		0.174783	0.002631	16.813333	82.562438	
Sat	No	0.158048	0.001581	19.661778	79.908965	
Yes		0.147906	0.003767	21.276667	101.387535	
Sun	No	0.160113	0.001793	20.506667	66.099980	
Yes		0.187250	0.023757	24.120000	109.046044	
Thur	No	0.160298	0.001503	17.113111	59.625081	
Yes		0.163863	0.001551	19.190588	69.808518	

Now, suppose you wanted to apply potentially different functions to one or more of the columns. To do this, pass a dictionary to agg that contains a mapping of column names to any of the function specifications listed so far:

```
In [79]: grouped.agg({"tip" : np.max, "size" : "sum"})
```

```
Out[79]:
```

	tip	size	
day	smoker		
Fri	No	3.50	9

```
Yes    4.73  31
Sat No   9.00 115
    Yes  10.00 104
Sun No   6.00 167
    Yes  6.50  49
Thur No   6.70 112
    Yes  5.00  40
```

```
In [80]: grouped.agg({"tip_pct" : ["min", "max", "mean", "std"],
....:           "size" : "sum"})
```

```
Out[80]:
```

	tip_pct	size			
	min	max	mean	std	sum
day	smoker				
Fri No	0.120385	0.187735	0.151650	0.028123	9
Yes	0.103555	0.263480	0.174783	0.051293	31
Sat No	0.056797	0.291990	0.158048	0.039767	115
Yes	0.035638	0.325733	0.147906	0.061375	104
Sun No	0.059447	0.252672	0.160113	0.042347	167
Yes	0.065660	0.710345	0.187250	0.154134	49
Thur No	0.072961	0.266312	0.160298	0.038774	112
Yes	0.090014	0.241255	0.163863	0.039389	40

A DataFrame will have hierarchical columns only if multiple functions are applied to at least one column.

Returning Aggregated Data Without Row Indexes

In all of the examples up until now, the aggregated data comes back with an index, potentially hierarchical, composed from the unique group key combinations. Since this isn't always desirable, you can disable this behavior in most cases by passing `as_index=False` to `groupby`:

```
In [81]: tips.groupby(["day", "smoker"], as_index=False).mean()
```

```
Out[81]:
```

	day	smoker	total_bill	tip	size	tip_pct
0	Fri	No	18.420000	2.812500	2.250000	0.151650
1	Fri	Yes	16.813333	2.714000	2.066667	0.174783
2	Sat	No	19.661778	3.102889	2.555556	0.158048
3	Sat	Yes	21.276667	2.875476	2.476190	0.147906
4	Sun	No	20.506667	3.167895	2.929825	0.160113
5	Sun	Yes	24.120000	3.516842	2.578947	0.187250
6	Thur	No	17.113111	2.673778	2.488889	0.160298
7	Thur	Yes	19.190588	3.030000	2.352941	0.163863

Of course, it's always possible to obtain the result in this format by calling `reset_index` on the result. Using the `as_index=False` argument avoids some unnecessary computations.

10.3 Apply: General split-apply-combine

The most general-purpose GroupBy method is `apply`, which is the subject of this section. `apply` splits the object being manipulated into pieces, invokes the passed function on each piece, and then attempts to concatenate the pieces.

Returning to the tipping dataset from before, suppose you wanted to select the top five `tip_pct` values by group. First, write a function that selects the rows with the largest values in a particular column:

```
In [82]: def top(df, n=5, column="tip_pct"):  
....:     return df.sort_values(column, ascending=False)[:n]
```

```
In [83]: top(tips, n=6)
```

```
Out[83]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
232	11.61	3.39	No	Sat	Dinner	2	0.291990

```
183  23.17 6.50 Yes Sun Dinner 4 0.280535
109  14.31 4.00 Yes Sat Dinner 2 0.279525
```

Now, if we group by smoker, say, and call apply with this function, we get the following:

In [84]: `tips.groupby("smoker").apply(top)`

Out[84]:

```
total_bill tip smoker day time size tip_pct
smoker
No   232   11.61 3.39 No Sat Dinner 2 0.291990
149   7.51 2.00 No Thur Lunch 2 0.266312
51    10.29 2.60 No Sun Dinner 2 0.252672
185   20.69 5.00 No Sun Dinner 5 0.241663
88    24.71 5.85 No Thur Lunch 2 0.236746
Yes  172   7.25 5.15 Yes Sun Dinner 2 0.710345
178   9.60 4.00 Yes Sun Dinner 2 0.416667
67    3.07 1.00 Yes Sat Dinner 1 0.325733
183   23.17 6.50 Yes Sun Dinner 4 0.280535
109   14.31 4.00 Yes Sat Dinner 2 0.279525
```

What has happened here? First, the tips DataFrame is split into groups based on the value of smoker. Then the top function is called on each group, and the results of each function call are glued together using pandas.concat, labeling the pieces with the group names. The result therefore has a hierarchical index with an inner level that contains index values from the original DataFrame.

If you pass a function to apply that takes other arguments or keywords, you can pass these after the function:

In [85]: `tips.groupby(["smoker", "day"]).apply(top, n=1, column="total_bill")`

Out[85]:

```
total_bill tip smoker day time size tip_pct
```

smoker day

```
No   Fri 94    22.75 3.25  No   Fri Dinner  2 0.142857
      Sat 212    48.33 9.00  No   Sat Dinner  4 0.186220
      Sun 156    48.17 5.00  No   Sun Dinner  6 0.103799
      Thur 142   41.19 5.00  No Thur Lunch   5 0.121389
Yes   Fri 95    40.17 4.73  Yes  Fri Dinner  4 0.117750
      Sat 170    50.81 10.00 Yes  Sat Dinner  3 0.196812
      Sun 182    45.35 3.50  Yes  Sun Dinner  3 0.077178
      Thur 197   43.11 5.00  Yes Thur Lunch   4 0.115982
```

Beyond these basic usage mechanics, getting the most out of apply may require some creativity. What occurs inside the function passed is up to you; it must either return a pandas object or a scalar value. The rest of this chapter will consist mainly of examples showing you how to solve various problems using groupby.

For example, you may recall that I earlier called describe on a GroupBy object:

```
In [86]: result = tips.groupby("smoker")["tip_pct"].describe()
```

```
In [87]: result
```

```
Out[87]:
```

```
count    mean    std    min    25%    50%    75% \
smoker
No    151.0  0.159328 0.039910 0.056797 0.136906 0.155625 0.185014
Yes   93.0   0.163196 0.085119 0.035638 0.106771 0.153846 0.195059
max
smoker
No    0.291990
Yes   0.710345
```

```
In [88]: result.unstack("smoker")
```

```
Out[88]:
```

```
smoker

count    No      151.000000
       Yes     93.000000
mean    No      0.159328
       Yes     0.163196
std     No      0.039910
       Yes     0.085119
min    No      0.056797
       Yes     0.035638
25%   No      0.136906
       Yes     0.106771
50%   No      0.155625
       Yes     0.153846
75%   No      0.185014
       Yes     0.195059
max    No      0.291990
       Yes     0.710345

dtype: float64
```

Inside GroupBy, when you invoke a method like describe, it is actually just a shortcut for:

```
def f(group):
    return group.describe()

grouped.apply(f)
```

Suppressing the Group Keys

In the preceding examples, you see that the resulting object has a hierarchical index formed from the group keys, along with the indexes of each piece of the original object. You can disable this by passing `group_keys=False` to `groupby`:

```
In [89]: tips.groupby("smoker", group_keys=False).apply(top)
```

Out[89]:

	total_bill	tip	smoker	day	time	size	tip_pct
232	11.61	3.39	No	Sat	Dinner	2	0.291990
149	7.51	2.00	No	Thur	Lunch	2	0.266312
51	10.29	2.60	No	Sun	Dinner	2	0.252672
185	20.69	5.00	No	Sun	Dinner	5	0.241663
88	24.71	5.85	No	Thur	Lunch	2	0.236746
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
109	14.31	4.00	Yes	Sat	Dinner	2	0.279525

Quantile and Bucket Analysis

As you may recall from [Chapter 8](#), pandas has some tools, in particular `pandas.cut` and `pandas.qcut`, for slicing data up into buckets with bins of your choosing, or by sample quantiles. Combining these functions with `groupby` makes it convenient to perform bucket or quantile analysis on a dataset. Consider a simple random dataset and an equal-length bucket categorization using `pandas.cut`:

```
In [90]: frame = pd.DataFrame({"data1": np.random.standard_normal(1000),  
....: "data2": np.random.standard_normal(1000)})
```

In [91]: `frame.head()`

Out[91]:

	data1	data2
0	-0.660524	-0.612905
1	0.862580	0.316447

```
2 -0.010032 0.838295
3 0.050009 -1.034423
4 0.670216 0.434304
```

In [92]: quartiles = pd.cut(frame["data1"], 4)

In [93]: quartiles.head(10)

Out[93]:

```
0 (-1.23, 0.489]
1 (0.489, 2.208]
2 (-1.23, 0.489]
3 (-1.23, 0.489]
4 (0.489, 2.208]
5 (0.489, 2.208]
6 (-1.23, 0.489]
7 (-1.23, 0.489]
8 (-2.956, -1.23]
9 (-1.23, 0.489]
```

Name: data1, dtype: category

Categories (4, interval[float64, right]): [(-2.956, -1.23] < (-1.23, 0.489] < (0.489, 2.208] <

(2.208, 3.928]]

The Categorical object returned by cut can be passed directly to groupby. So we could compute a set of group statistics for the quartiles, like so:

```
In [94]: def get_stats(group):
....:     return pd.DataFrame(
....:         {"min": group.min(), "max": group.max(),
....:          "count": group.count(), "mean": group.mean()})
```

```
....: )
```

```
In [95]: grouped = frame.groupby(quartiles)
```

```
In [96]: grouped.apply(get_stats)
```

```
Out[96]:
```

	min	max	count	mean
data1				
(-2.956, -1.23]	data1	-2.949343	-1.230179	94 -1.658818
	data2	-3.399312	1.670835	94 -0.033333
(-1.23, 0.489]	data1	-1.228918	0.488675	598 -0.329524
	data2	-2.989741	3.260383	598 -0.002622
(0.489, 2.208]	data1	0.489965	2.200997	298 1.065727
	data2	-3.745356	2.954439	298 0.078249
(2.208, 3.928]	data1	2.212303	3.927528	10 2.644253
	data2	-1.929776	1.765640	10 0.024750

Keep in mind the same result could have been computed more simply with:

```
In [97]: grouped.agg(["min", "max", "count", "mean"])
```

```
Out[97]:
```

	data1		data2		\		
	min	max	count	mean	min	max	count
data1							
(-2.956, -1.23]	-2.949343	-1.230179	94	-1.658818	-3.399312	1.670835	94
(-1.23, 0.489]	-1.228918	0.488675	598	-0.329524	-2.989741	3.260383	598
(0.489, 2.208]	0.489965	2.200997	298	1.065727	-3.745356	2.954439	298
(2.208, 3.928]	2.212303	3.927528	10	2.644253	-1.929776	1.765640	10
				mean			

```
data1  
(-2.956, -1.23] -0.033333  
(-1.23, 0.489] -0.002622  
(0.489, 2.208] 0.078249  
(2.208, 3.928] 0.024750
```

These were equal-length buckets; to compute equal-size buckets based on sample quantiles, use pandas.qcut. We can pass 4 as the number of bucket compute sample quartiles, and pass labels=False to obtain just the quartile indices instead of intervals:

```
In [98]: quartiles_samp = pd.qcut(frame["data1"], 4, labels=False)
```

```
In [99]: quartiles_samp.head()
```

```
Out[99]:
```

```
0 1  
1 3  
2 2  
3 2  
4 3
```

```
Name: data1, dtype: int64
```

```
In [100]: grouped = frame.groupby(quartiles_samp)
```

```
In [101]: grouped.apply(get_stats)
```

```
Out[101]:
```

		min	max	count	mean
	data1				
0	data1	-2.949343	-0.685484	250	-1.212173
	data2	-3.399312	2.628441	250	-0.027045
1	data1	-0.683066	-0.030280	250	-0.368334
	data2	-2.630247	3.260383	250	-0.027845

```
2  data1 -0.027734 0.618965 250 0.295812
   data2 -3.056990 2.458842 250 0.014450
3  data1 0.623587 3.927528 250 1.248875
   data2 -3.745356 2.954439 250 0.115899
```

Example: Filling Missing Values with Group-Specific Values

When cleaning up missing data, in some cases you will remove data observations using dropna, but in others you may want to fill in the null (NA) values using a fixed value or some value derived from the data. fillna is the right tool to use; for example, here I fill in the null values with the mean:

```
In [102]: s = pd.Series(np.random.standard_normal(6))
```

```
In [103]: s[::2] = np.nan
```

```
In [104]: s
```

```
Out[104]:
```

```
0      NaN
1    0.227290
2      NaN
3   -2.153545
4      NaN
5   -0.375842
dtype: float64
```

```
In [105]: s.fillna(s.mean())
```

```
Out[105]:
```

```
0  -0.767366
1   0.227290
2  -0.767366
3   -2.153545
```

```
4 -0.767366
5 -0.375842
dtype: float64
```

Suppose you need the fill value to vary by group. One way to do this is to group the data and use apply with a function that calls fillna on each data chunk. Here is some sample data on US states divided into eastern and western regions:

```
In [106]: states = ["Ohio", "New York", "Vermont", "Florida",
.....:      "Oregon", "Nevada", "California", "Idaho"]
```

```
In [107]: group_key = ["East", "East", "East", "East",
.....:      "West", "West", "West", "West"]
```

```
In [108]: data = pd.Series(np.random.standard_normal(8), index=states)
```

```
In [109]: data
```

```
Out[109]:
```

```
Ohio      0.329939
New York   0.981994
Vermont    1.105913
Florida    -1.613716
Oregon     1.561587
Nevada     0.406510
California  0.359244
Idaho      -0.614436
dtype: float64
```

Let's set some values in the data to be missing:

```
In [110]: data[["Vermont", "Nevada", "Idaho"]] = np.nan
```

```
In [111]: data
```

```
Out[111]:
```

```
Ohio      0.329939  
New York  0.981994  
Vermont    NaN  
Florida   -1.613716  
Oregon    1.561587  
Nevada    NaN  
California 0.359244  
Idaho     NaN  
dtype: float64
```

```
In [112]: data.groupby(group_key).size()
```

```
Out[112]:
```

```
East    4  
West    4  
dtype: int64
```

```
In [113]: data.groupby(group_key).count()
```

```
Out[113]:
```

```
East    3  
West    2  
dtype: int64
```

```
In [114]: data.groupby(group_key).mean()
```

```
Out[114]:
```

```
East   -0.100594  
West   0.960416  
dtype: float64
```

We can fill the NA values using the group means, like so:

In [115]: def fill_mean(group):

```
.....:     return group.fillna(group.mean())
```

In [116]: data.groupby(group_key).apply(fill_mean)

Out[116]:

```
Ohio      0.329939  
New York   0.981994  
Vermont    -0.100594  
Florida    -1.613716  
Oregon     1.561587  
Nevada     0.960416  
California 0.359244  
Idaho      0.960416  
dtype: float64
```

In another case, you might have predefined fill values in your code that vary by group. Since the groups have a name attribute set internally, we can use that:

In [117]: fill_values = {"East": 0.5, "West": -1}

In [118]: def fill_func(group):

```
.....:     return group.fillna(fill_values[group.name])
```

In [119]: data.groupby(group_key).apply(fill_func)

Out[119]:

```
Ohio      0.329939
```

```
New York    0.981994
Vermont    0.500000
Florida   -1.613716
Oregon     1.561587
Nevada    -1.000000
California 0.359244
Idaho     -1.000000
dtype: float64
```

Example: Random Sampling and Permutation

Suppose you wanted to draw a random sample (with or without replacement) from a large dataset for Monte Carlo simulation purposes or some other application. There are a number of ways to perform the “draws”; here we use the sample method for Series.

To demonstrate, here’s a way to construct a deck of English-style playing cards:

```
suits = ["H", "S", "C", "D"] # Hearts, Spades, Clubs, Diamonds
```

```
card_val = (list(range(1, 11)) + [10] * 3) * 4
```

```
base_names = ["A"] + list(range(2, 11)) + ["J", "K", "Q"]
```

```
cards = []
```

```
for suit in suits:
```

```
    cards.extend(str(num) + suit for num in base_names)
```

```
deck = pd.Series(card_val, index=cards)
```

Now we have a Series of length 52 whose index contains card names, and values are the ones used in blackjack and other games (to keep things simple, I let the ace "A" be 1):

```
In [121]: deck.head(13)
```

```
Out[121]:
```

```
AH    1
```

```
2H 2
3H 3
4H 4
5H 5
6H 6
7H 7
8H 8
9H 9
10H 10
JH 10
KH 10
QH 10
dtype: int64
```

Now, based on what I said before, drawing a hand of five cards from the deck could be written as:

```
In [122]: def draw(deck, n=5):
.....:     return deck.sample(n)
```

```
In [123]: draw(deck)
```

```
Out[123]:
```

```
4D 4
QH 10
8S 8
7D 7
9C 9
dtype: int64
```

Suppose you wanted two random cards from each suit. Because the suit is the last character of each card name, we can group based on this and use apply:

In [124]: def get_suit(card):

```
.....: # last letter is suit  
.....: return card[-1]
```

In [125]: deck.groupby(get_suit).apply(draw, n=2)

Out[125]:

C 6C 6

KC 10

D 7D 7

3D 3

H 7H 7

9H 9

S 2S 2

QS 10

dtype: int64

Alternatively, we could pass group_keys=False to drop the outer suit index, leaving in just the selected cards:

In [126]: deck.groupby(get_suit, group_keys=False).apply(draw, n=2)

Out[126]:

AC 1

3C 3

5D 5

4D 4

10H 10

7H 7

QS 10

```
7S    7
```

```
dtype: int64
```

Example: Group Weighted Average and Correlation

Under the split-apply-combine paradigm of groupby, operations between columns in a DataFrame or two Series, such as a group weighted average, are possible. As an example, take this dataset containing group keys, values, and some weights:

```
In [127]: df = pd.DataFrame({"category": ["a", "a", "a", "a",
.....:                 "b", "b", "b", "b"],
.....:                 "data": np.random.standard_normal(8),
.....:                 "weights": np.random.uniform(size=8)})
```

```
In [128]: df
```

```
Out[128]:
```

	category	data	weights
0	a	-1.691656	0.955905
1	a	0.511622	0.012745
2	a	-0.401675	0.137009
3	a	0.968578	0.763037
4	b	-1.818215	0.492472
5	b	0.279963	0.832908
6	b	-0.200819	0.658331
7	b	-0.217221	0.612009

The weighted average by category would then be:

```
In [129]: grouped = df.groupby("category")
```

```
In [130]: def get_wavg(group):
```

```
.....:     return np.average(group["data"], weights=group["weights"])
```

In [131]: grouped.apply(get_wavg)

Out[131]:

category

a -0.495807

b -0.357273

dtype: float64

As another example, consider a financial dataset originally obtained from Yahoo! Finance containing end-of-day prices for a few stocks and the S&P 500 index (the SPX symbol):

In [132]: close_px = pd.read_csv("examples/stock_px.csv", parse_dates=True,

```
.....:             index_col=0)
```

In [133]: close_px.info()

<class 'pandas.core.frame.DataFrame'>

DatetimeIndex: 2214 entries, 2003-01-02 to 2011-10-14

Data columns (total 4 columns):

Column Non-Null Count Dtype

0 AAPL 2214 non-null float64

1 MSFT 2214 non-null float64

2 XOM 2214 non-null float64

3 SPX 2214 non-null float64

dtypes: float64(4)

memory usage: 86.5 KB

In [134]: close_px.tail(4)

Out[134]:

```
AAPL MSFT XOM SPX  
2011-10-11 400.29 27.00 76.27 1195.54  
2011-10-12 402.19 26.96 77.16 1207.25  
2011-10-13 408.43 27.18 76.37 1203.66  
2011-10-14 422.00 27.27 78.11 1224.58
```

The DataFrame info() method here is a convenient way to get an overview of the contents of a DataFrame.

One task of interest might be to compute a DataFrame consisting of the yearly correlations of daily returns (computed from percent changes) with SPX. As one way to do this, we first create a function that computes the pair-wise correlation of each column with the "SPX" column:

```
In [135]: def spx_corr(group):  
.....:     return group.corrwith(group["SPX"])
```

Next, we compute percent change on close_px using pct_change:

```
In [136]: rets = close_px.pct_change().dropna()
```

Lastly, we group these percent changes by year, which can be extracted from each row label with a one-line function that returns the year attribute of each datetime label:

```
In [137]: def get_year(x):  
.....:     return x.year
```

```
In [138]: by_year = rets.groupby(get_year)
```

```
In [139]: by_year.apply(spx_corr)
```

Out[139]:

```
AAPL  MSFT  XOM  SPX

2003 0.541124 0.745174 0.661265 1.0
2004 0.374283 0.588531 0.557742 1.0
2005 0.467540 0.562374 0.631010 1.0
2006 0.428267 0.406126 0.518514 1.0
2007 0.508118 0.658770 0.786264 1.0
2008 0.681434 0.804626 0.828303 1.0
2009 0.707103 0.654902 0.797921 1.0
2010 0.710105 0.730118 0.839057 1.0
2011 0.691931 0.800996 0.859975 1.0
```

You could also compute intercolumn correlations. Here we compute the annual correlation between Apple and Microsoft:

```
In [140]: def corr_aapl_msft(group):
.....:     return group["AAPL"].corr(group["MSFT"])
```

```
In [141]: by_year.apply(corr_aapl_msft)
```

```
Out[141]:
```

```
2003  0.480868
2004  0.259024
2005  0.300093
2006  0.161735
2007  0.417738
2008  0.611901
2009  0.432738
2010  0.571946
2011  0.581987
dtype: float64
```

Example: Group-Wise Linear Regression

In the same theme as the previous example, you can use groupby to perform more complex group-wise statistical analysis, as long as the function returns a pandas object or scalar value. For example, I can define the following regress function (using the statsmodels econometrics library), which executes an ordinary least squares (OLS) regression on each chunk of data:

```
import statsmodels.api as sm

def regress(data, yvar=None, xvars=None):
    Y = data[yvar]
    X = data[xvars]
    X["intercept"] = 1.
    result = sm.OLS(Y, X).fit()
    return result.params
```

You can install statsmodels with conda if you don't have it already:

```
conda install statsmodels
```

Now, to run a yearly linear regression of AAPL on SPX returns, execute:

```
In [143]: by_year.apply(regress, yvar="AAPL", xvars=["SPX"])
```

```
Out[143]:
```

```
SPX intercept
2003 1.195406 0.000710
2004 1.363463 0.004201
2005 1.766415 0.003246
2006 1.645496 0.000080
2007 1.198761 0.003438
```

```
2008 0.968016 -0.001110
2009 0.879103 0.002954
2010 1.052608 0.001261
2011 0.806605 0.001514
```

10.4 Group Transforms and “Unwrapped” GroupBys

In [Section 10.3, “Apply: General split-apply-combine,”](#) we looked at the `apply` method in grouped operations for performing transformations. There is another built-in method called `transform`, which is similar to `apply` but imposes more constraints on the kind of function you can use:

- It can produce a scalar value to be broadcast to the shape of the group.
- It can produce an object of the same shape as the input group.
- It must not mutate its input.

Let’s consider a simple example for illustration:

```
In [144]: df = pd.DataFrame({'key': ['a', 'b', 'c'] * 4,
.....:     'value': np.arange(12.)})
```

```
In [145]: df
```

```
Out[145]:
```

	key	value
0	a	0.0
1	b	1.0
2	c	2.0
3	a	3.0
4	b	4.0
5	c	5.0
6	a	6.0
7	b	7.0
8	c	8.0
9	a	9.0

```
10 b 10.0  
11 c 11.0
```

Here are the group means by key:

```
In [146]: g = df.groupby('key')['value']
```

```
In [147]: g.mean()
```

```
Out[147]:
```

```
key  
a    4.5  
b    5.5  
c    6.5
```

```
Name: value, dtype: float64
```

Suppose instead we wanted to produce a Series of the same shape as df['value'] but with values replaced by the average grouped by 'key'. We can pass a function that computes the mean of a single group to transform:

```
In [148]: def get_mean(group):
```

```
.....:     return group.mean()
```

```
In [149]: g.transform(get_mean)
```

```
Out[149]:
```

```
0    4.5  
1    5.5  
2    6.5  
3    4.5  
4    5.5  
5    6.5
```

```
6  4.5
7  5.5
8  6.5
9  4.5
10 5.5
11 6.5
```

Name: value, dtype: float64

For built-in aggregation functions, we can pass a string alias as with the GroupBy agg method:

In [150]: g.transform('mean')

Out[150]:

```
0  4.5
1  5.5
2  6.5
3  4.5
4  5.5
5  6.5
6  4.5
7  5.5
8  6.5
9  4.5
10 5.5
11 6.5
```

Name: value, dtype: float64

Like apply, transform works with functions that return Series, but the result must be the same size as the input. For example, we can multiply each group by 2 using a helper function:

```
In [151]: def times_two(group):
```

```
.....:     return group * 2
```

```
In [152]: g.transform(times_two)
```

```
Out[152]:
```

```
0    0.0
```

```
1    2.0
```

```
2    4.0
```

```
3    6.0
```

```
4    8.0
```

```
5   10.0
```

```
6   12.0
```

```
7   14.0
```

```
8   16.0
```

```
9   18.0
```

```
10  20.0
```

```
11  22.0
```

```
Name: value, dtype: float64
```

As a more complicated example, we can compute the ranks in descending order for each group:

```
In [153]: def get_ranks(group):
```

```
.....:     return group.rank(ascending=False)
```

```
In [154]: g.transform(get_ranks)
```

```
Out[154]:
```

```
0    4.0
```

```
1 4.0
2 4.0
3 3.0
4 3.0
5 3.0
6 2.0
7 2.0
8 2.0
9 1.0
10 1.0
11 1.0
```

Name: value, dtype: float64

Consider a group transformation function composed from simple aggregations:

```
In [155]: def normalize(x):
.....:     return (x - x.mean()) / x.std()
```

We can obtain equivalent results in this case using either transform or apply:

```
In [156]: g.transform(normalize)
```

Out[156]:

```
0 -1.161895
1 -1.161895
2 -1.161895
3 -0.387298
4 -0.387298
5 -0.387298
6 0.387298
```

```
7  0.387298
8  0.387298
9  1.161895
10 1.161895
11 1.161895

Name: value, dtype: float64
```

In [157]: g.apply(normalize)

Out[157]:

```
0 -1.161895
1 -1.161895
2 -1.161895
3 -0.387298
4 -0.387298
5 -0.387298
6  0.387298
7  0.387298
8  0.387298
9  1.161895
10 1.161895
11 1.161895
```

Name: value, dtype: float64

Built-in aggregate functions like 'mean' or 'sum' are often much faster than a general apply function. These also have a "fast path" when used with transform. This allows us to perform what is called an *unwrapped* group operation:

In [158]: g.transform('mean')

Out[158]:

```
0  4.5
1  5.5
```

```
2  6.5
3  4.5
4  5.5
5  6.5
6  4.5
7  5.5
8  6.5
9  4.5
10 5.5
11 6.5
```

Name: value, dtype: float64

In [159]: normalized = (df['value'] - g.transform('mean')) / g.transform('std')

In [160]: normalized

Out[160]:

```
0 -1.161895
1 -1.161895
2 -1.161895
3 -0.387298
4 -0.387298
5 -0.387298
6  0.387298
7  0.387298
8  0.387298
9  1.161895
10 1.161895
11 1.161895
```

Name: value, dtype: float64

Here, we are doing arithmetic between the outputs of multiple GroupBy operations instead of writing a function and passing it to groupby(...).apply. That is what is meant by “unwrapped.”

While an unwrapped group operation may involve multiple group aggregations, the overall benefit of vectorized operations often outweighs this.

10.5 Pivot Tables and Cross-Tabulation

A *pivot table* is a data summarization tool frequently found in spreadsheet programs and other data analysis software. It aggregates a table of data by one or more keys, arranging the data in a rectangle with some of the group keys along the rows and some along the columns. Pivot tables in Python with pandas are made possible through the groupby facility described in this chapter, combined with reshape operations utilizing hierarchical indexing. DataFrame also has a pivot_table method, and there is also a top-level pandas.pivot_table function. In addition to providing a convenience interface to groupby, pivot_table can add partial totals, also known as *margins*.

Returning to the tipping dataset, suppose you wanted to compute a table of group means (the default pivot_table aggregation type) arranged by day and smoker on the rows:

In [161]: tips.head()

Out[161]:

	total_bill	tip	smoker	day	time	size	tip_pct
0	16.99	1.01	No	Sun	Dinner	2	0.059447
1	10.34	1.66	No	Sun	Dinner	3	0.160542
2	21.01	3.50	No	Sun	Dinner	3	0.166587
3	23.68	3.31	No	Sun	Dinner	2	0.139780
4	24.59	3.61	No	Sun	Dinner	4	0.146808

In [162]: tips.pivot_table(index=["day", "smoker"])

Out[162]:

	size	tip	tip_pct	total_bill	
day	smoker				
Fri	No	2.250000	2.812500	0.151650	18.420000
	Yes	2.066667	2.714000	0.174783	16.813333
Sat	No	2.555556	3.102889	0.158048	19.661778
	Yes	2.476190	2.875476	0.147906	21.276667
Sun	No	2.929825	3.167895	0.160113	20.506667

```
Yes  2.578947 3.516842 0.187250 24.120000
Thur No  2.488889 2.673778 0.160298 17.113111
Yes  2.352941 3.030000 0.163863 19.190588
```

This could have been produced with groupby directly, using `tips.groupby(["day", "smoker"]).mean()`. Now, suppose we want to take the average of only tip_pct and size, and additionally group by time. I'll put smoker in the table columns and time and day in the rows:

```
In [163]: tips.pivot_table(index=["time", "day"], columns="smoker",
.....:           values=["tip_pct", "size"])
```

Out[163]:

	size		tip_pct		
smoker	No	Yes	No	Yes	
time	day				
Dinner	Fri	2.000000	2.222222	0.139622	0.165347
Sat		2.555556	2.476190	0.158048	0.147906
Sun		2.929825	2.578947	0.160113	0.187250
Thur		2.000000	NaN	0.159744	NaN
Lunch	Fri	3.000000	1.833333	0.187735	0.188937
Thur		2.500000	2.352941	0.160311	0.163863

We could augment this table to include partial totals by passing `margins=True`. This has the effect of adding All row and column labels, with corresponding values being the group statistics for all the data within a single tier:

```
In [164]: tips.pivot_table(index=["time", "day"], columns="smoker",
.....:           values=["tip_pct", "size"], margins=True)
```

Out[164]:

	size		tip_pct			
smoker	No	Yes	All	No	Yes	All

```

time day

Dinner Fri 2.000000 2.222222 2.166667 0.139622 0.165347 0.158916
      Sat 2.555556 2.476190 2.517241 0.158048 0.147906 0.153152
      Sun 2.929825 2.578947 2.842105 0.160113 0.187250 0.166897
      Thur 2.000000   NaN 2.000000 0.159744   NaN 0.159744
Lunch Fri 3.000000 1.833333 2.000000 0.187735 0.188937 0.188765
      Thur 2.500000 2.352941 2.459016 0.160311 0.163863 0.161301
All     2.668874 2.408602 2.569672 0.159328 0.163196 0.160803

```

Here, the All values are means without taking into account smoker versus non-smoker (the All columns) or any of the two levels of grouping on the rows (the All row).

To use an aggregation function other than mean, pass it to the aggfunc keyword argument. For example, "count" or len will give you a cross-tabulation (count or frequency) of group sizes (though "count" will exclude null values from the count within data groups, while len will not):

```
In [165]: tips.pivot_table(index=["time", "smoker"], columns="day",
```

```
.....:      values="tip_pct", aggfunc=len, margins=True)
```

```
Out[165]:
```

	day	Fri	Sat	Sun	Thur	All
time	smoker					
Dinner	No	3.0	45.0	57.0	1.0	106
	Yes	9.0	42.0	19.0	NaN	70
Lunch	No	1.0	NaN	NaN	44.0	45
	Yes	6.0	NaN	NaN	17.0	23
All		19.0	87.0	76.0	62.0	244

If some combinations are empty (or otherwise NA), you may wish to pass a fill_value:

```
In [166]: tips.pivot_table(index=["time", "size", "smoker"], columns="day",
```

```

.....:      values="tip_pct", fill_value=0)

Out[166]:
day       Fri   Sat   Sun   Thur
time  size smoker
Dinner 1  No    0.000000 0.137931 0.000000 0.000000
          Yes   0.000000 0.325733 0.000000 0.000000
          2  No    0.139622 0.162705 0.168859 0.159744
          Yes   0.171297 0.148668 0.207893 0.000000
          3  No    0.000000 0.154661 0.152663 0.000000
...
          ...   ...   ...
Lunch  3  Yes   0.000000 0.000000 0.000000 0.204952
          4  No    0.000000 0.000000 0.000000 0.138919
          Yes   0.000000 0.000000 0.000000 0.155410
          5  No    0.000000 0.000000 0.000000 0.121389
          6  No    0.000000 0.000000 0.000000 0.173706
[21 rows x 4 columns]

```

See [Table 10-2](#) for a summary of pivot_table options.

Argument	Description
values	Column name or names to aggregate; by default, aggregates all numeric columns
index	Column names or other group keys to group on the rows of the resulting pivot table
columns	Column names or other group keys to group on the columns of the resulting pivot table

Argument	Description
aggfunc	Aggregation function or list of functions ("mean" by default); can be any function valid in a groupby context
fill_value	Replace missing values in the result table
dropna	If True, do not include columns whose entries are all NA
margins	Add row/column subtotals and grand total (False by default)
margins_name	Name to use for the margin row/column labels when passing margins=True; defaults to "All"
observed	With Categorical group keys, if True, show only the observed category values in the keys rather than all categories

Table 10-2. pivot_table options

Cross-Tabulations: Crosstab

A *cross-tabulation* (or *crosstab* for short) is a special case of a pivot table that computes group frequencies. Here is an example:

In [167]: from io import StringIO

In [168]: data = """Sample Nationality Handedness

```
.....: 1 USA Right-handed
.....: 2 Japan Left-handed
.....: 3 USA Right-handed
.....: 4 Japan Right-handed
.....: 5 Japan Left-handed
.....: 6 Japan Right-handed
.....: 7 USA Right-handed
.....: 8 USA Left-handed
```

```
.....: 9 Japan Right-handed  
.....: 10 USA Right-handed"""  
.....:
```

```
In [169]: data = pd.read_table(StringIO(data), sep="\s+")
```

```
In [170]: data
```

```
Out[170]:
```

```
Sample Nationality Handedness  
0 1 USA Right-handed  
1 2 Japan Left-handed  
2 3 USA Right-handed  
3 4 Japan Right-handed  
4 5 Japan Left-handed  
5 6 Japan Right-handed  
6 7 USA Right-handed  
7 8 USA Left-handed  
8 9 Japan Right-handed  
9 10 USA Right-handed
```

As part of some survey analysis, we might want to summarize this data by nationality and handedness. You could use `pivot_table` to do this, but the `pandas.crosstab` function can be more convenient:

```
In [171]: pd.crosstab(data["Nationality"], data["Handedness"], margins=True)
```

```
Out[171]:
```

```
Handedness Left-handed Right-handed All  
Nationality  
Japan 2 3 5
```

USA	1	4	5
All	3	7	10

The first two arguments to crosstab can each be an array or Series or a list of arrays. As in the tips data:

```
In [172]: pd.crosstab([tips["time"], tips["day"]], tips["smoker"], margins=True)
```

```
Out[172]:
```

smoker	No	Yes	All
--------	----	-----	-----

time	day
------	-----

Dinner	Fri	3	9	12
--------	-----	---	---	----

Sat	45	42	87
-----	----	----	----

Sun	57	19	76
-----	----	----	----

Thur	1	0	1
------	---	---	---

Lunch	Fri	1	6	7
-------	-----	---	---	---

Thur	44	17	61
------	----	----	----

All	151	93	244
-----	-----	----	-----

10.6 Conclusion

Mastering pandas's data grouping tools can help with data cleaning and modeling or statistical analysis work. In [Chapter 13](#) we will look at several more example use cases for groupby on real data.

In the next chapter, we turn our attention to time series data.